



This work is licensed under a [Creative Commons
Attribution 4.0 International License](#).



MACHINE LEARNING for SENSOR DATA



Faktion

Jos Polfliet, VP of Applied AI

FAKTION – we put thought in everything

**WE BUILD
DEEP LEARNING,
MACHINE LEARNING AND
ARTIFICIAL INTELLIGENCE
SOLUTIONS**

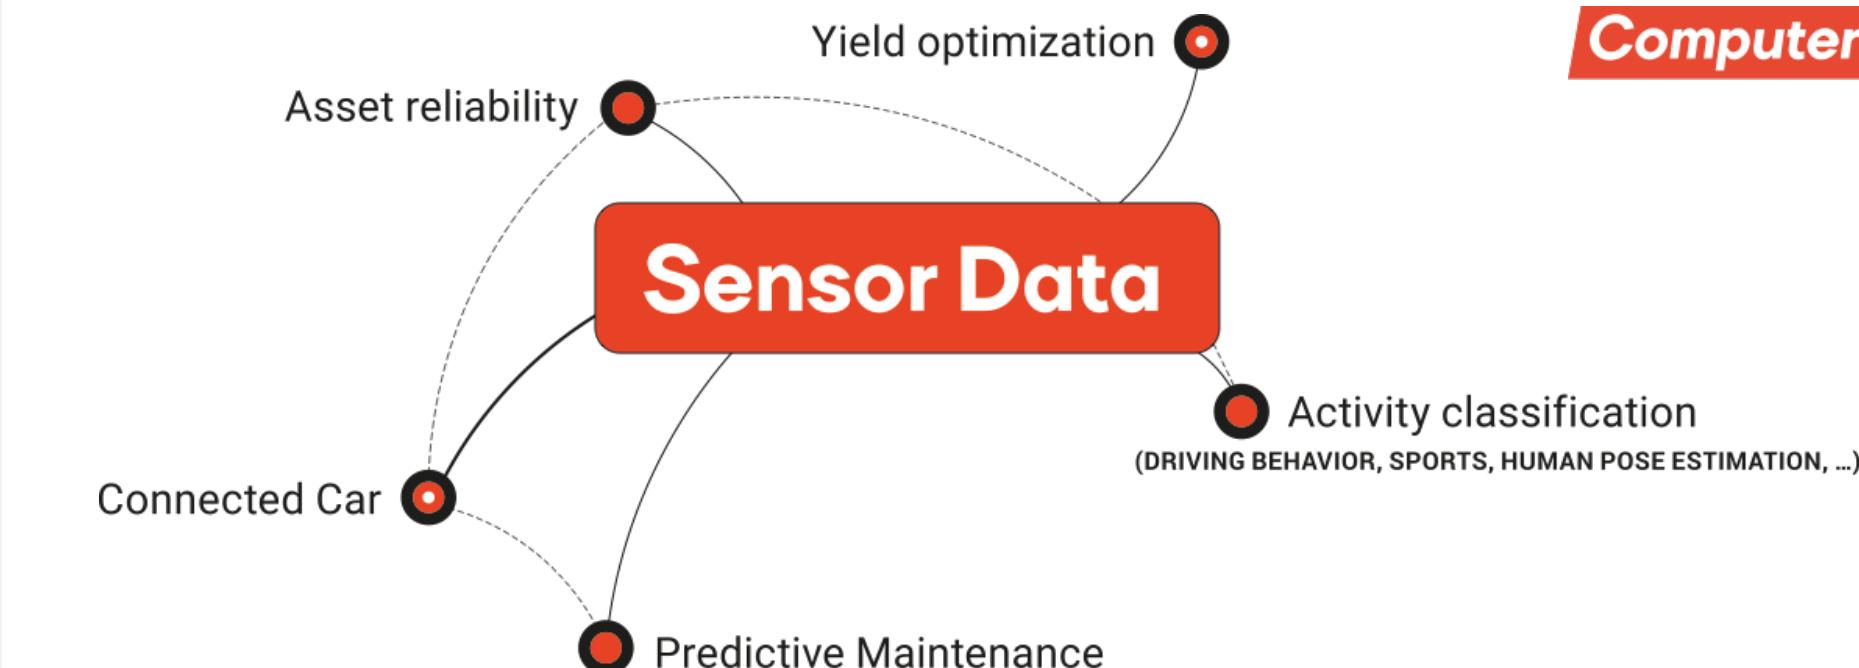
 *Artificial Intelligence is confusing. We know. Truth is, there is money to be made by selling hype. Contact us when you need a partner that delivers results instead.*

FAKTION — “we put thought in everything”



Faktion is an applied AI service provider.
We build custom software applications
with Machine and Deep Learning technology.

Due to the strong extent of our experience,
we share that by means of strategic consulting.



**We build reusable and
scalable AI sensor data
platforms that can easily
be extended to dozens of
other use cases**

HOW

CAN WE BUILD THIS



Let's use Machine Learning for sensor data.

THE FUNDAMENTAL PROBLEM

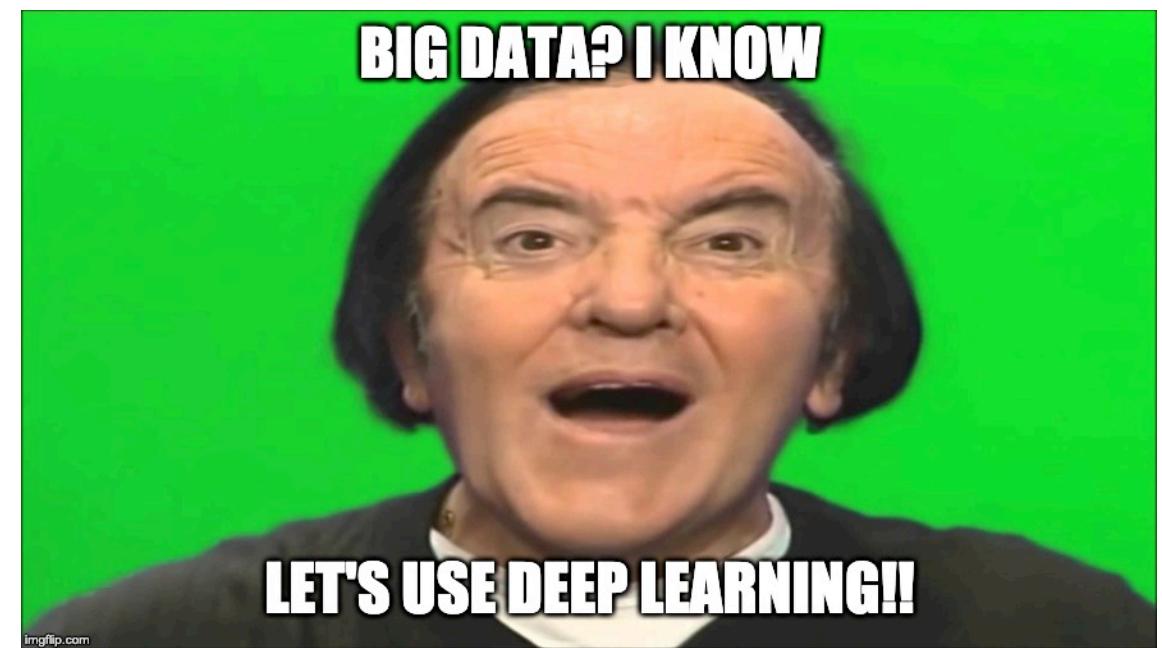
A journey of discovery

What you get

25x Temperature, 25x pressure, 25x flow at 100Hz

75 sensors * 100 Hz
* 3600 s * 24 h * 30
days * 3 months =
58,320,000,000
rows

Datetime	Sensor	Value
01/02/2019 00:00:00.00	Temperature 1	120.89121
01/02/2019 00:00:00.01	Temperature 1	120.89118
,,,		
31/05/2019 23:59:59.99	Temperature 1	116.56119
01/02/2019 00:00:00.00	Temperature 2	117.4215
01/02/2019 00:00:00.01	Temperature 2	117.4071
,,,		
31/05/2019 23:59:59.99	Temperature 2	119.14111
...



What not to do

25x Temperature, 25x pressure, 25x flow at 100Hz for 60s

$100 \text{ Hz} * 3600 \text{ s} * 75 \text{ sensors} = 27,000,000 \text{ columns}$

24 hours * 30 days
* 3 months
= 2,160 rows

Datetime	T	00	...	F_19_600	Target
01/02/2019 00:00					GOOD
01/02/2019 01:00					BAD
31/05/2019 23:00					GOOD





What would you like to learn today?

Learn ▾

Practice

Projects

Pricing

For Business

0 XP



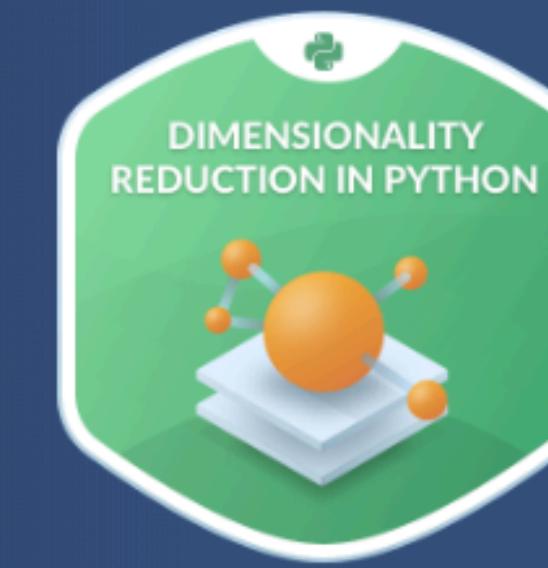
▼



INTERACTIVE COURSE

Dimensionality Reduction in Python

Start Course For Free



⌚ 4 hours | ► 16 Videos | ↻ 58 Exercises | 🚩 562 Participants | 💼 4,700 XP

Course Description

High-dimensional datasets can be overwhelming and leave you not knowing where to start. Typically, you'd visually explore a new dataset first, but when you have too many dimensions the classical approaches will seem insufficient. Fortunately, there are visualization techniques designed specifically for high dimensional data and you'll be introduced to these in this course. After exploring the data, you'll often find that many features hold little information because they don't show any variance or because they are duplicates of other features. You'll learn how to detect these features and drop them from the dataset so that you can focus on the informative ones. In a next step, you might want to build a model on these features, and it may turn out that some don't have any effect on the thing you're trying to predict. You'll learn how to detect and drop these irrelevant features too, in order to reduce dimensionality and thus complexity. Finally, you'll learn how feature extraction techniques can reduce dimensionality for you through the calculation of uncorrelated principal components.

1 Exploring high dimensional data FREE

0%

You'll be introduced to the concept of dimensionality reduction and will learn when and why this is important. You'll learn the difference between feature selection and feature extraction and will apply both techniques for data exploration. The chapter ends with a lesson on t-SNE, a powerful feature extraction technique that will allow you to visualize a high-dimensional dataset.

[VIEW CHAPTER DETAILS ▾](#)

[Continue Chapter](#)



Jeroen Boeye

Machine Learning Engineer @
Faktion

Jeroen is a machine learning engineer working at Faktion, an AI company from Belgium.

He uses both R and Python for his analyses and has a PhD background in computational biology. His experience mostly lies in working with structured data, produced by sensors or digital processes.

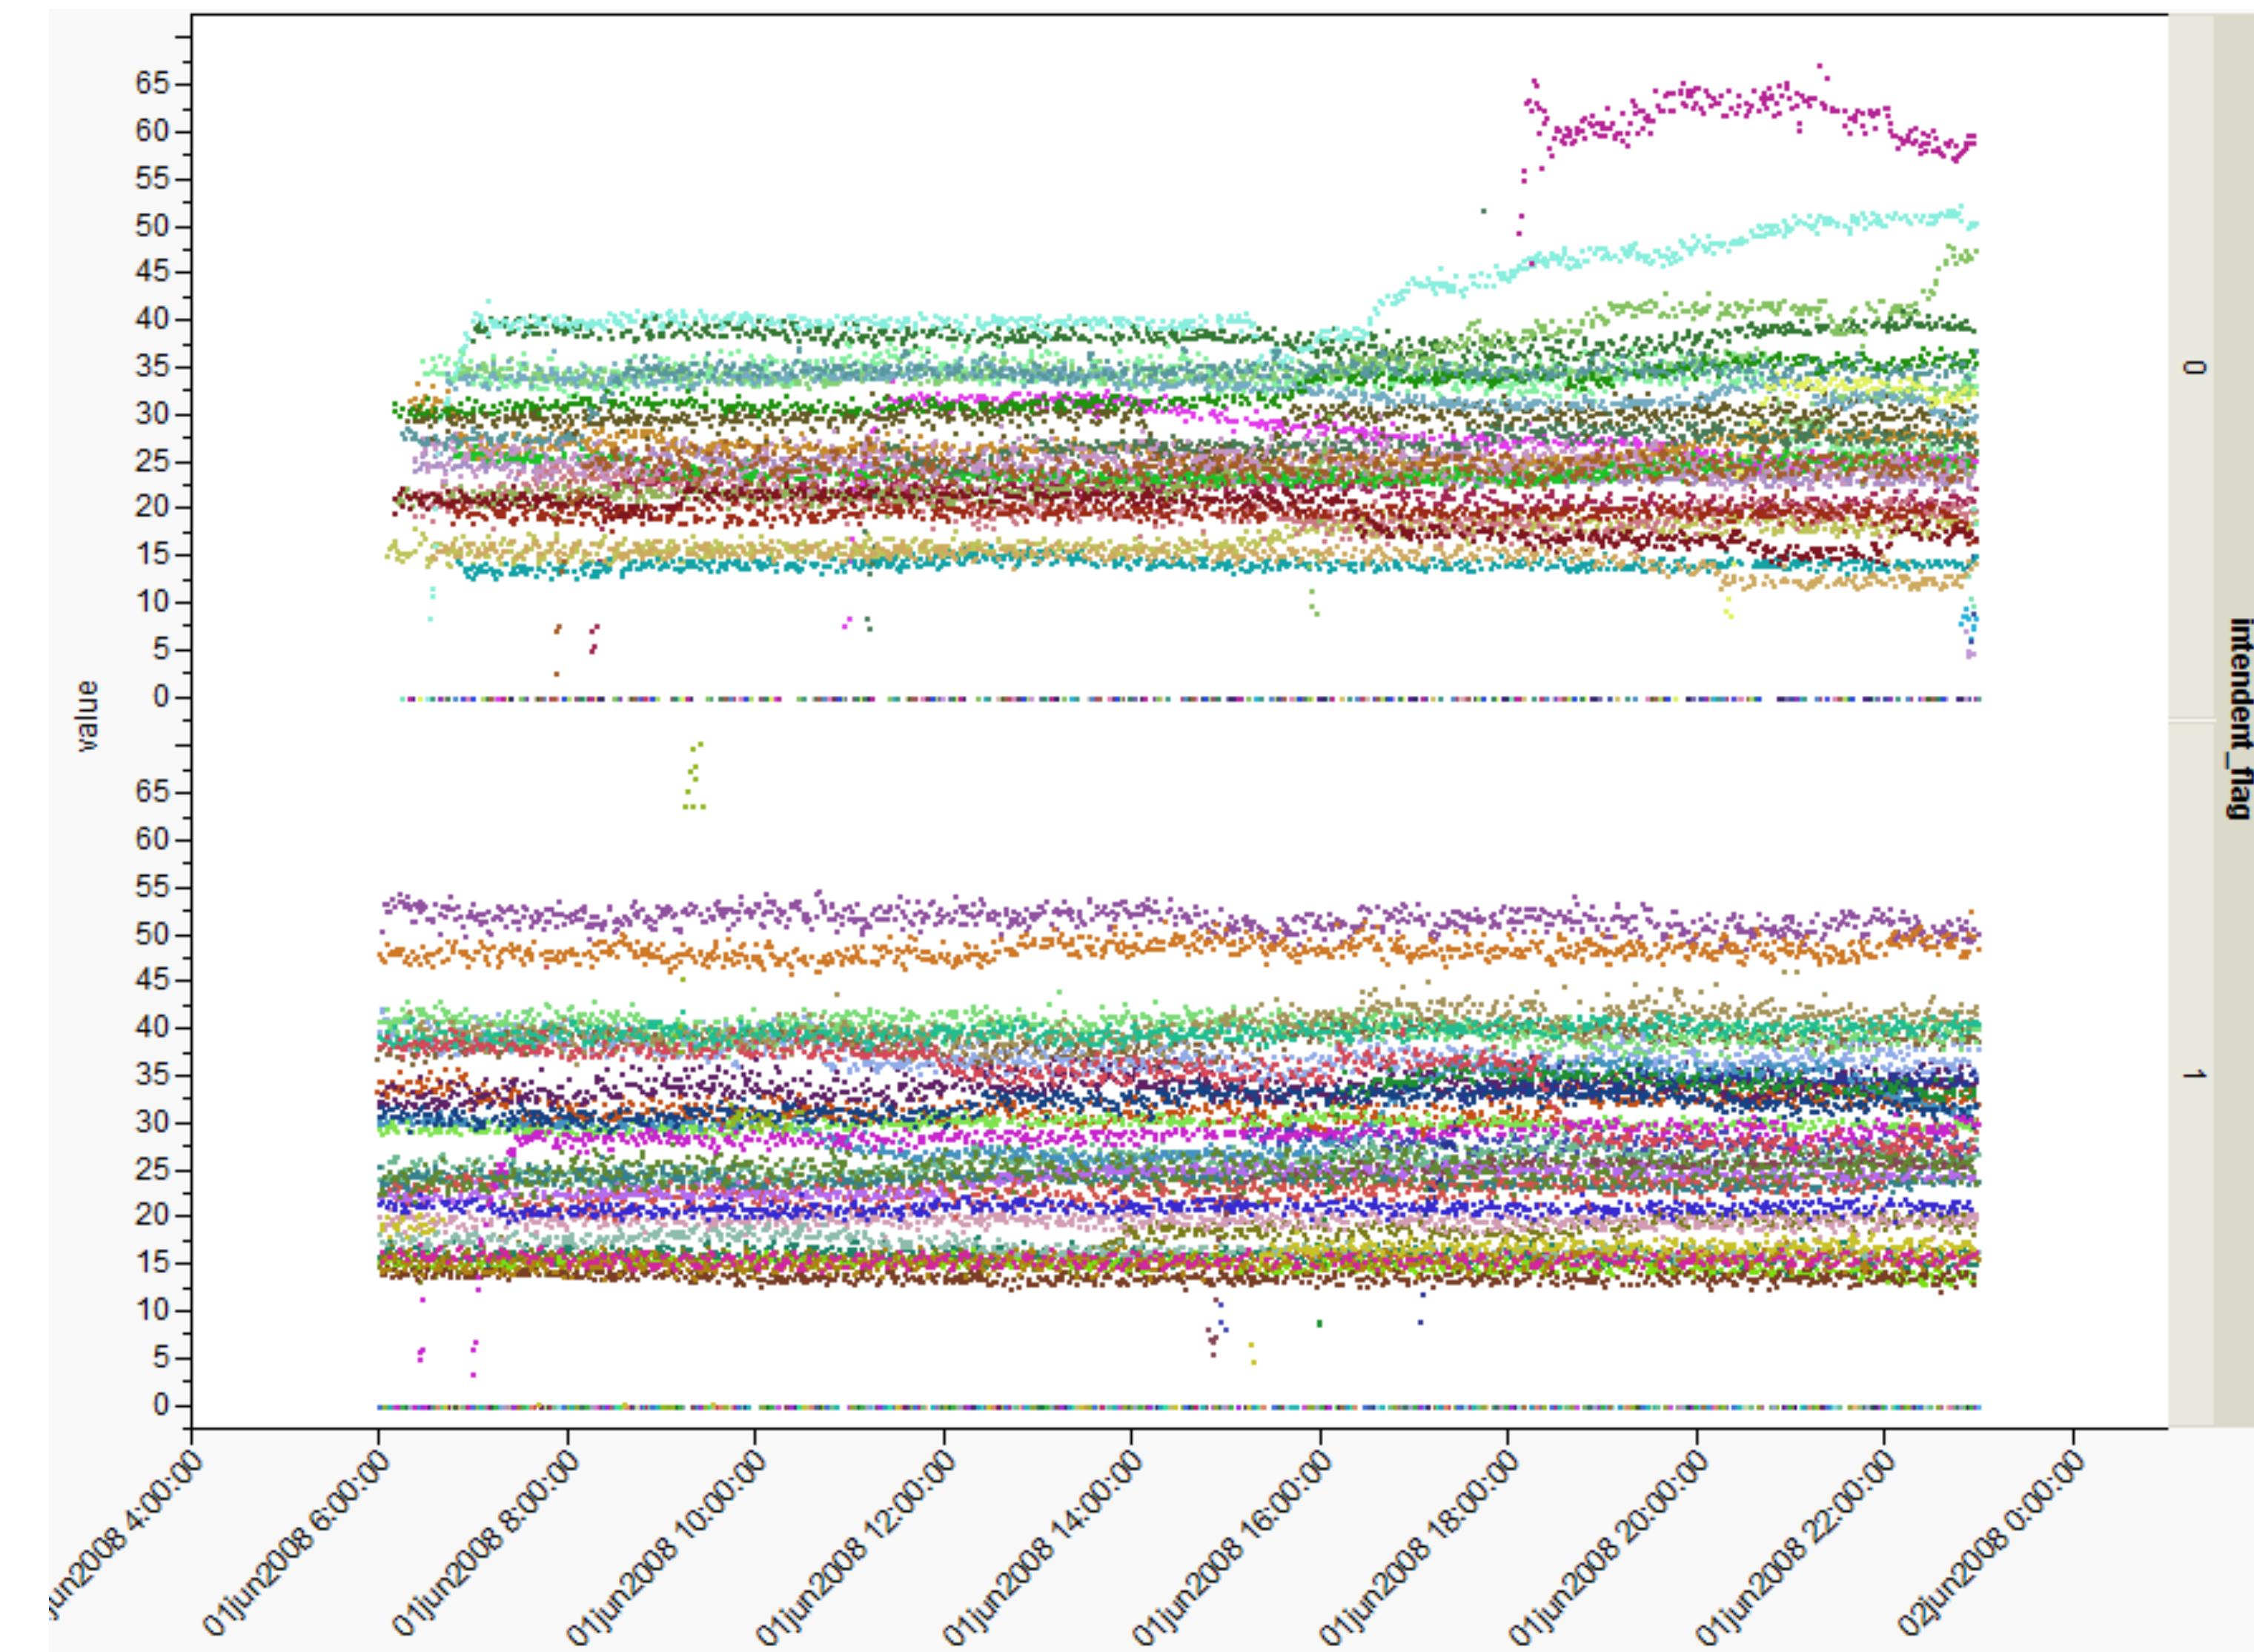
[See More](#)

Pro-tip:

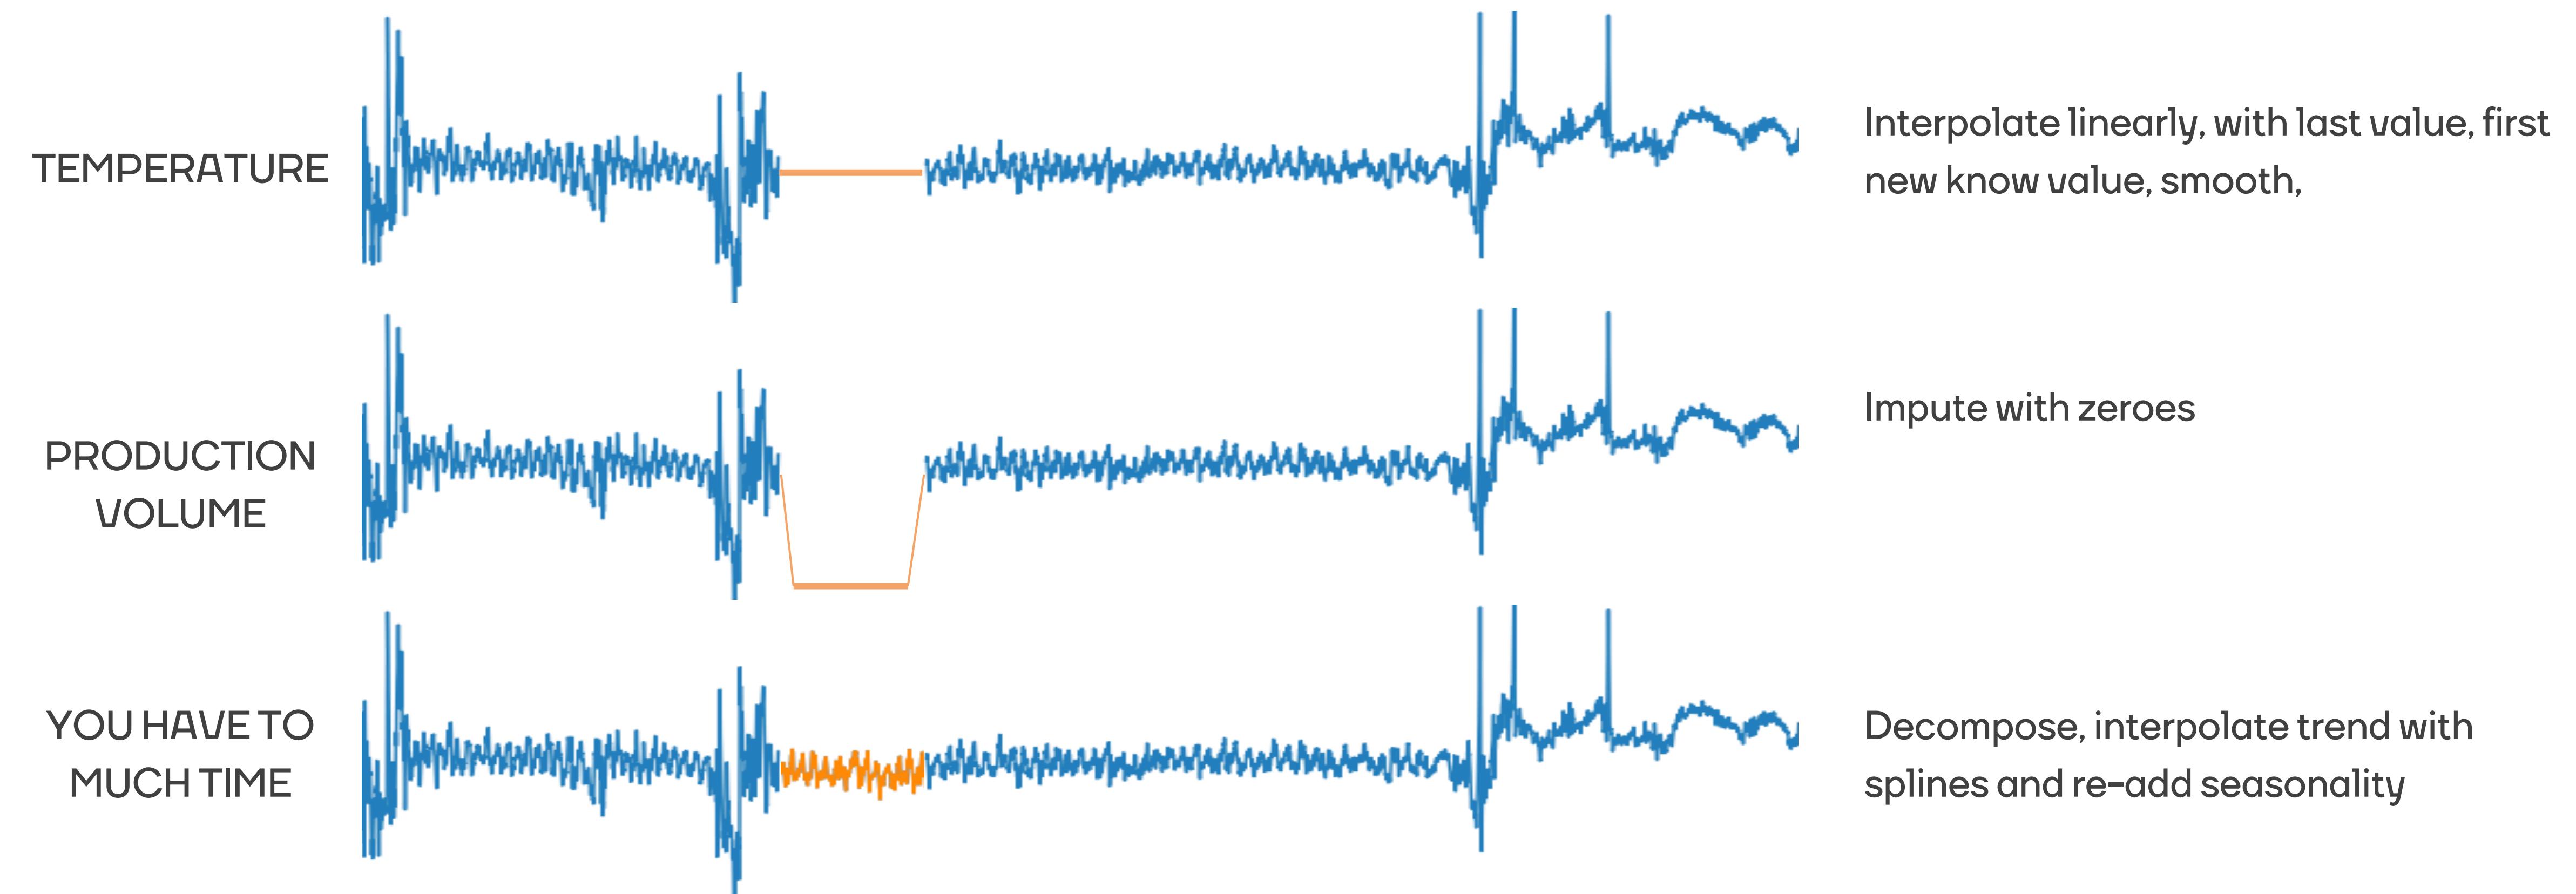
<https://www.datacamp.com/courses/dimensionality-reduction-in-python>

Time series data mining

Step 1: Prepare timeseries



Pre-processing - Interpolation



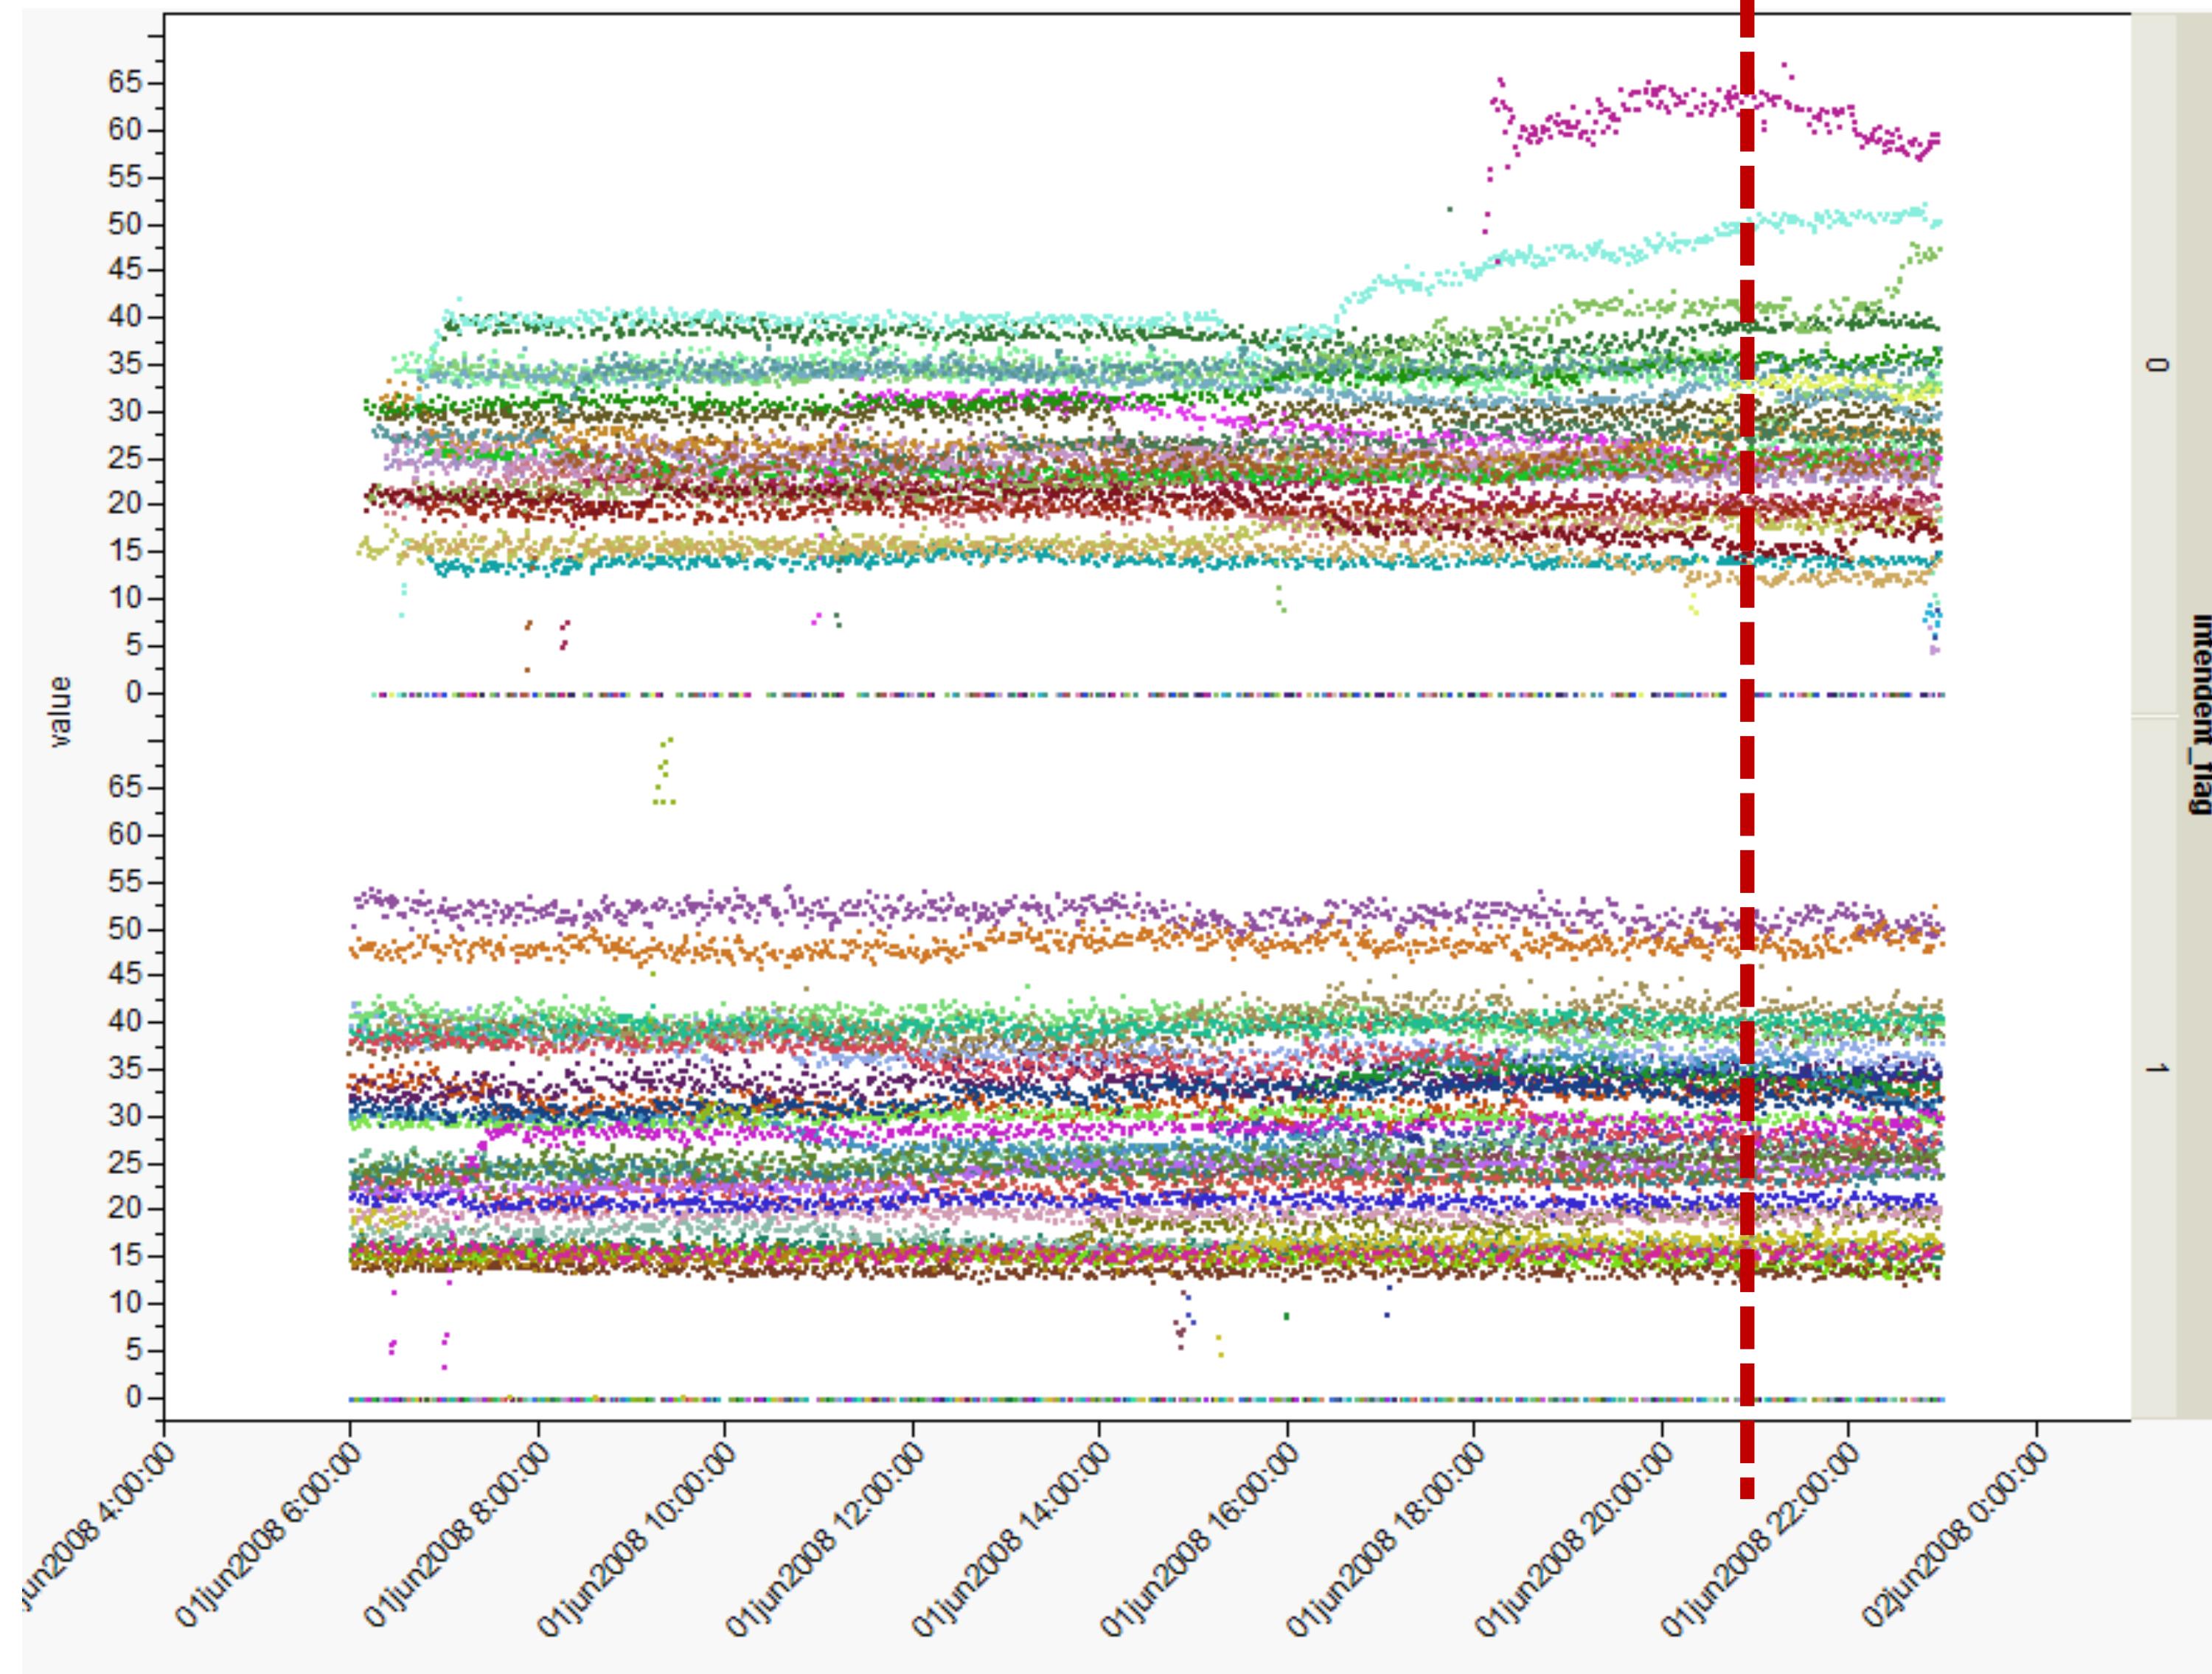
Interpolate linearly, with last value, first new know value, smooth,

Impute with zeroes

Decompose, interpolate trend with splines and re-add seasonality

Bucketization

TRIP TIME - 2 hours



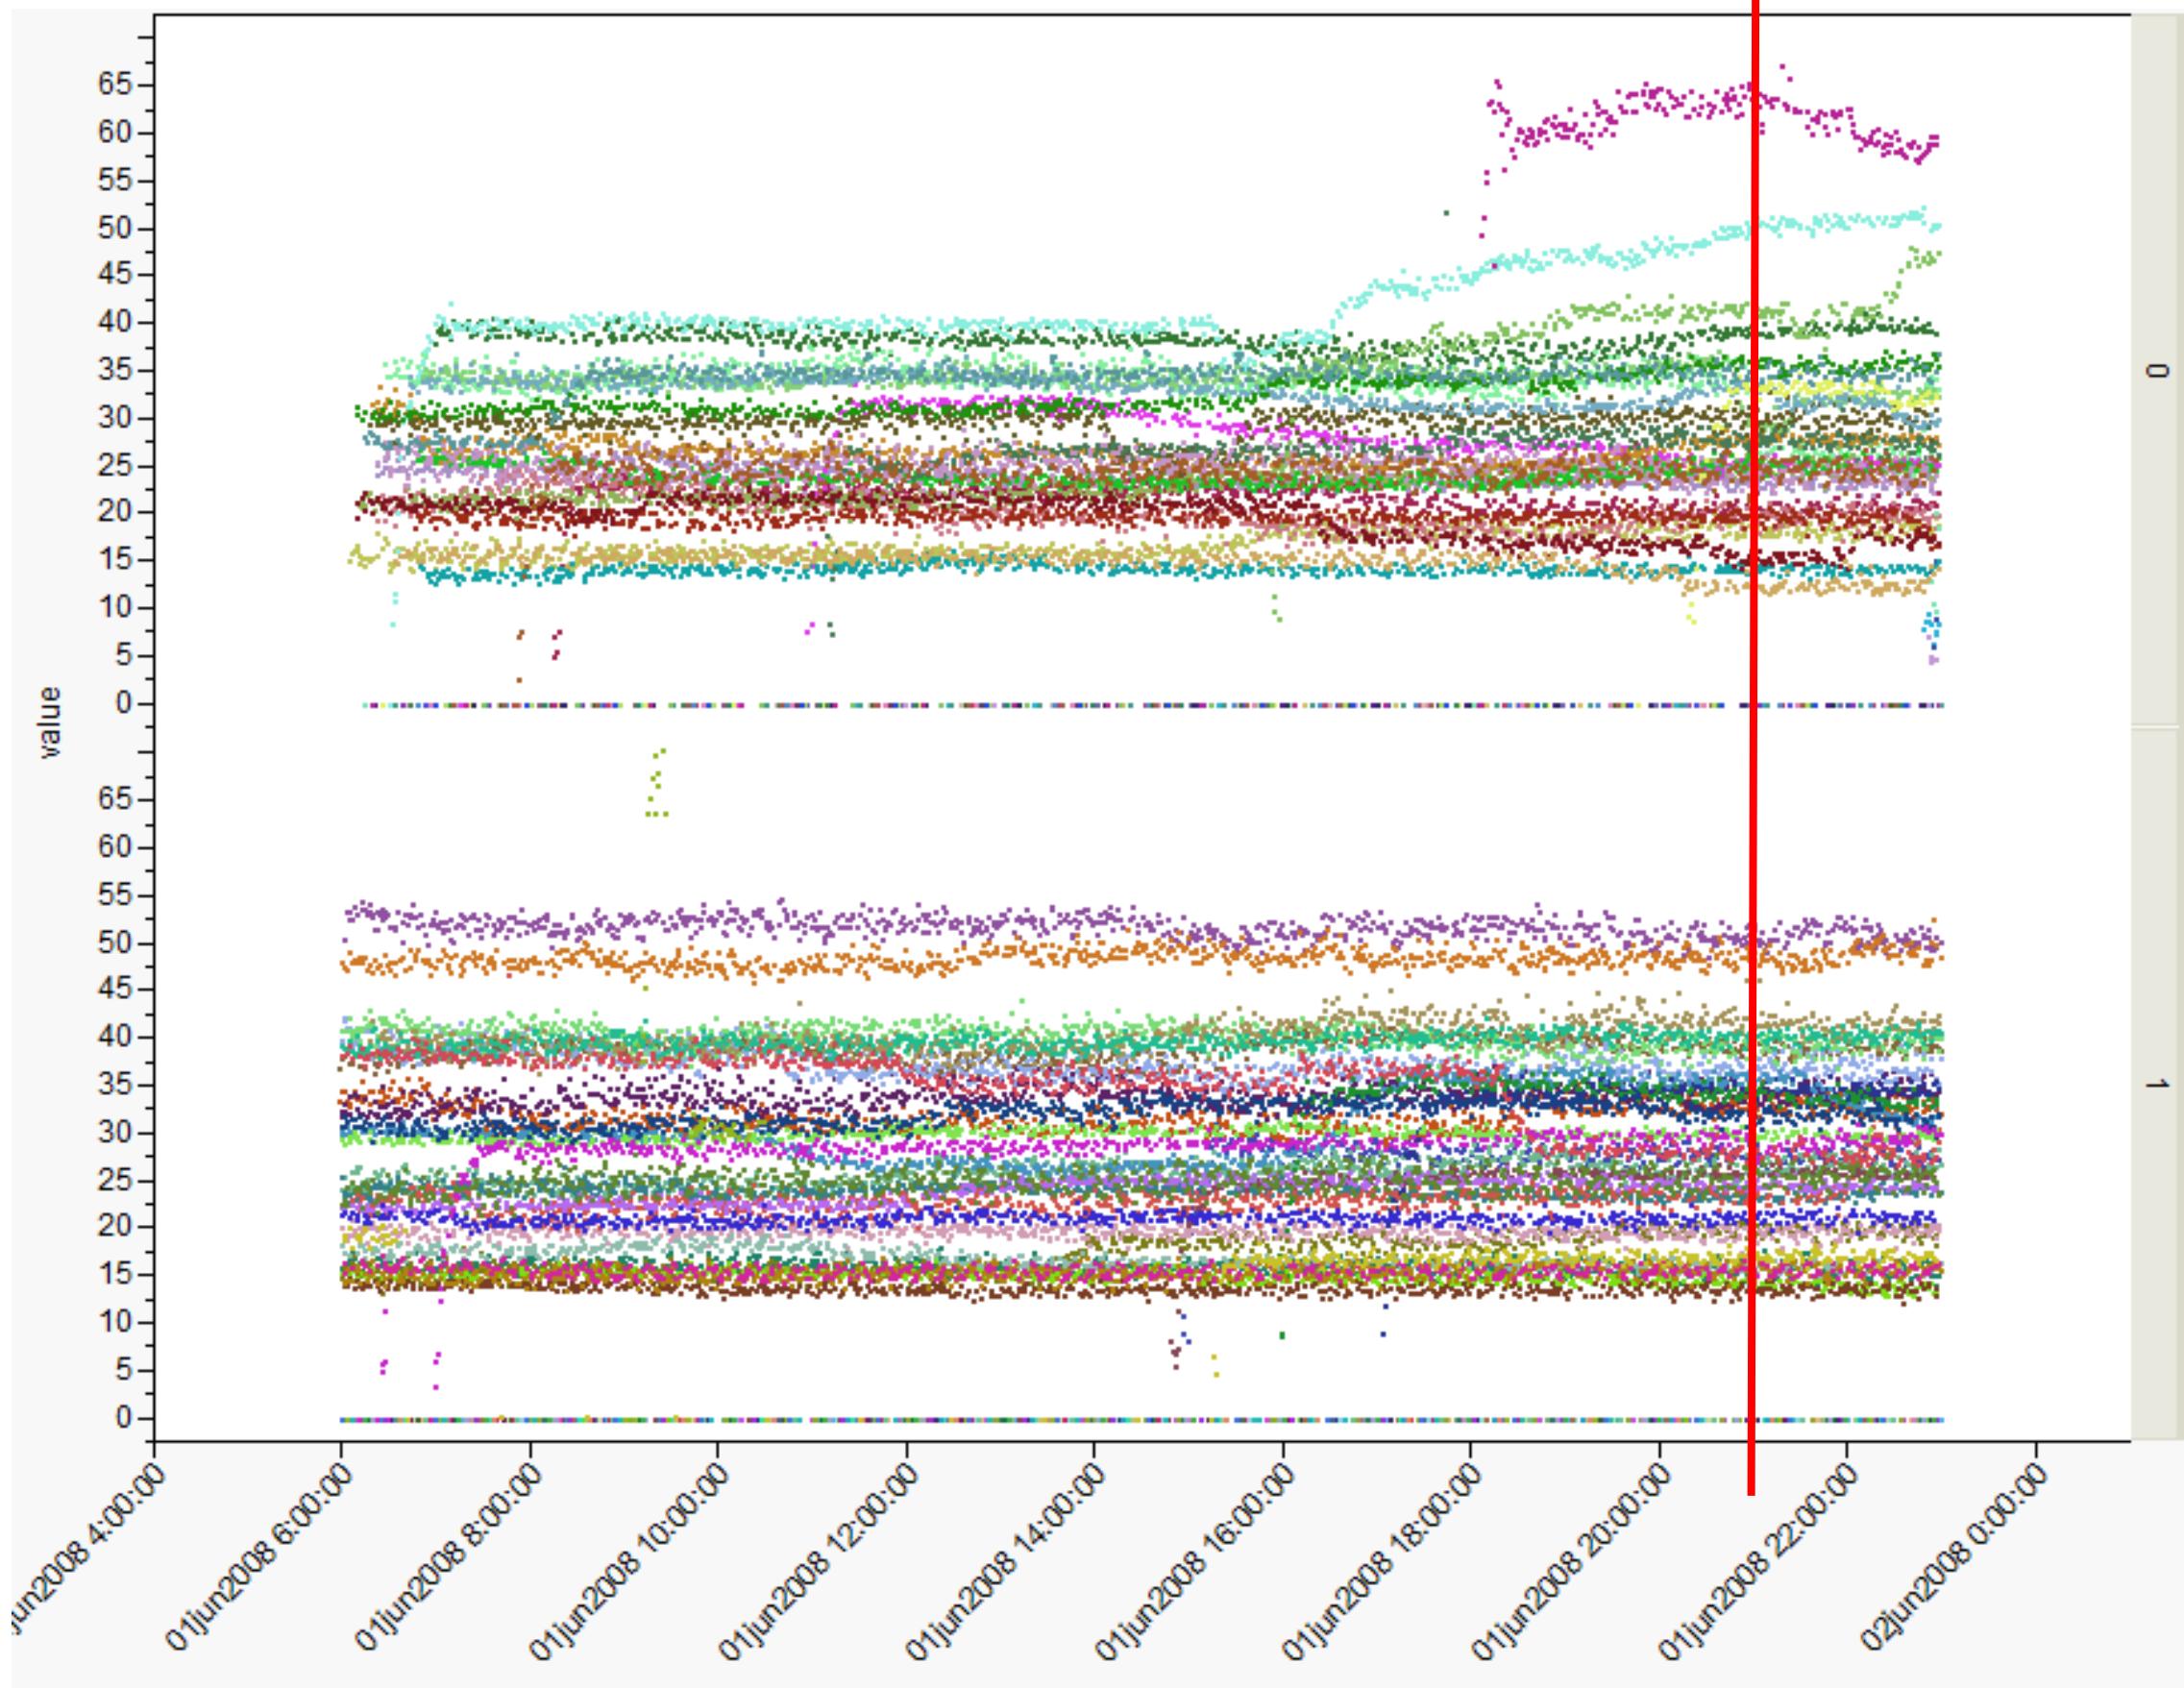
Cleaning

SOME OTHER MAGIC

- Removing outliers
- Changing reference systems
- Kalman Filter smoothing
- Combining multiple readings into 1
- ...

Time series data mining

Step 1: Prepare timeseries



Step 2: Calculate characteristics

```
for segment in segments:  
    for var in segment:  
        features_for_this_segment = [  
            var.mean(),  
            var.median(),  
            var.max (),  
            var.min (),  
            var.stdev (),  
            var.kurtosis (),  
            var.skewness(),  
            *var.percentiles([0.01, 0.02, 0.05, 0.10, ...]),  
        ]
```

<code>Series.abs()</code>	Return a Series/DataFrame with absolute numeric value of each element.
<code>Series.all([axis, bool_only, skipna, level])</code>	Return whether all elements are True, potentially over an axis.
<code>Series.any([axis, bool_only, skipna, level])</code>	Return whether any element is True, potentially over an axis.
<code>Series.autocorr([lag])</code>	Compute the lag-N autocorrelation.
<code>Series.between(left, right[, inclusive])</code>	Return boolean Series equivalent to left <= series <= right.
<code>Series.clip([lower, upper, axis, inplace])</code>	Trim values at input threshold(s).
<code>Series.clip_lower(threshold[, axis, inplace])</code>	(DEPRECATED) Trim values below a given threshold.
<code>Series.clip_upper(threshold[, axis, inplace])</code>	(DEPRECATED) Trim values above a given threshold.
<code>Series.corr(other[, method, min_periods])</code>	Compute correlation with other Series, excluding missing values.
<code>Series.count([level])</code>	Return number of non-NA/null observations in the Series.
<code>Series.cov(other[, min_periods])</code>	Compute covariance with Series, excluding missing values.
<code>Series.cummax([axis, skipna])</code>	Return cumulative maximum over a DataFrame or Series axis.
<code>Series.cummin([axis, skipna])</code>	Return cumulative minimum over a DataFrame or Series axis.
<code>Series.cumprod([axis, skipna])</code>	Return cumulative product over a DataFrame or Series axis.
<code>Series.cumsum([axis, skipna])</code>	Return cumulative sum over a DataFrame or Series axis.
<code>Series.describe([percentiles, include, exclude])</code>	Generate descriptive statistics that summarize the central tendency, dispersion and shape of a dataset's distribution, excluding NaN values.
<code>Series.diff([periods])</code>	First discrete difference of element.
<code>Series.factorize([sort, na_sentinel])</code>	Encode the object as an enumerated type or categorical variable.
<code>Series.kurt([axis, skipna, level, numeric_only])</code>	Return unbiased kurtosis over requested axis using Fisher's definition of kurtosis (kurtosis of normal == 0.0).
<code>Series.mad([axis, skipna, level])</code>	Return the mean absolute deviation of the values for the requested axis.
<code>Series.max([axis, skipna, level, numeric_only])</code>	Return the maximum of the values for the requested axis.
<code>Series.mean([axis, skipna, level, numeric_only])</code>	Return the mean of the values for the requested axis.
<code>Series.median([axis, skipna, level, ...])</code>	Return the median of the values for the requested axis.
<code>Series.min([axis, skipna, level, numeric_only])</code>	Return the minimum of the values for the requested axis.
<code>Series.mode([dropna])</code>	Return the mode(s) of the dataset.
<code>Series.nlargest([n, keep])</code>	Return the largest n elements.
<code>Series.nsmallest([n, keep])</code>	Return the smallest n elements.
<code>Series.pct_change([periods, fill_method, ...])</code>	Percentage change between the current and a prior element.
<code>Series.prod([axis, skipna, level, ...])</code>	Return the product of the values for the requested axis.
<code>Series.quantile([q, interpolation])</code>	Return value at the given quantile.
<code>Series.rank([axis, method, numeric_only, ...])</code>	Compute numerical data ranks (1 through n) along axis.
<code>Series.sem([axis, skipna, level, ddof, ...])</code>	Return unbiased standard error of the mean over requested axis.
<code>Series.skew([axis, skipna, level, numeric_only])</code>	Return unbiased skew over requested axis Normalized by N-1.
<code>Series.std([axis, skipna, level, ddof, ...])</code>	Return sample standard deviation over requested axis.
<code>Series.sum([axis, skipna, level, ...])</code>	Return the sum of the values for the requested axis.
<code>Series.var([axis, skipna, level, ddof, ...])</code>	Return unbiased variance over requested axis.
<code>Series.kurtosis([axis, skipna, level, ...])</code>	Return unbiased kurtosis over requested axis using Fisher's definition of kurtosis (kurtosis of normal == 0.0).
<code>Series.unique()</code>	Return unique values of Series object.
<code>Series.nunique([dropna])</code>	Return number of unique elements in the object.
<code>Series.is_unique</code>	Return boolean if values in the object are unique.
<code>Series.is_monotonic</code>	Return boolean if values in the object are monotonic_increasing.
<code>Series.is_monotonic_increasing</code>	Return boolean if values in the object are monotonic_increasing.
<code>Series.is_monotonic_decreasing</code>	Return boolean if values in the object are monotonic_decreasing.
<code>Series.value_counts([normalize, sort, ...])</code>	Return a Series containing counts of unique values.
<code>Series.compound([axis, skipna, level])</code>	Return the compound percentage of the values for the requested axis.



Time series data mining

Step 1: Prepare timeseries

Step 2: Calculate characteristics

Step 3: Train model

Characteristics become features (+- 12000)

160 observations

Random forests, XGBoost, Linear Discriminant

Analysis, ... or another

high dimensional classification technique

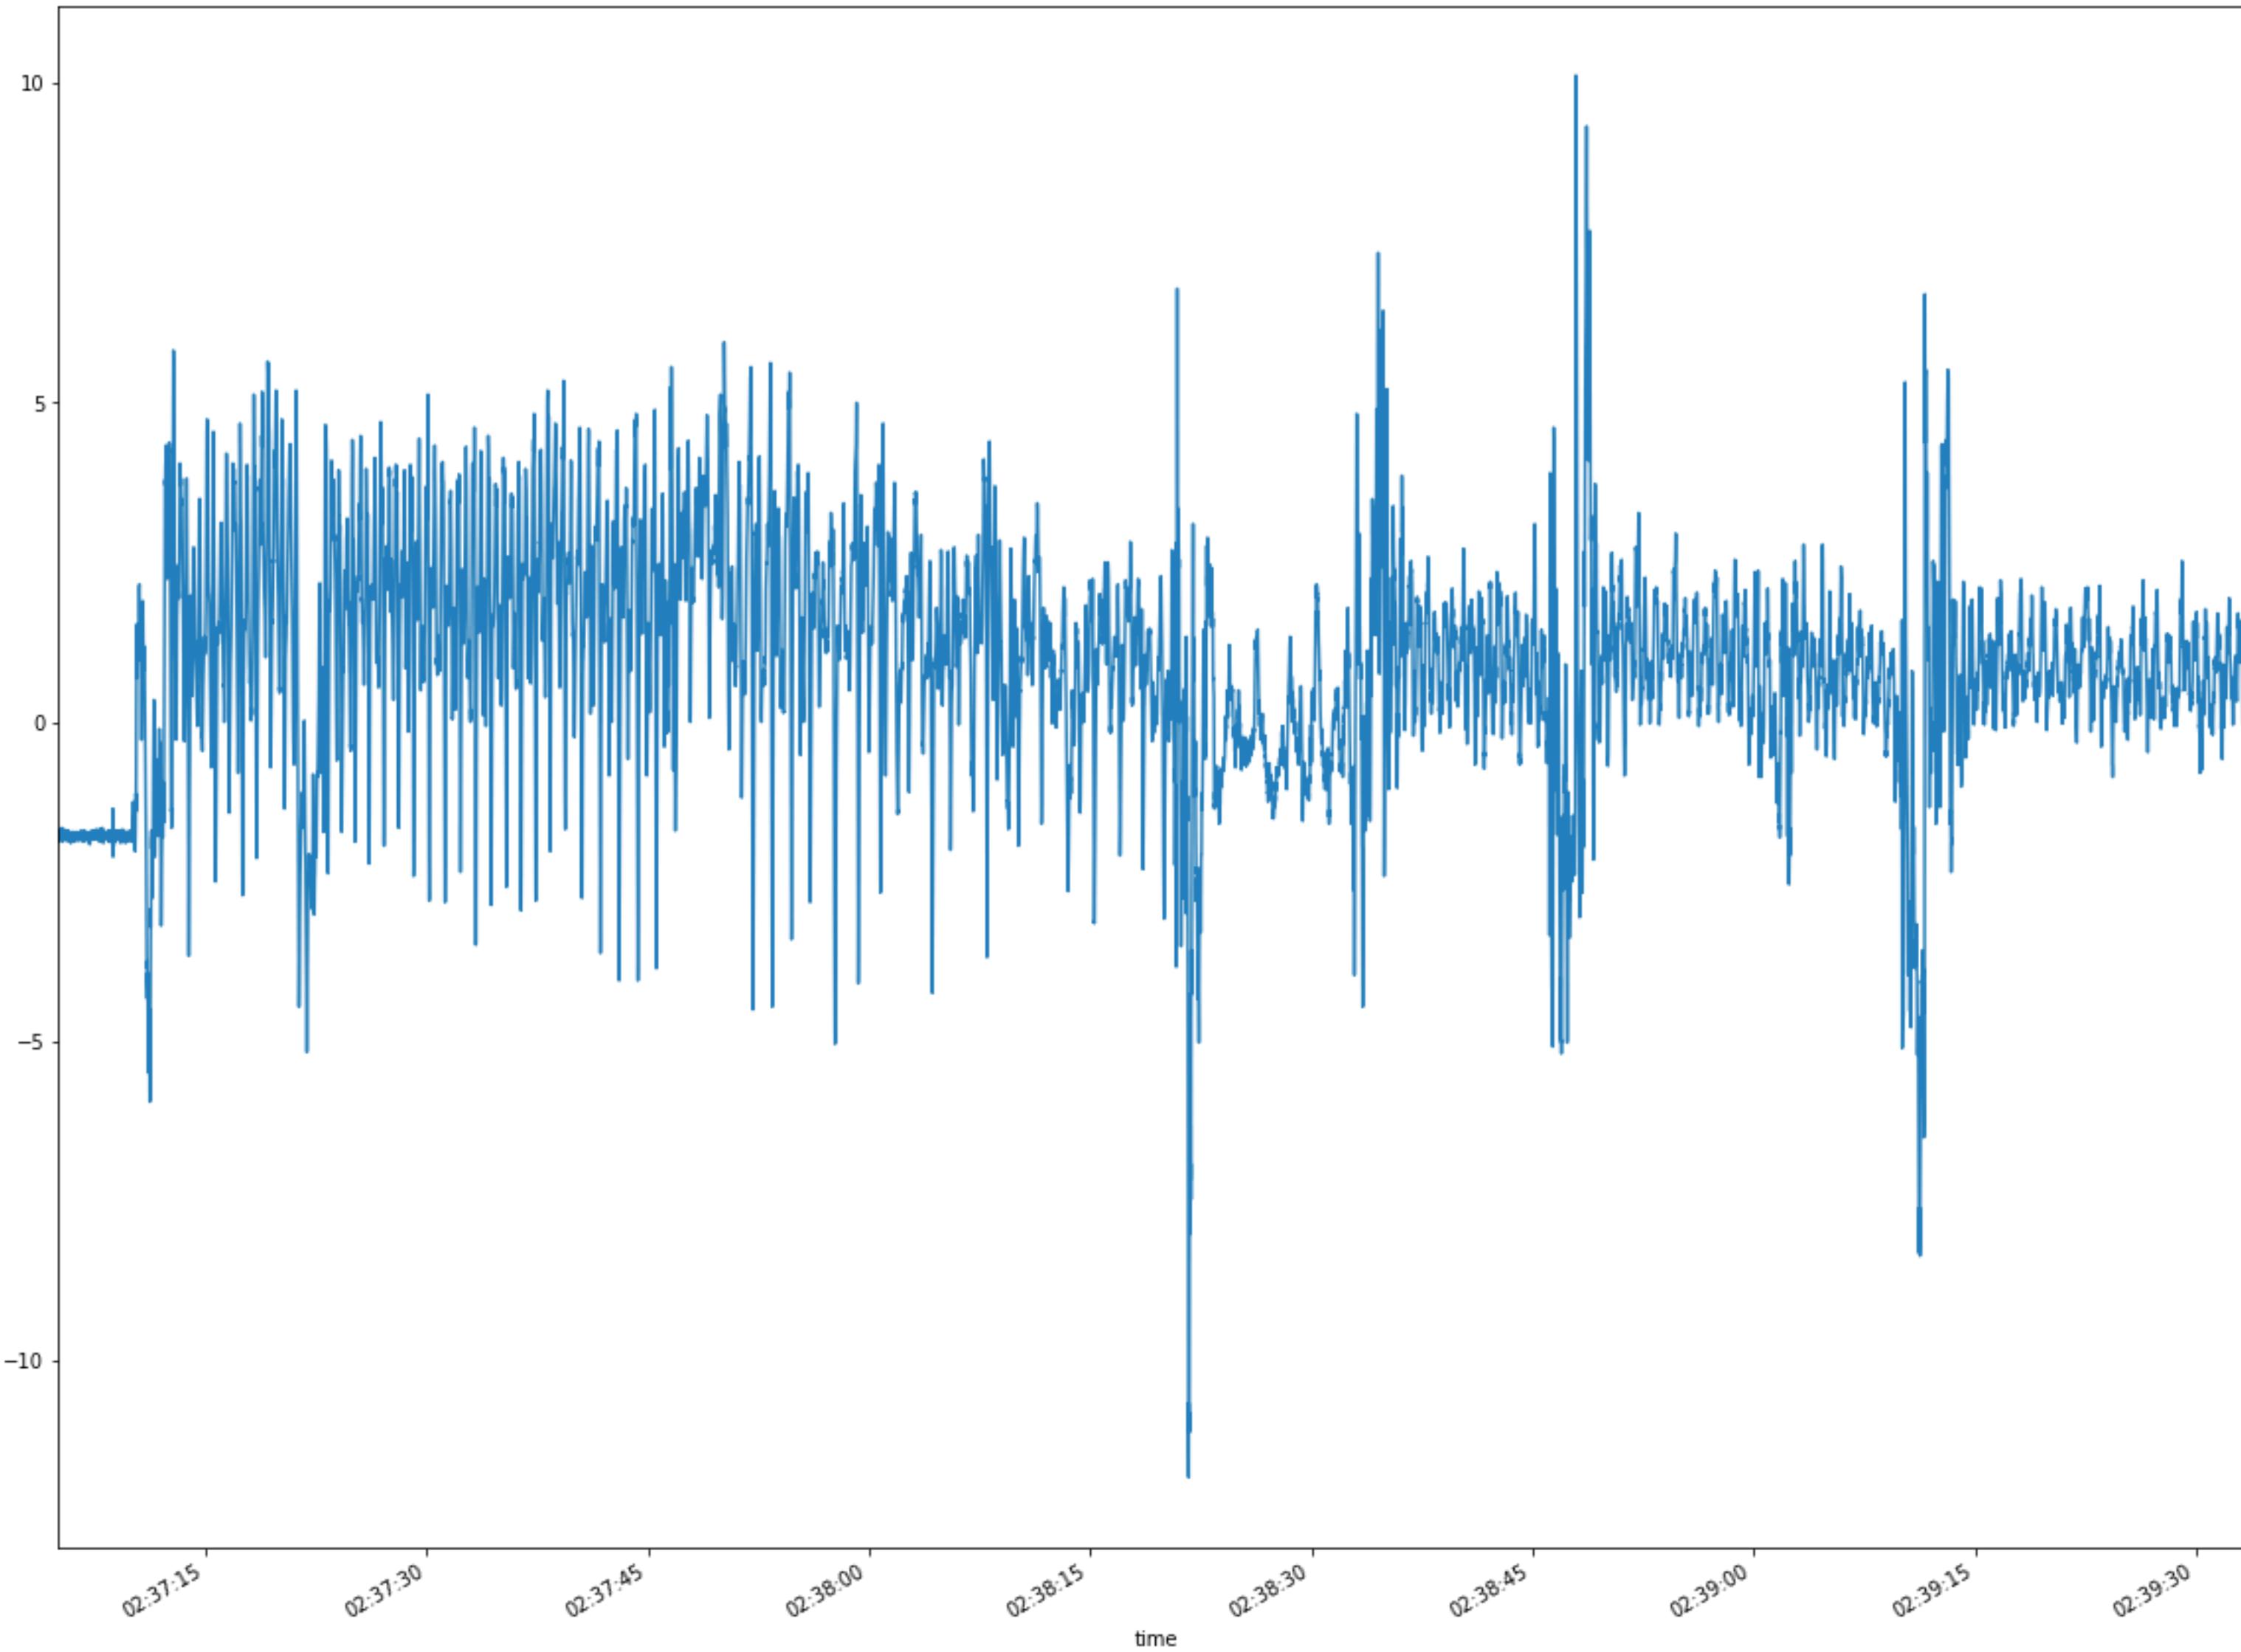
MORE FEATURES

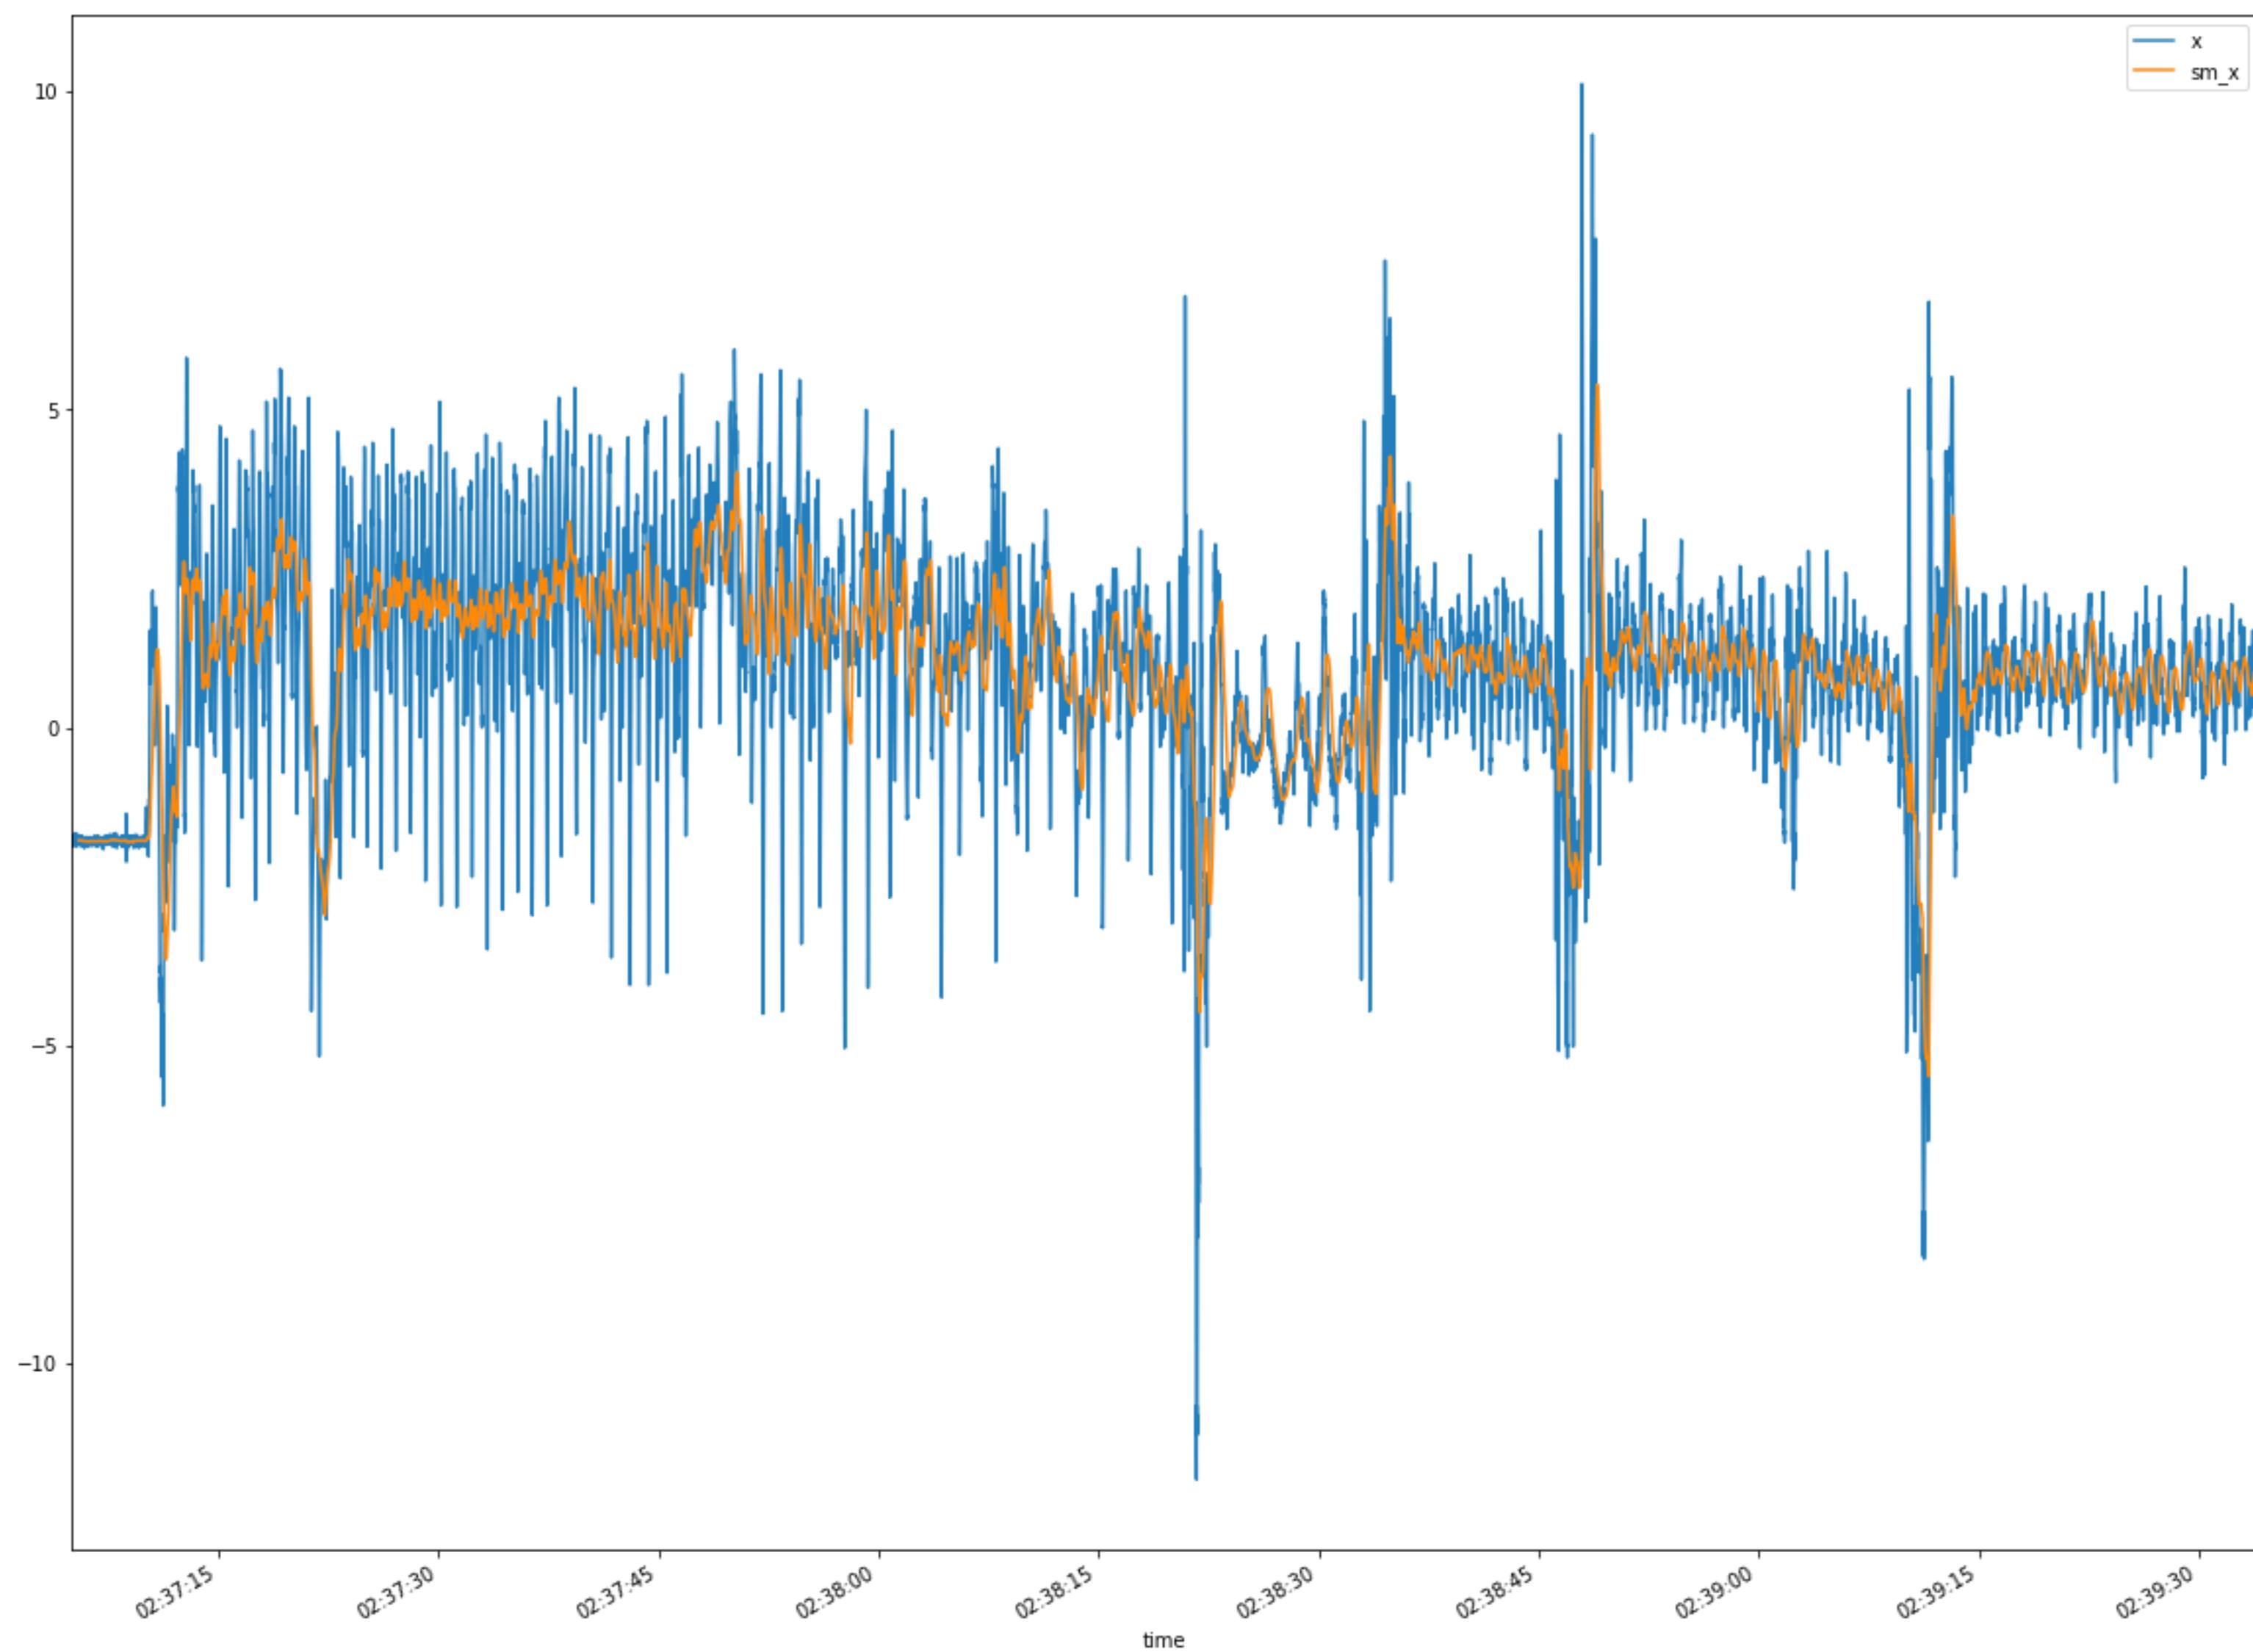
→ *Specific applications require specific models.*

CALCULATE ALL THE FEATURES

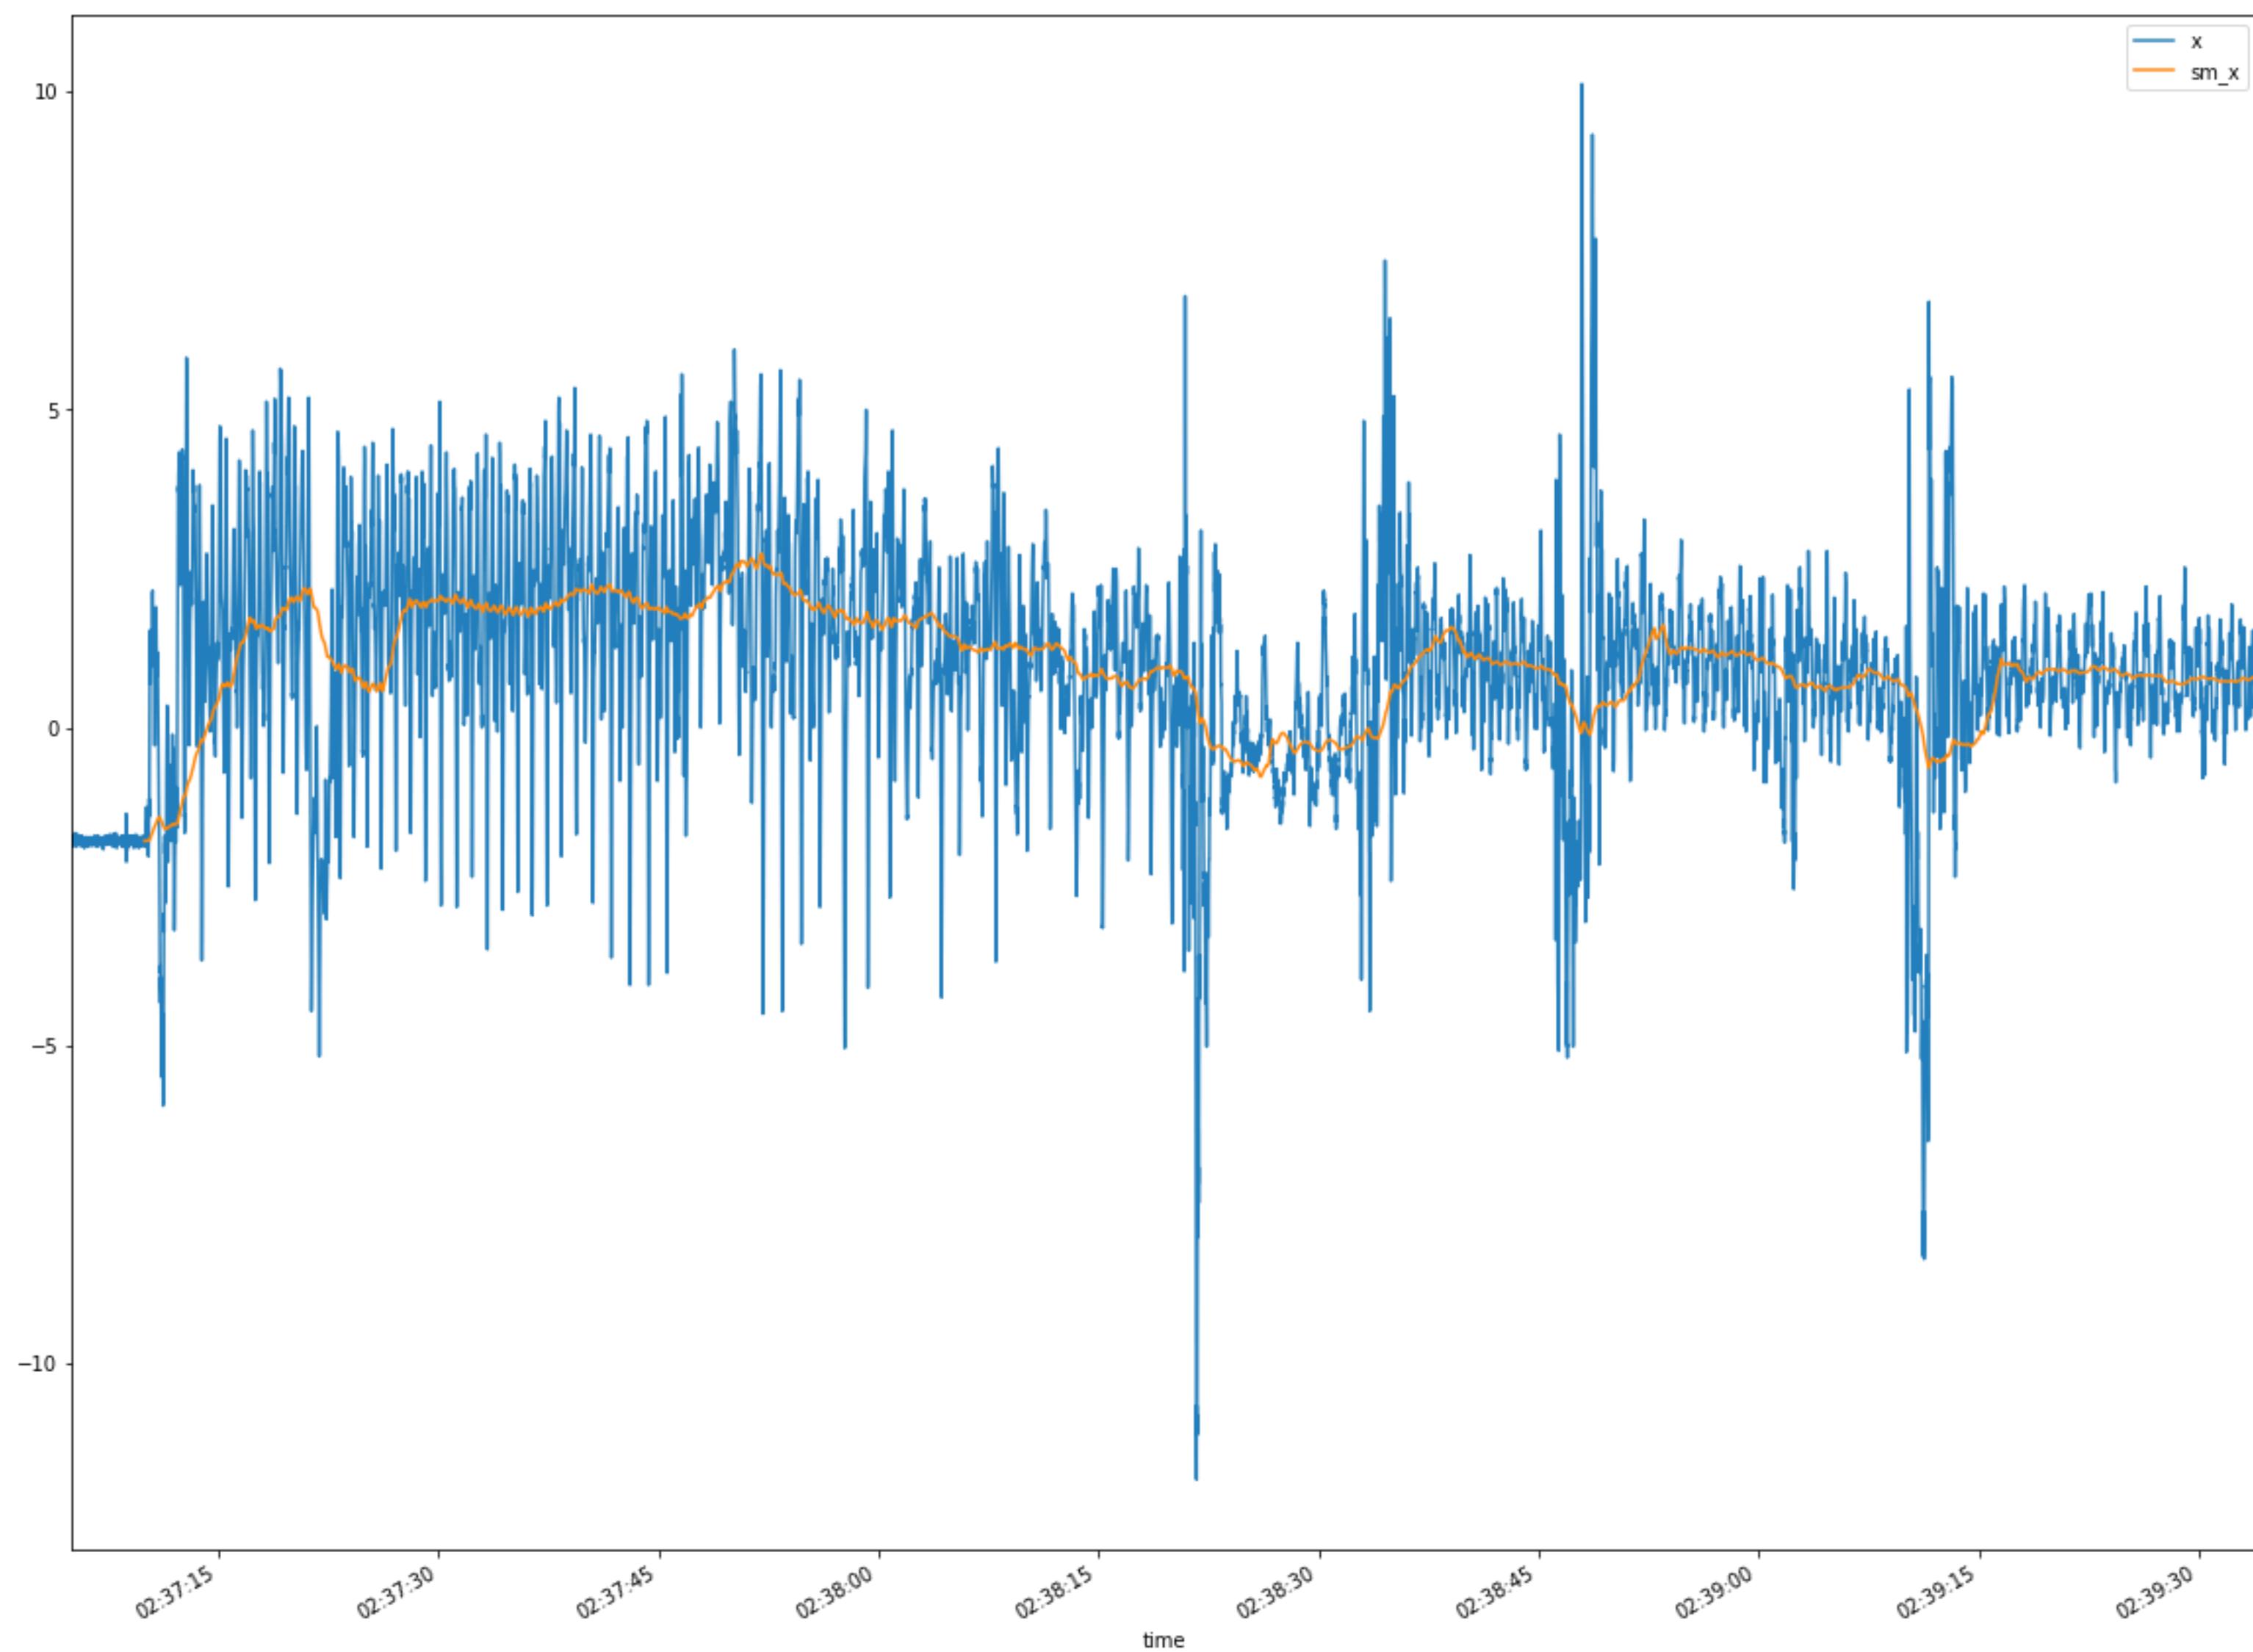


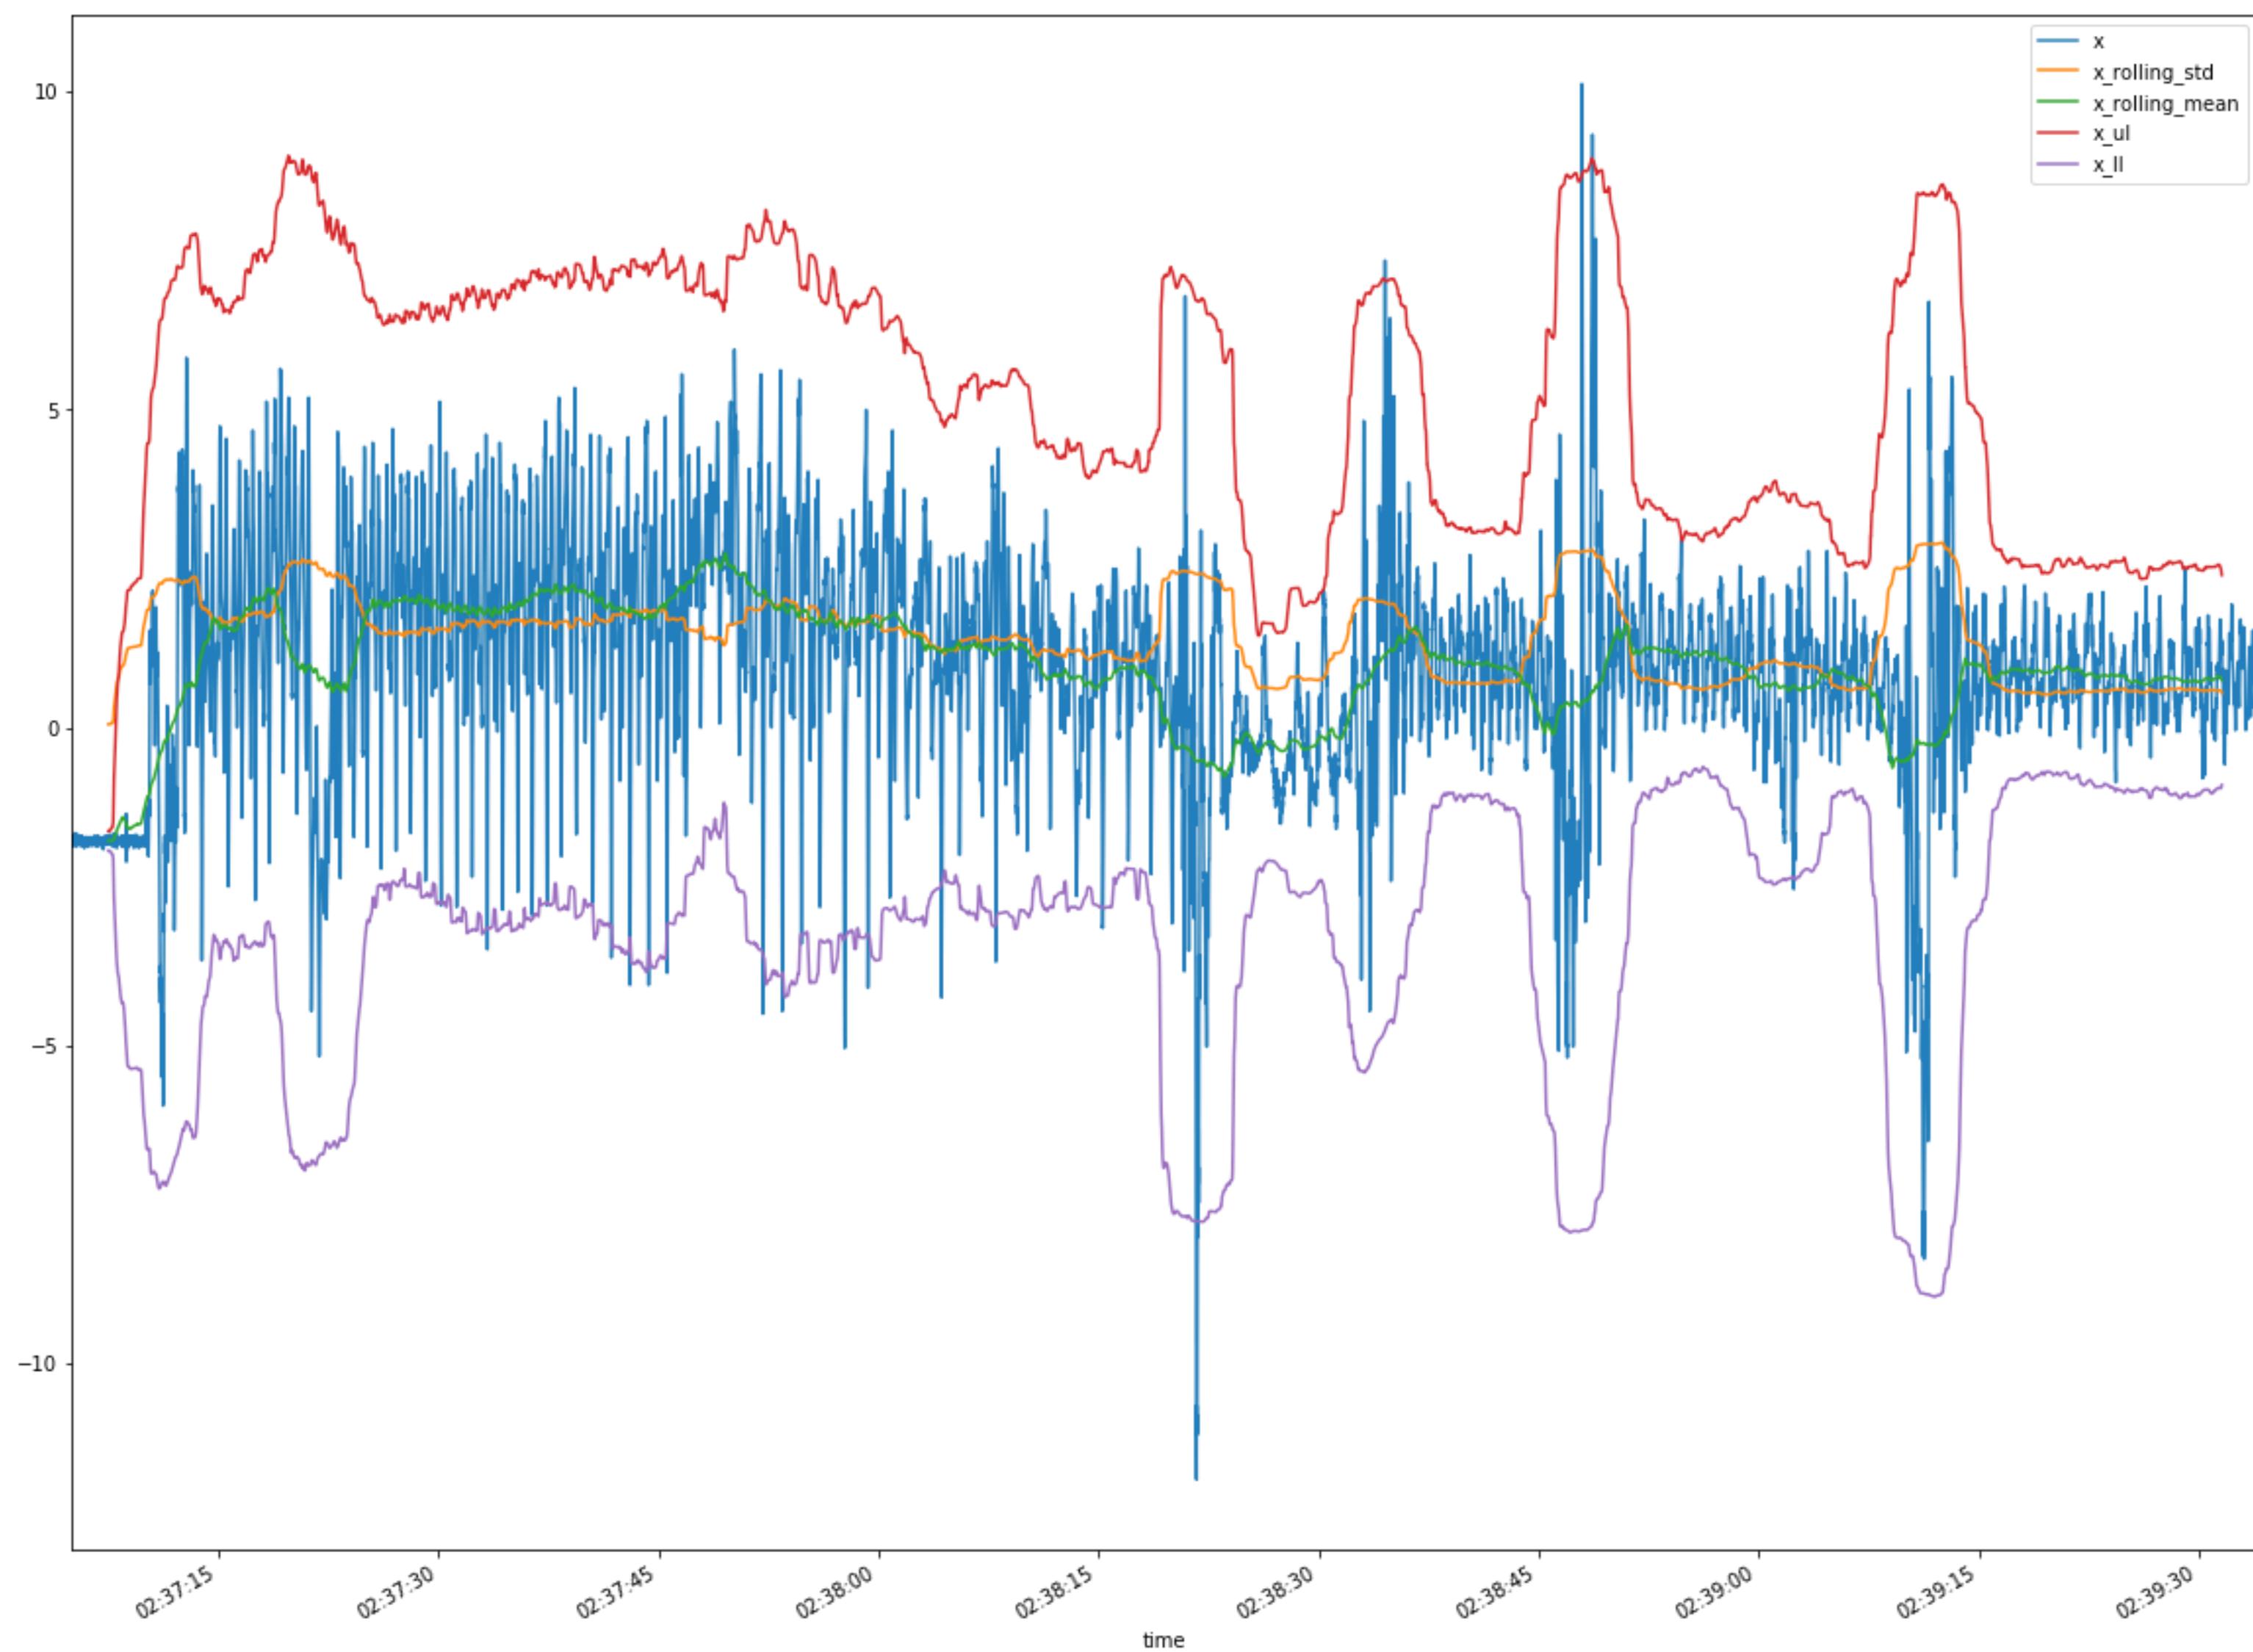
Peak detection algorithms





Rolling mean with window=50





<http://localhost:8888/notebooks/plot%20sensor%20data%20example.ipynb>

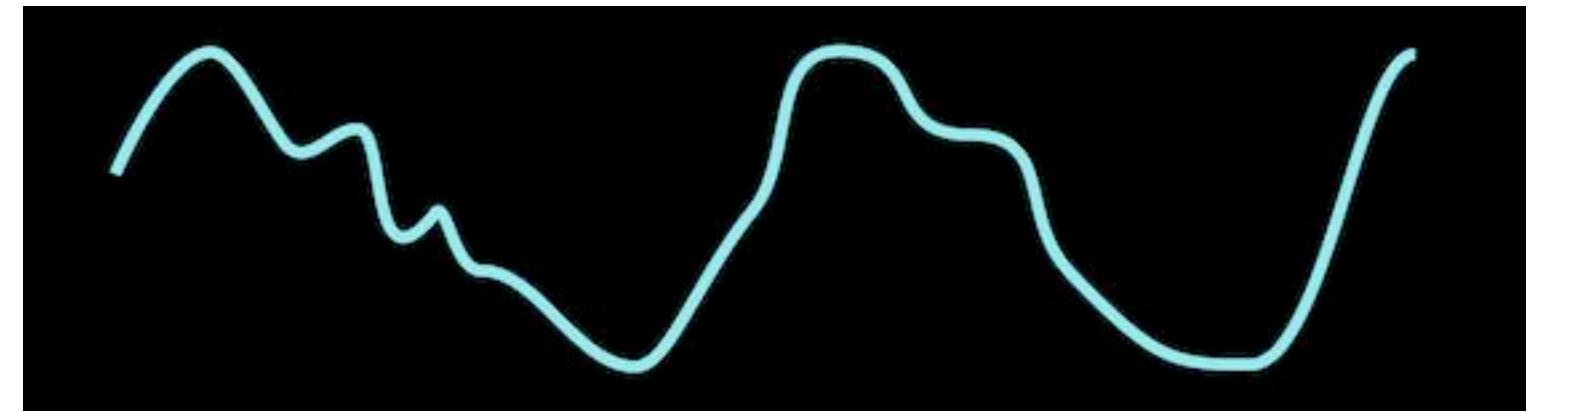
D
F

JOSEPH FOURIER



Joseph Fourier was a French mathematician and physicist who lived from 1768 and 1830 invented Fourier analysis, which is an important component of analyzing sensor data

FUNCTION



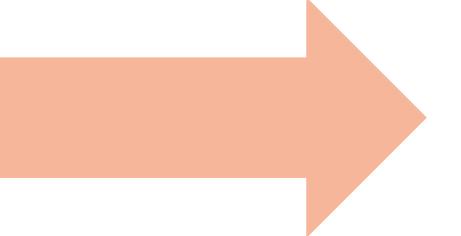
BEFORE

DATA

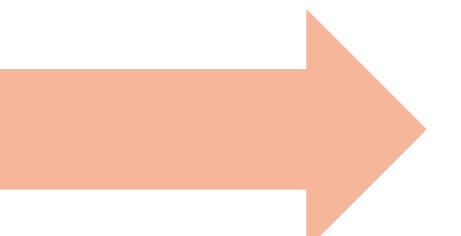
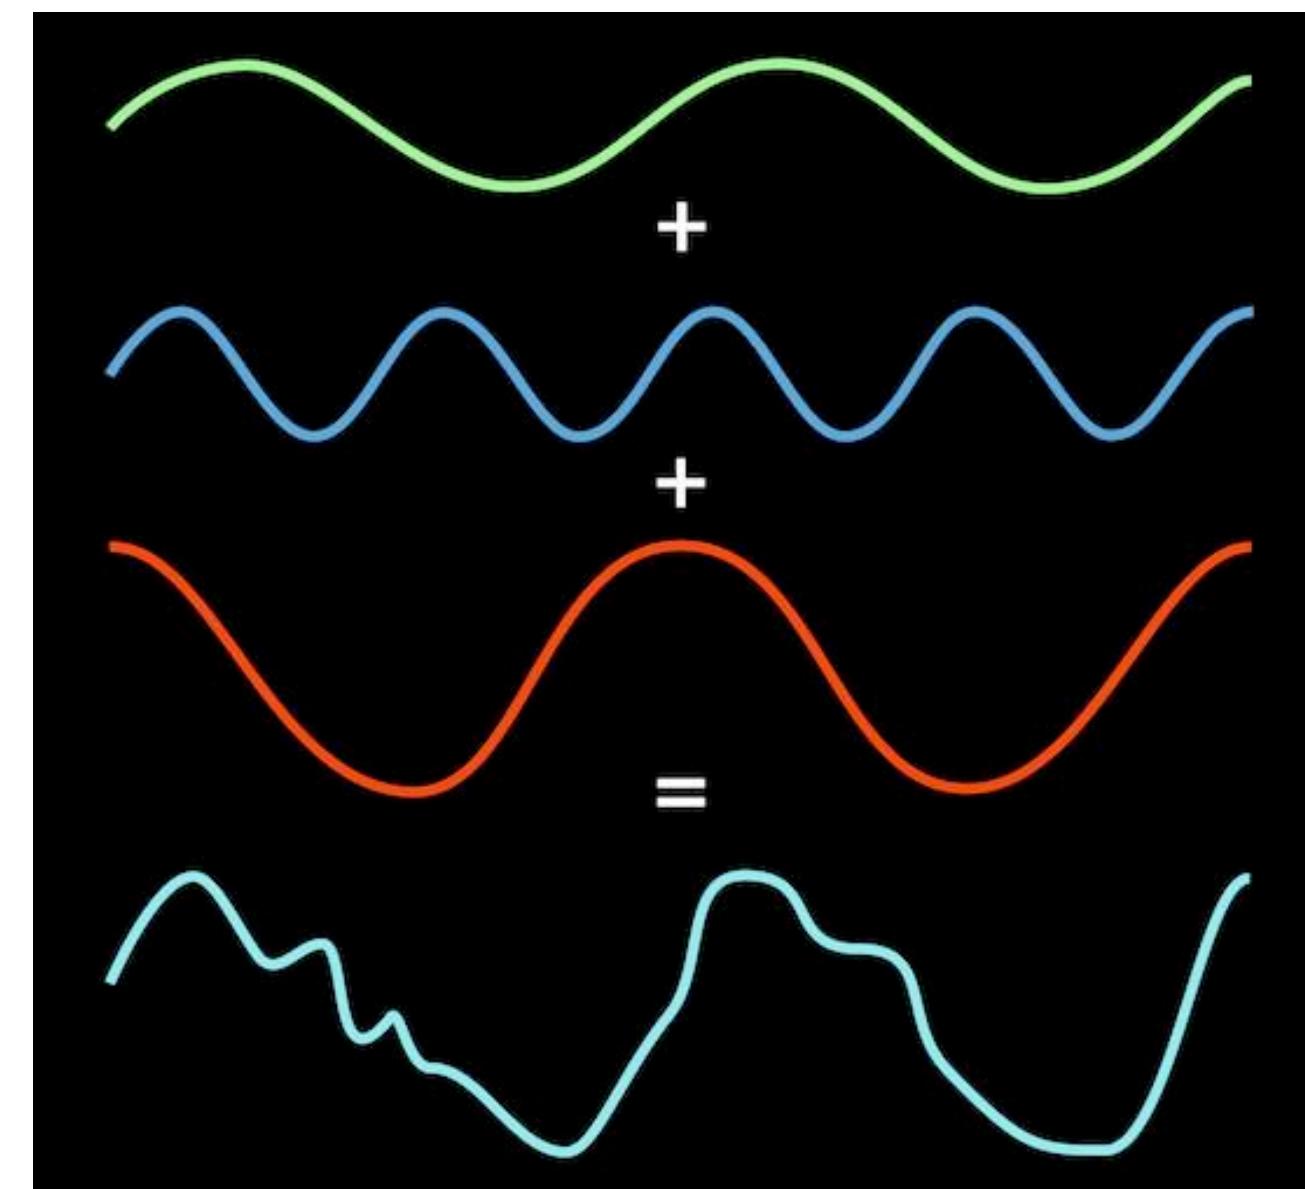
1000 times and values

DFT

$$X_k = \sum_{n=0}^{N-1} x_n \cdot e^{-\frac{i2\pi}{N} kn}$$



AFTER

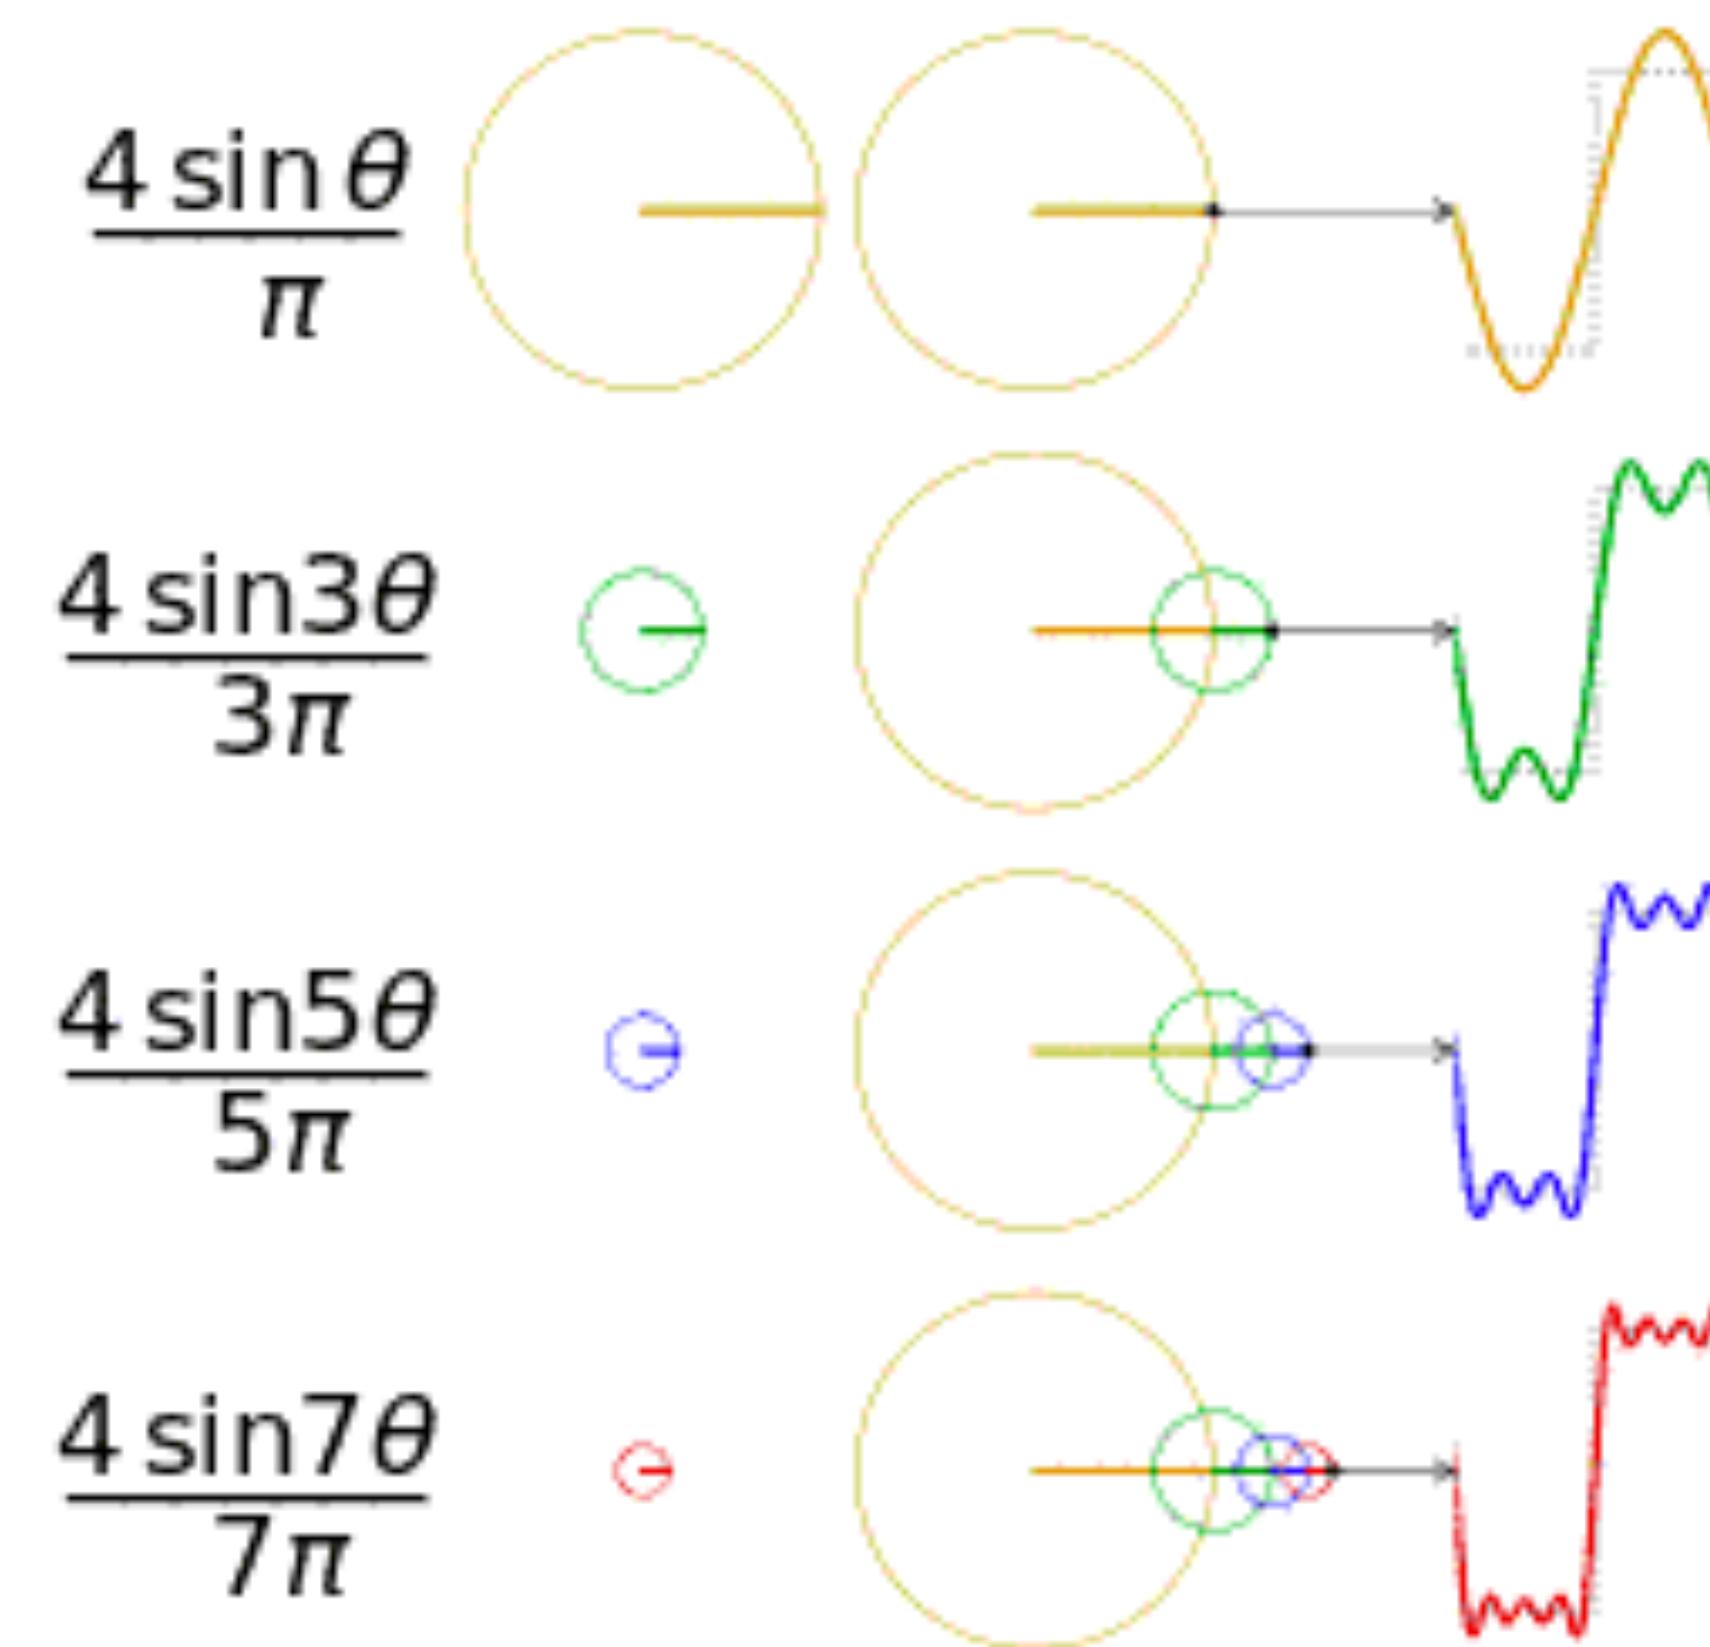


1000 frequencies and amplitudes

Recap or Intro to FOURIER analysis



SINE WAVE



Decomposition

Hi, Dr. Elizabeth?
Yeah, uh... I accidentally took
the Fourier transform of my cat...

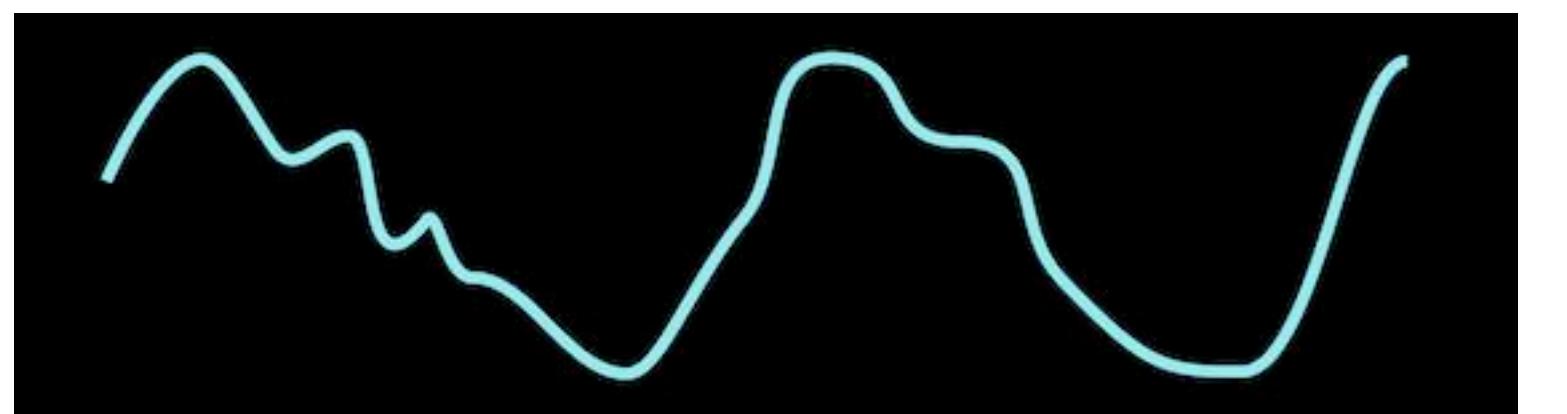


Meow!

A wavy line representing the sound of a cat's meow, with the word "Meow!" written above it.



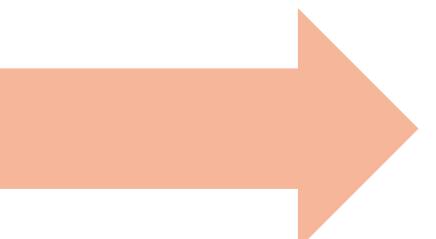
FUNCTION



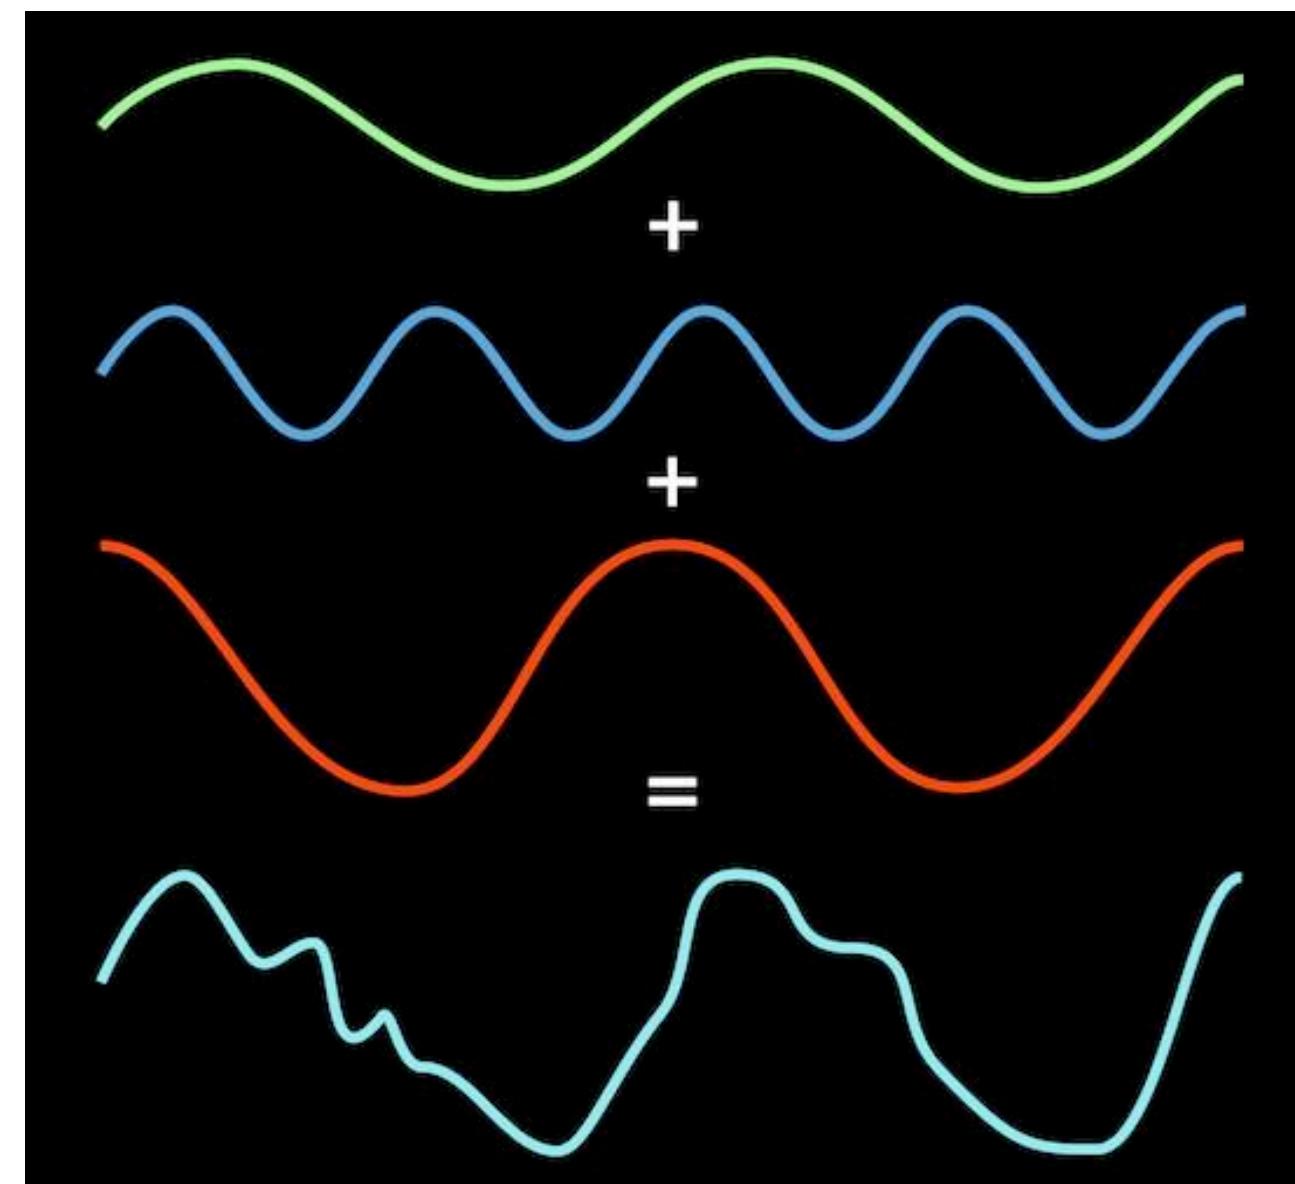
BEFORE

DFT

$$X_k = \sum_{n=0}^{N-1} x_n \cdot e^{-\frac{i2\pi}{N} kn}$$

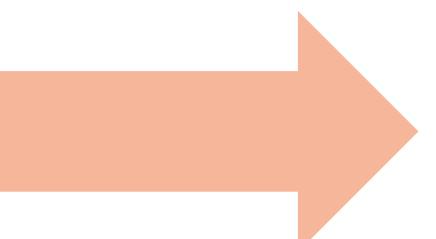


AFTER



DATA

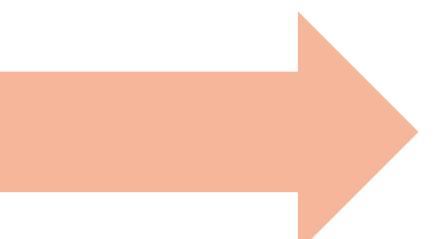
1000 times and values



1000 frequencies and amplitudes

MODEL

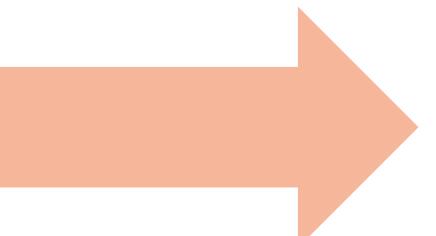
Use 1000 values for modelling



Use top 3 freqs and amplitudes
for modelling

MP3 ALGO

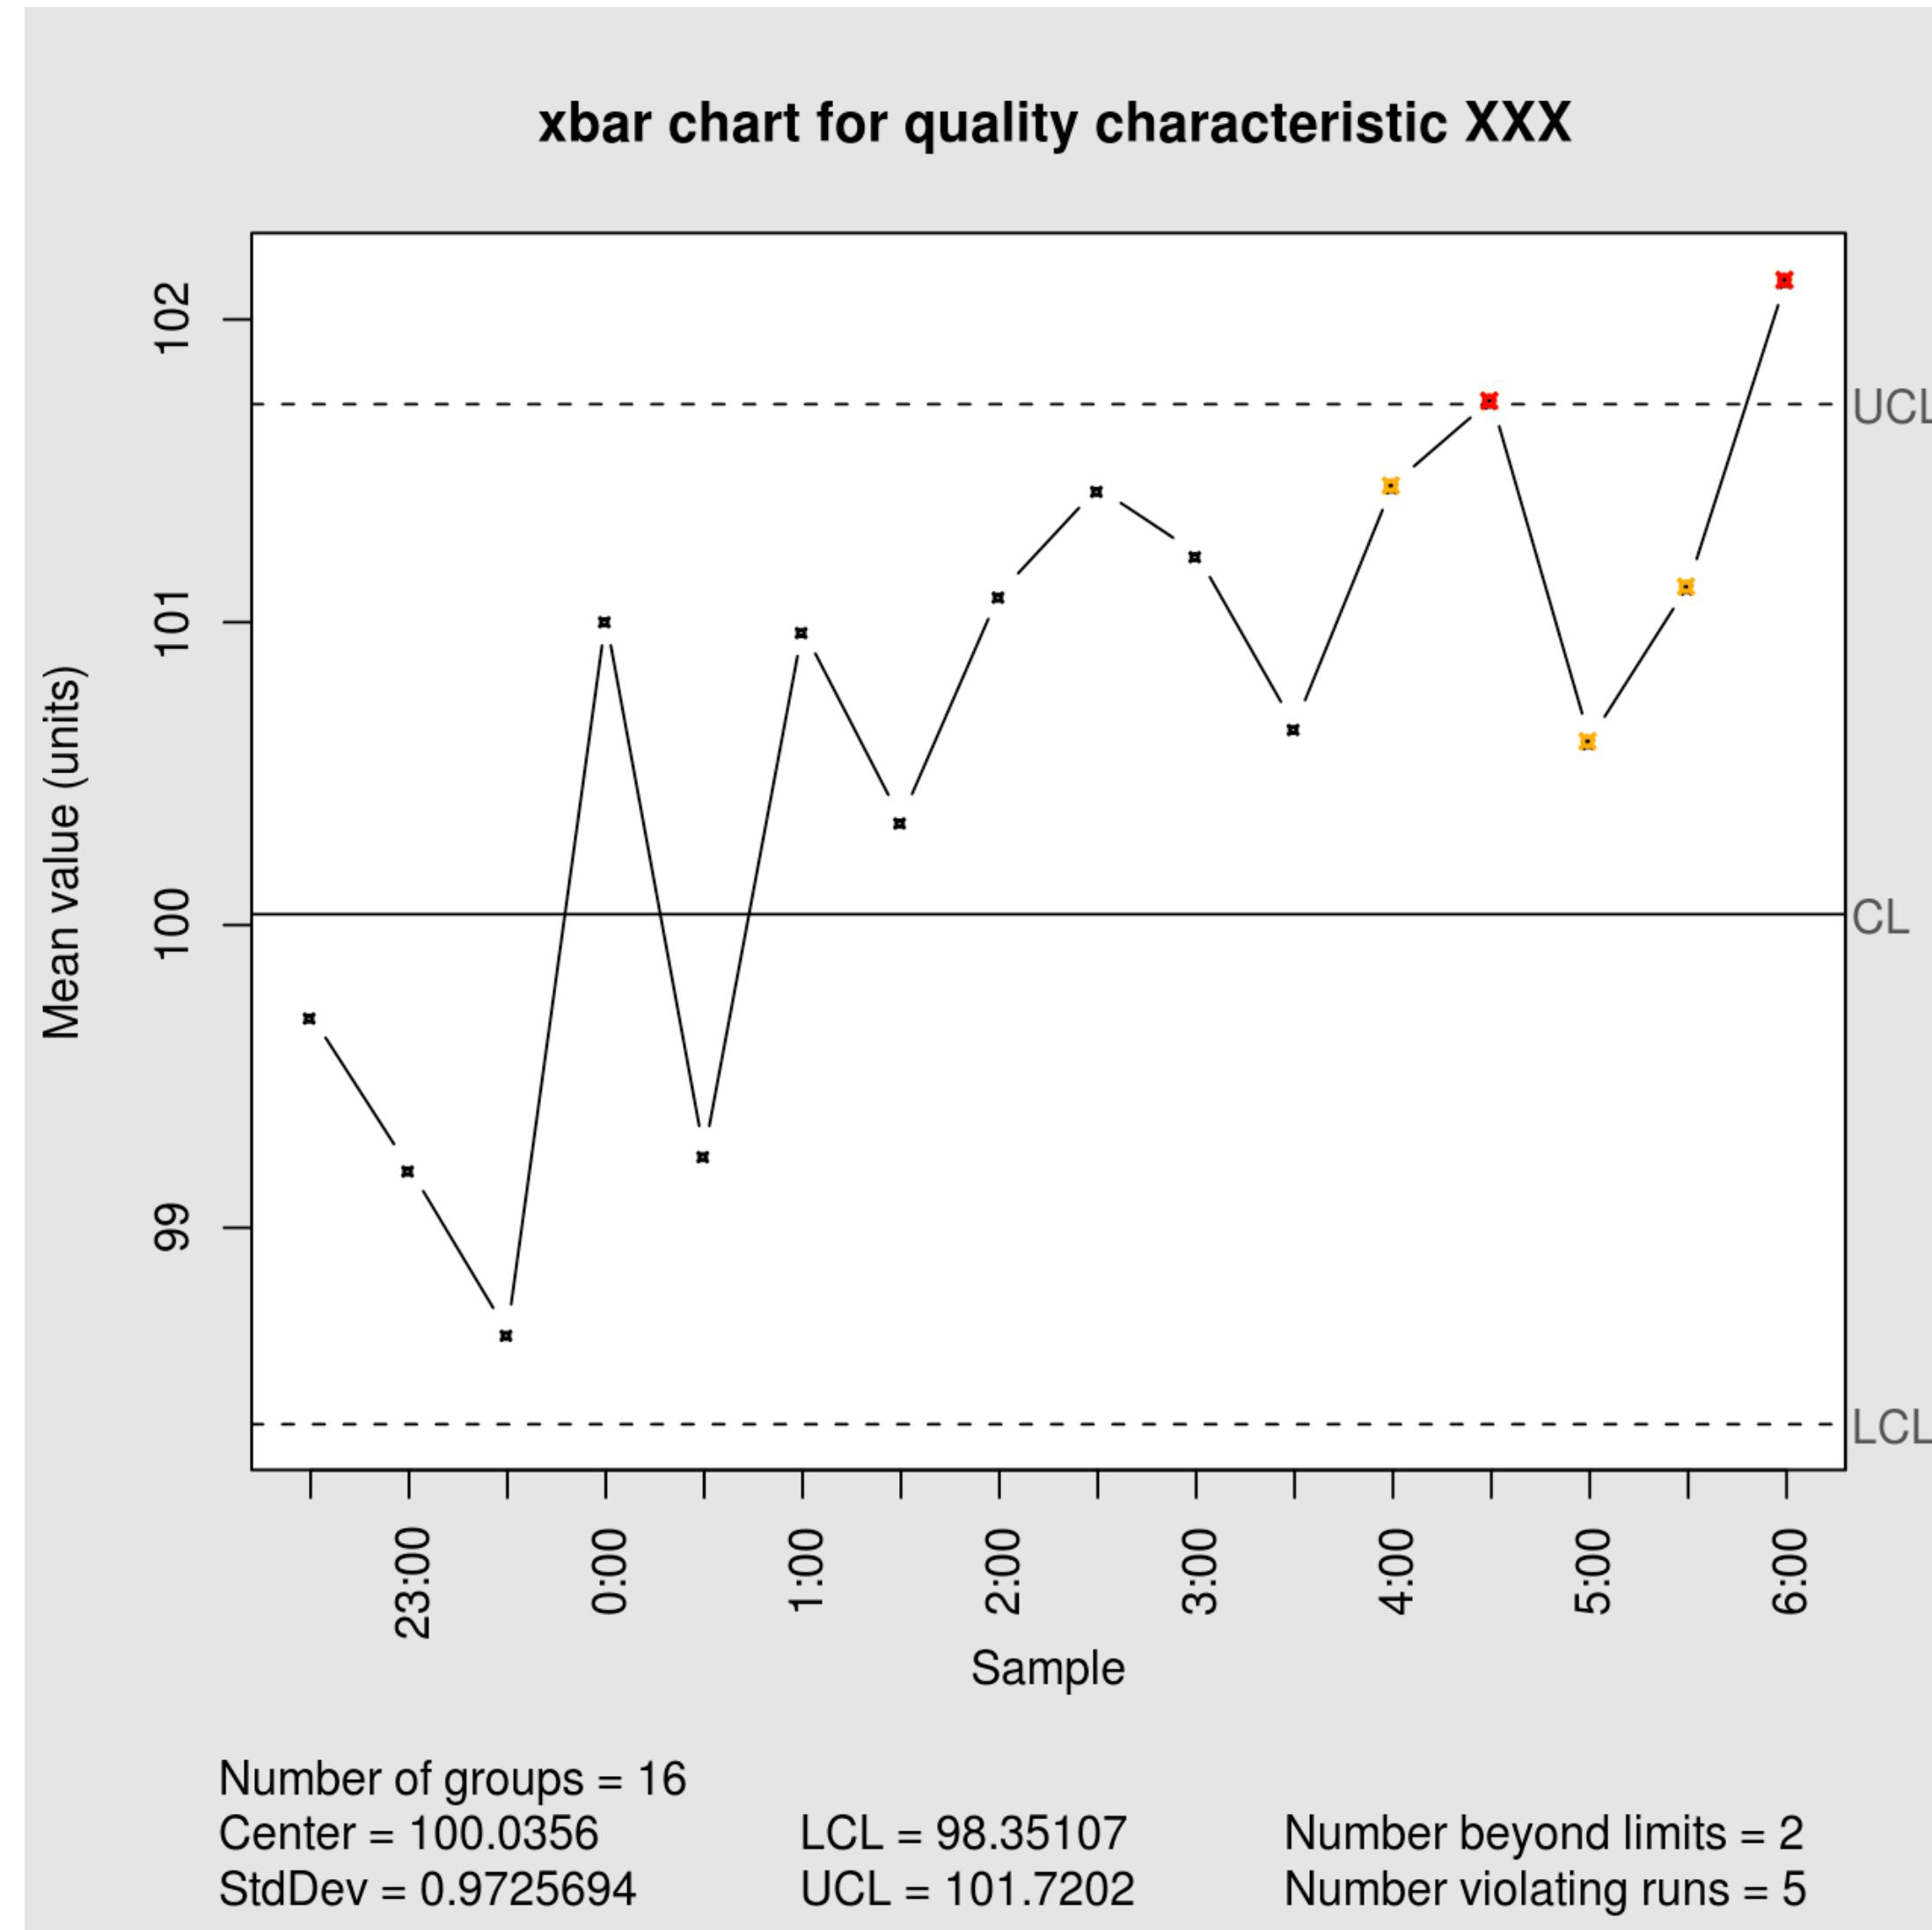
Store 1000 values in WAV file



Store 50 values in MP3 file

<http://localhost:8888/notebooks/plot%20fourier.ipynb#10.1.-Analyzing-the-frequency-components-of-a-signal-with-a-Fast-Fourier-Transform>

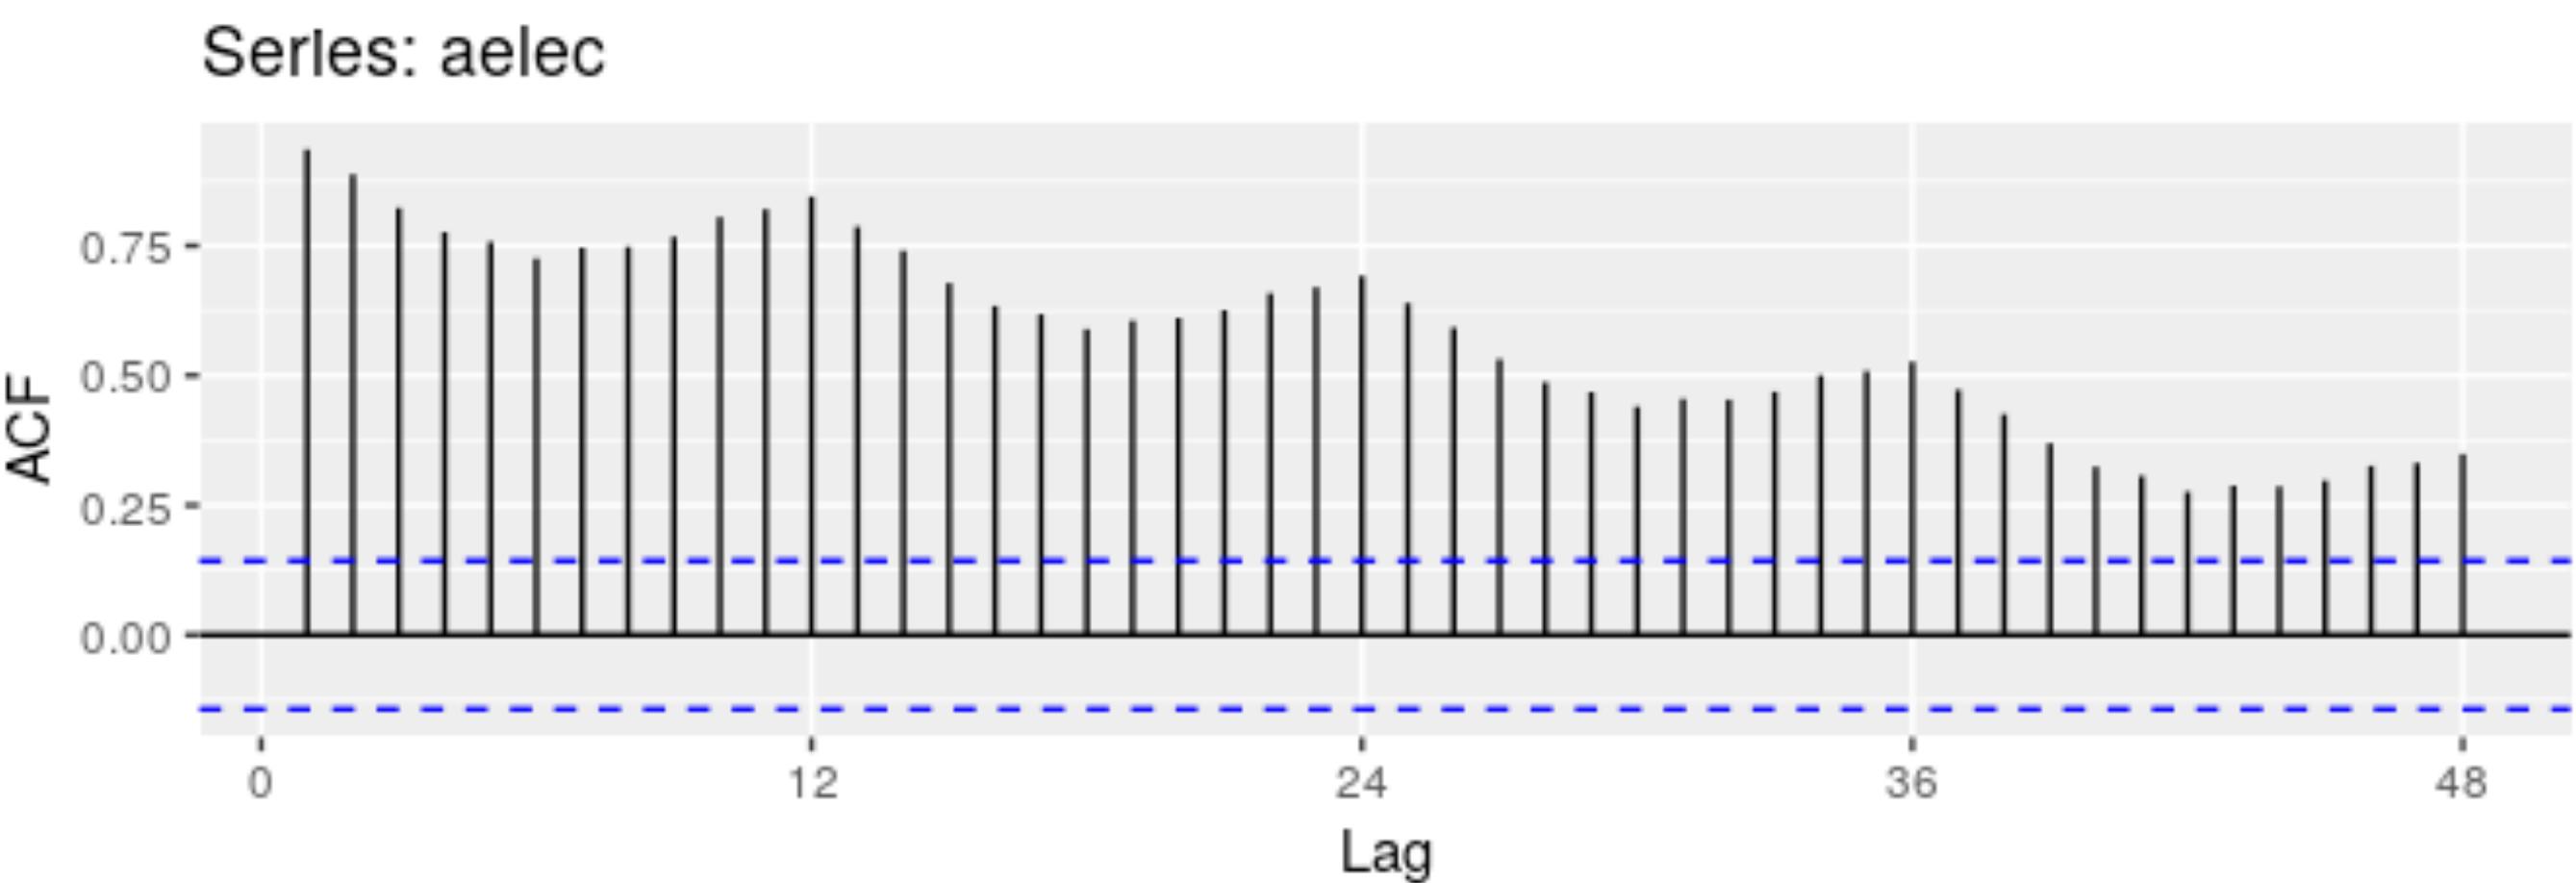
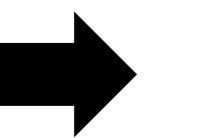
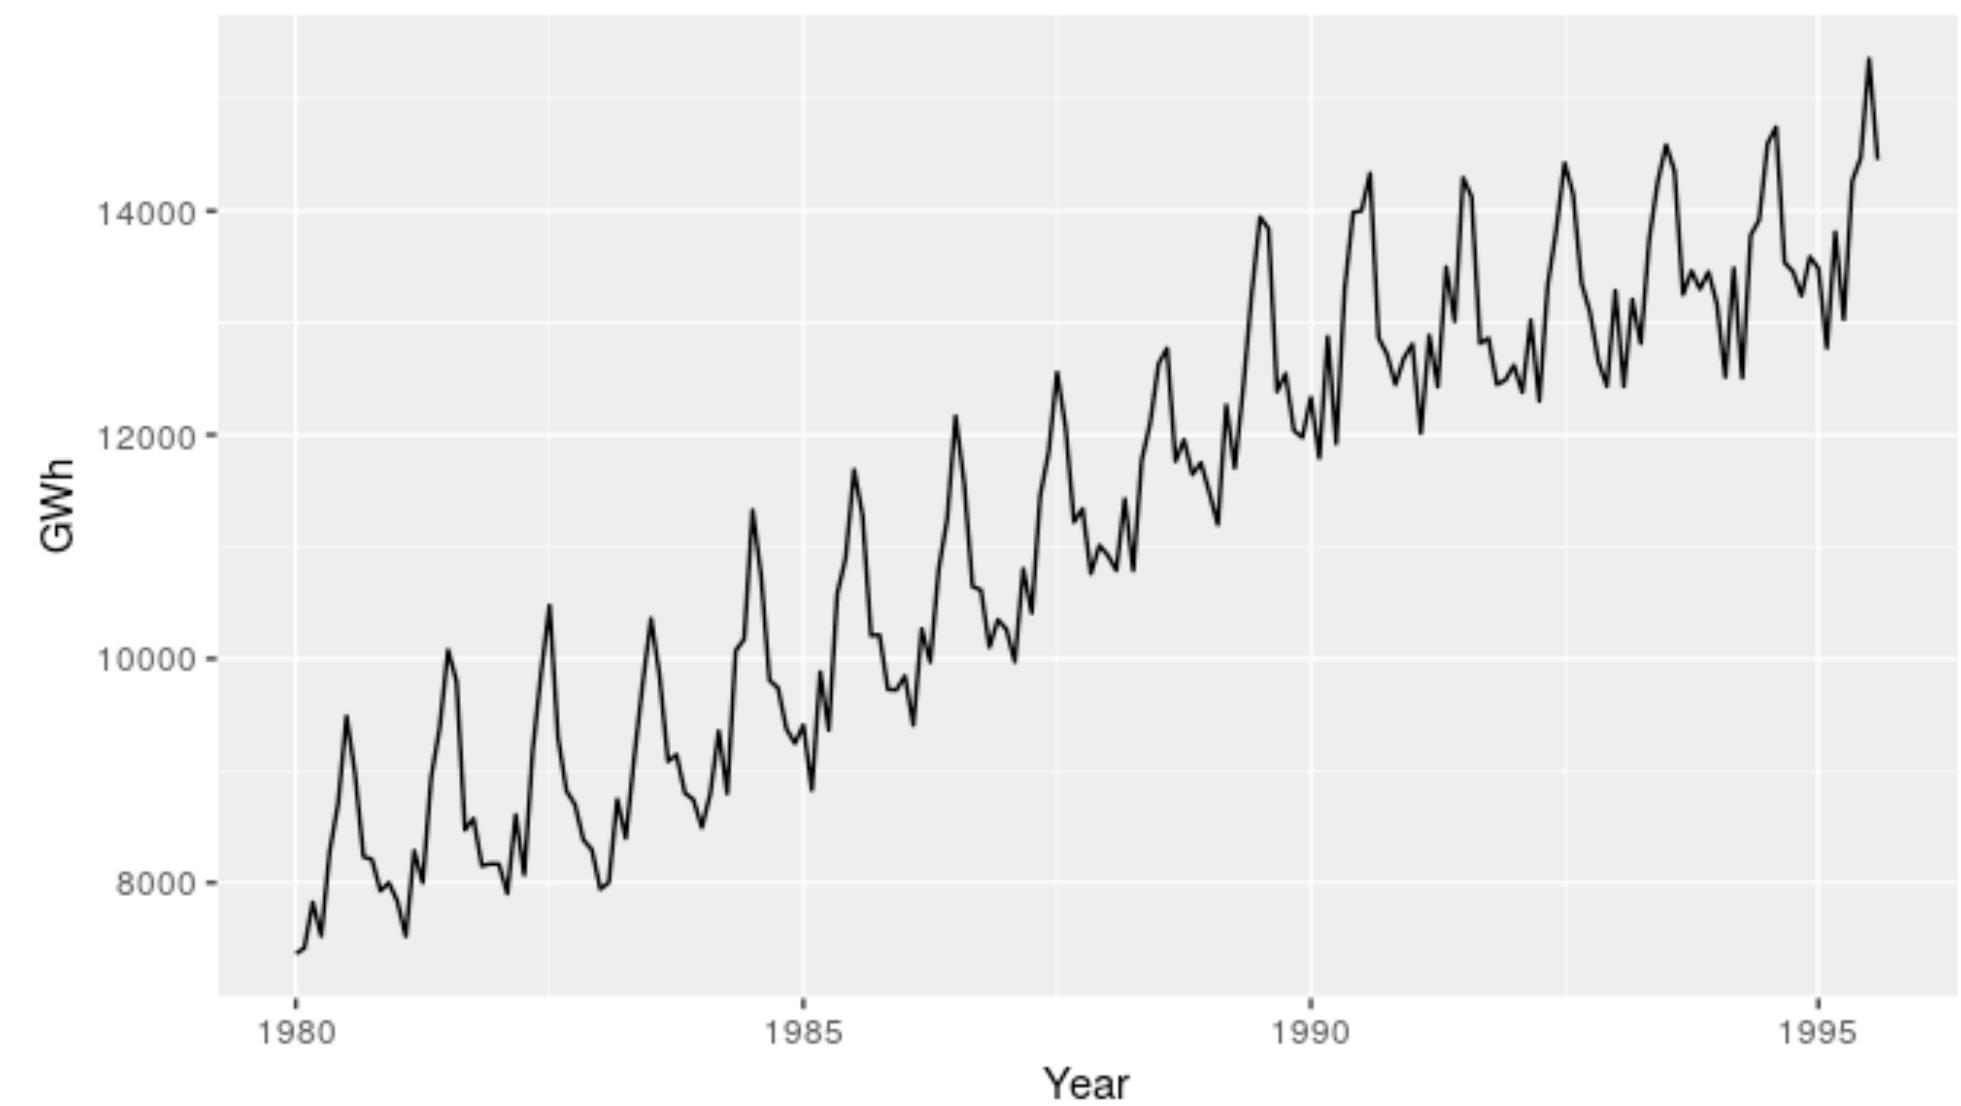
SHEWHART CONTROL CHARTS



Feature	Value
7 consecutive increasing	0
# > σ	1
# > $3^*\sigma$	1
5 consecutive above μ	1

AUTOCORRELATION

$$r_k = \frac{\sum_{t=k+1}^T (y_t - \bar{y})(y_{t-k} - \bar{y})}{\sum_{t=1}^T (y_t - \bar{y})^2},$$

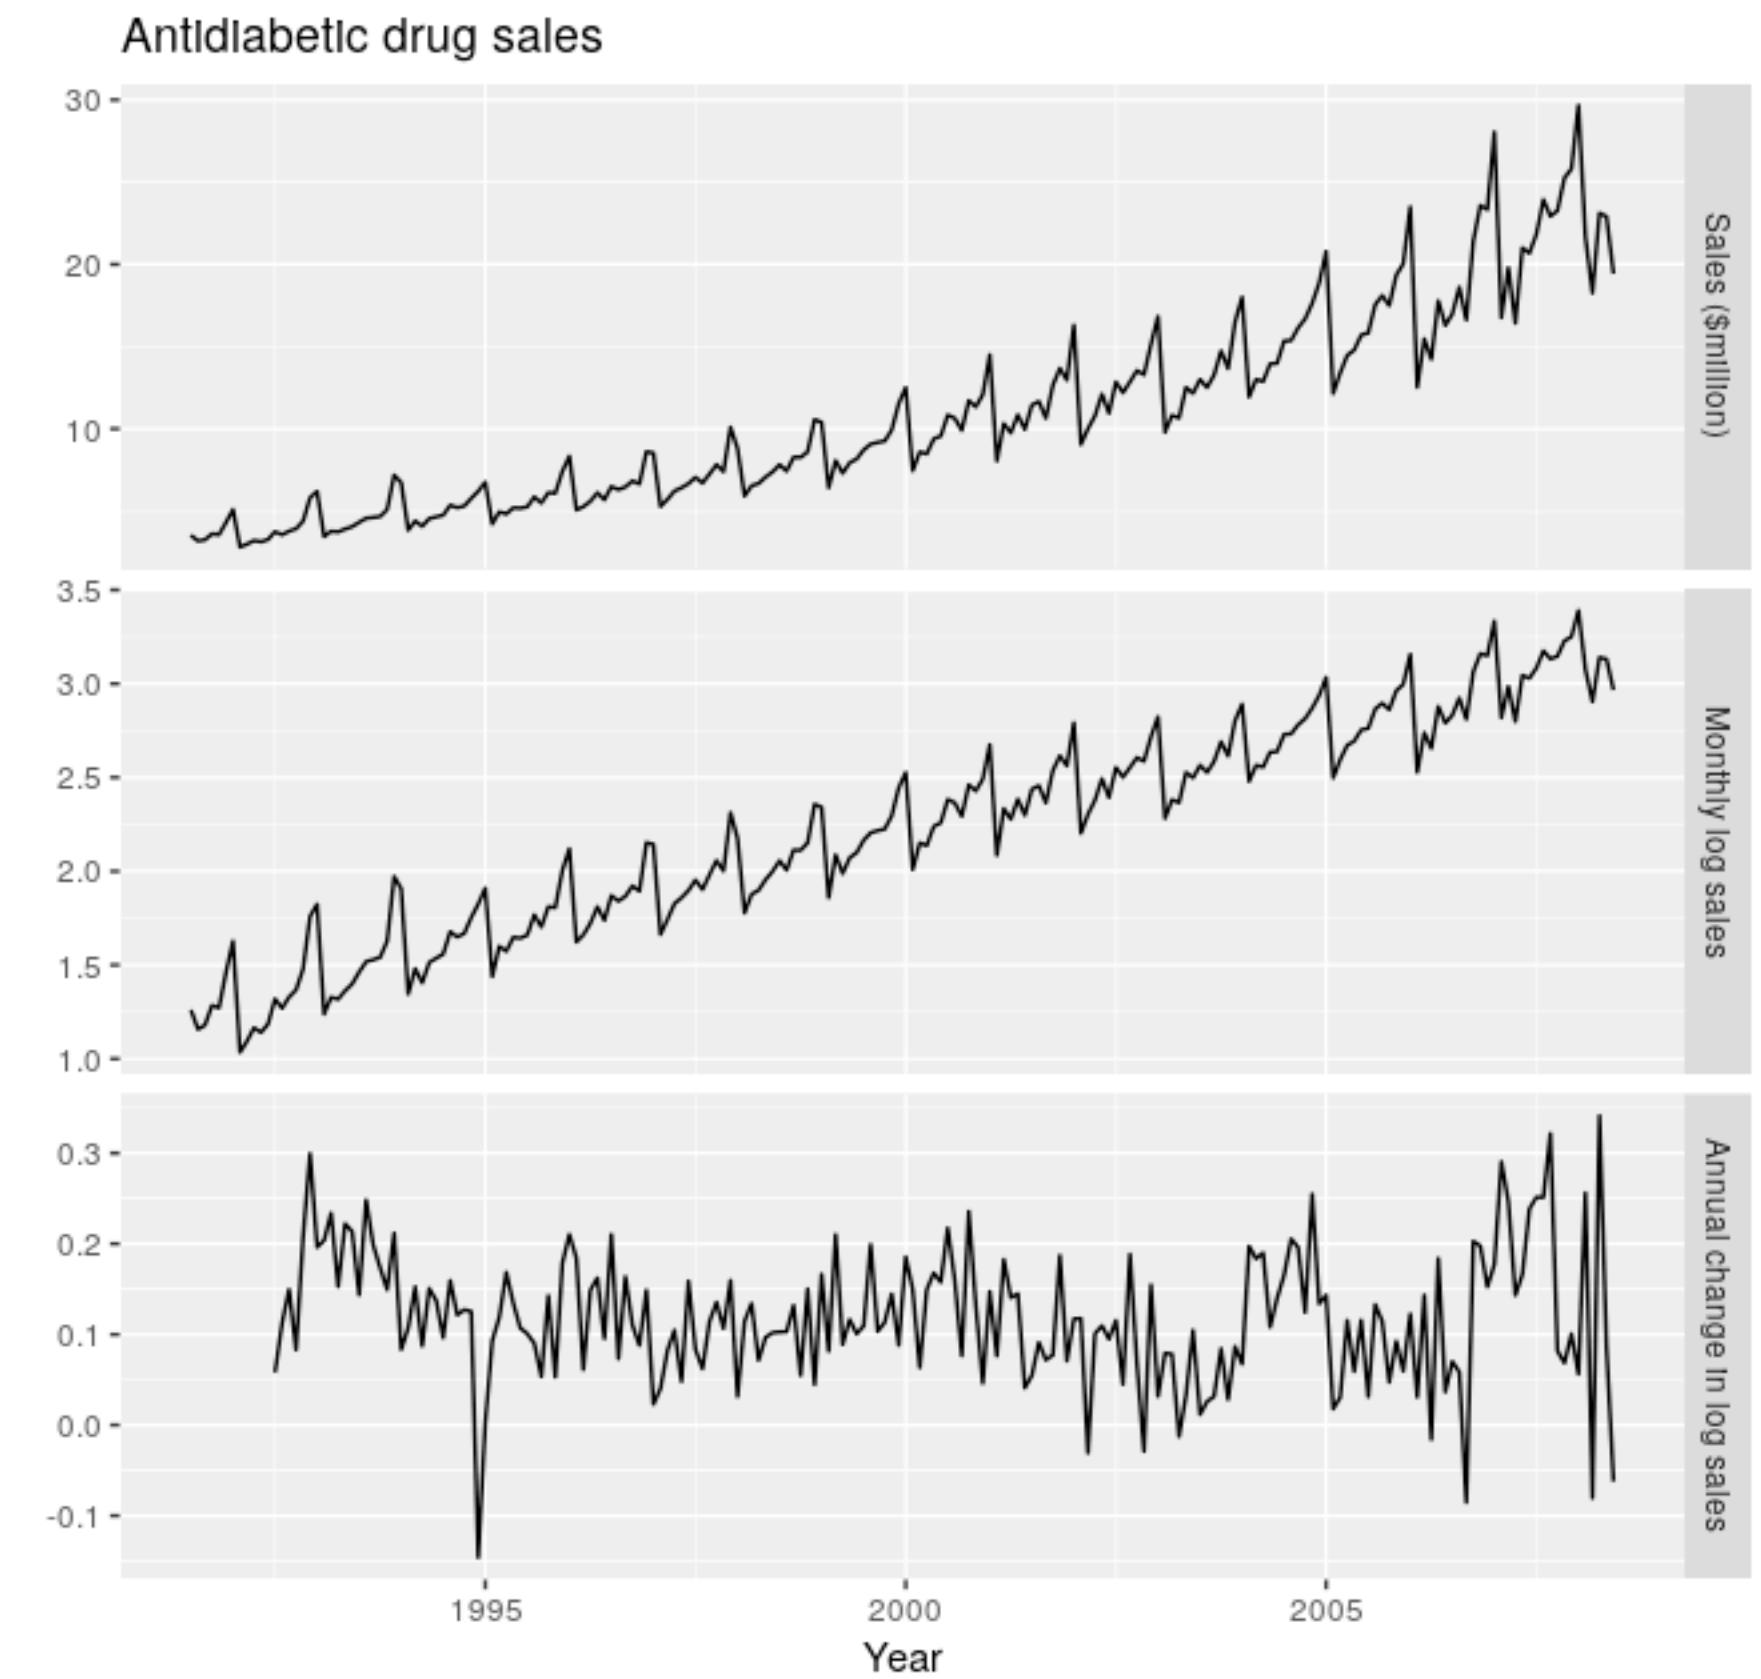


Forecasting using ARIMA

AutoRegressive Integrated Moving Average.

$$y'_t = c + \phi_1 y'_{t-1} + \cdots + \phi_p y'_{t-p} + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t,$$

- **AR:** *Autoregression*. A model that uses the dependent relationship between an observation and some number of lagged observations.
- **I:** *Integrated*. The use of differencing of raw observations (e.g. subtracting an observation from an observation at the previous time step) in order to make the time series stationary.
- **MA:** *Moving Average*. A model that uses the dependency between an observation and a residual error from a moving average model applied to lagged observations.

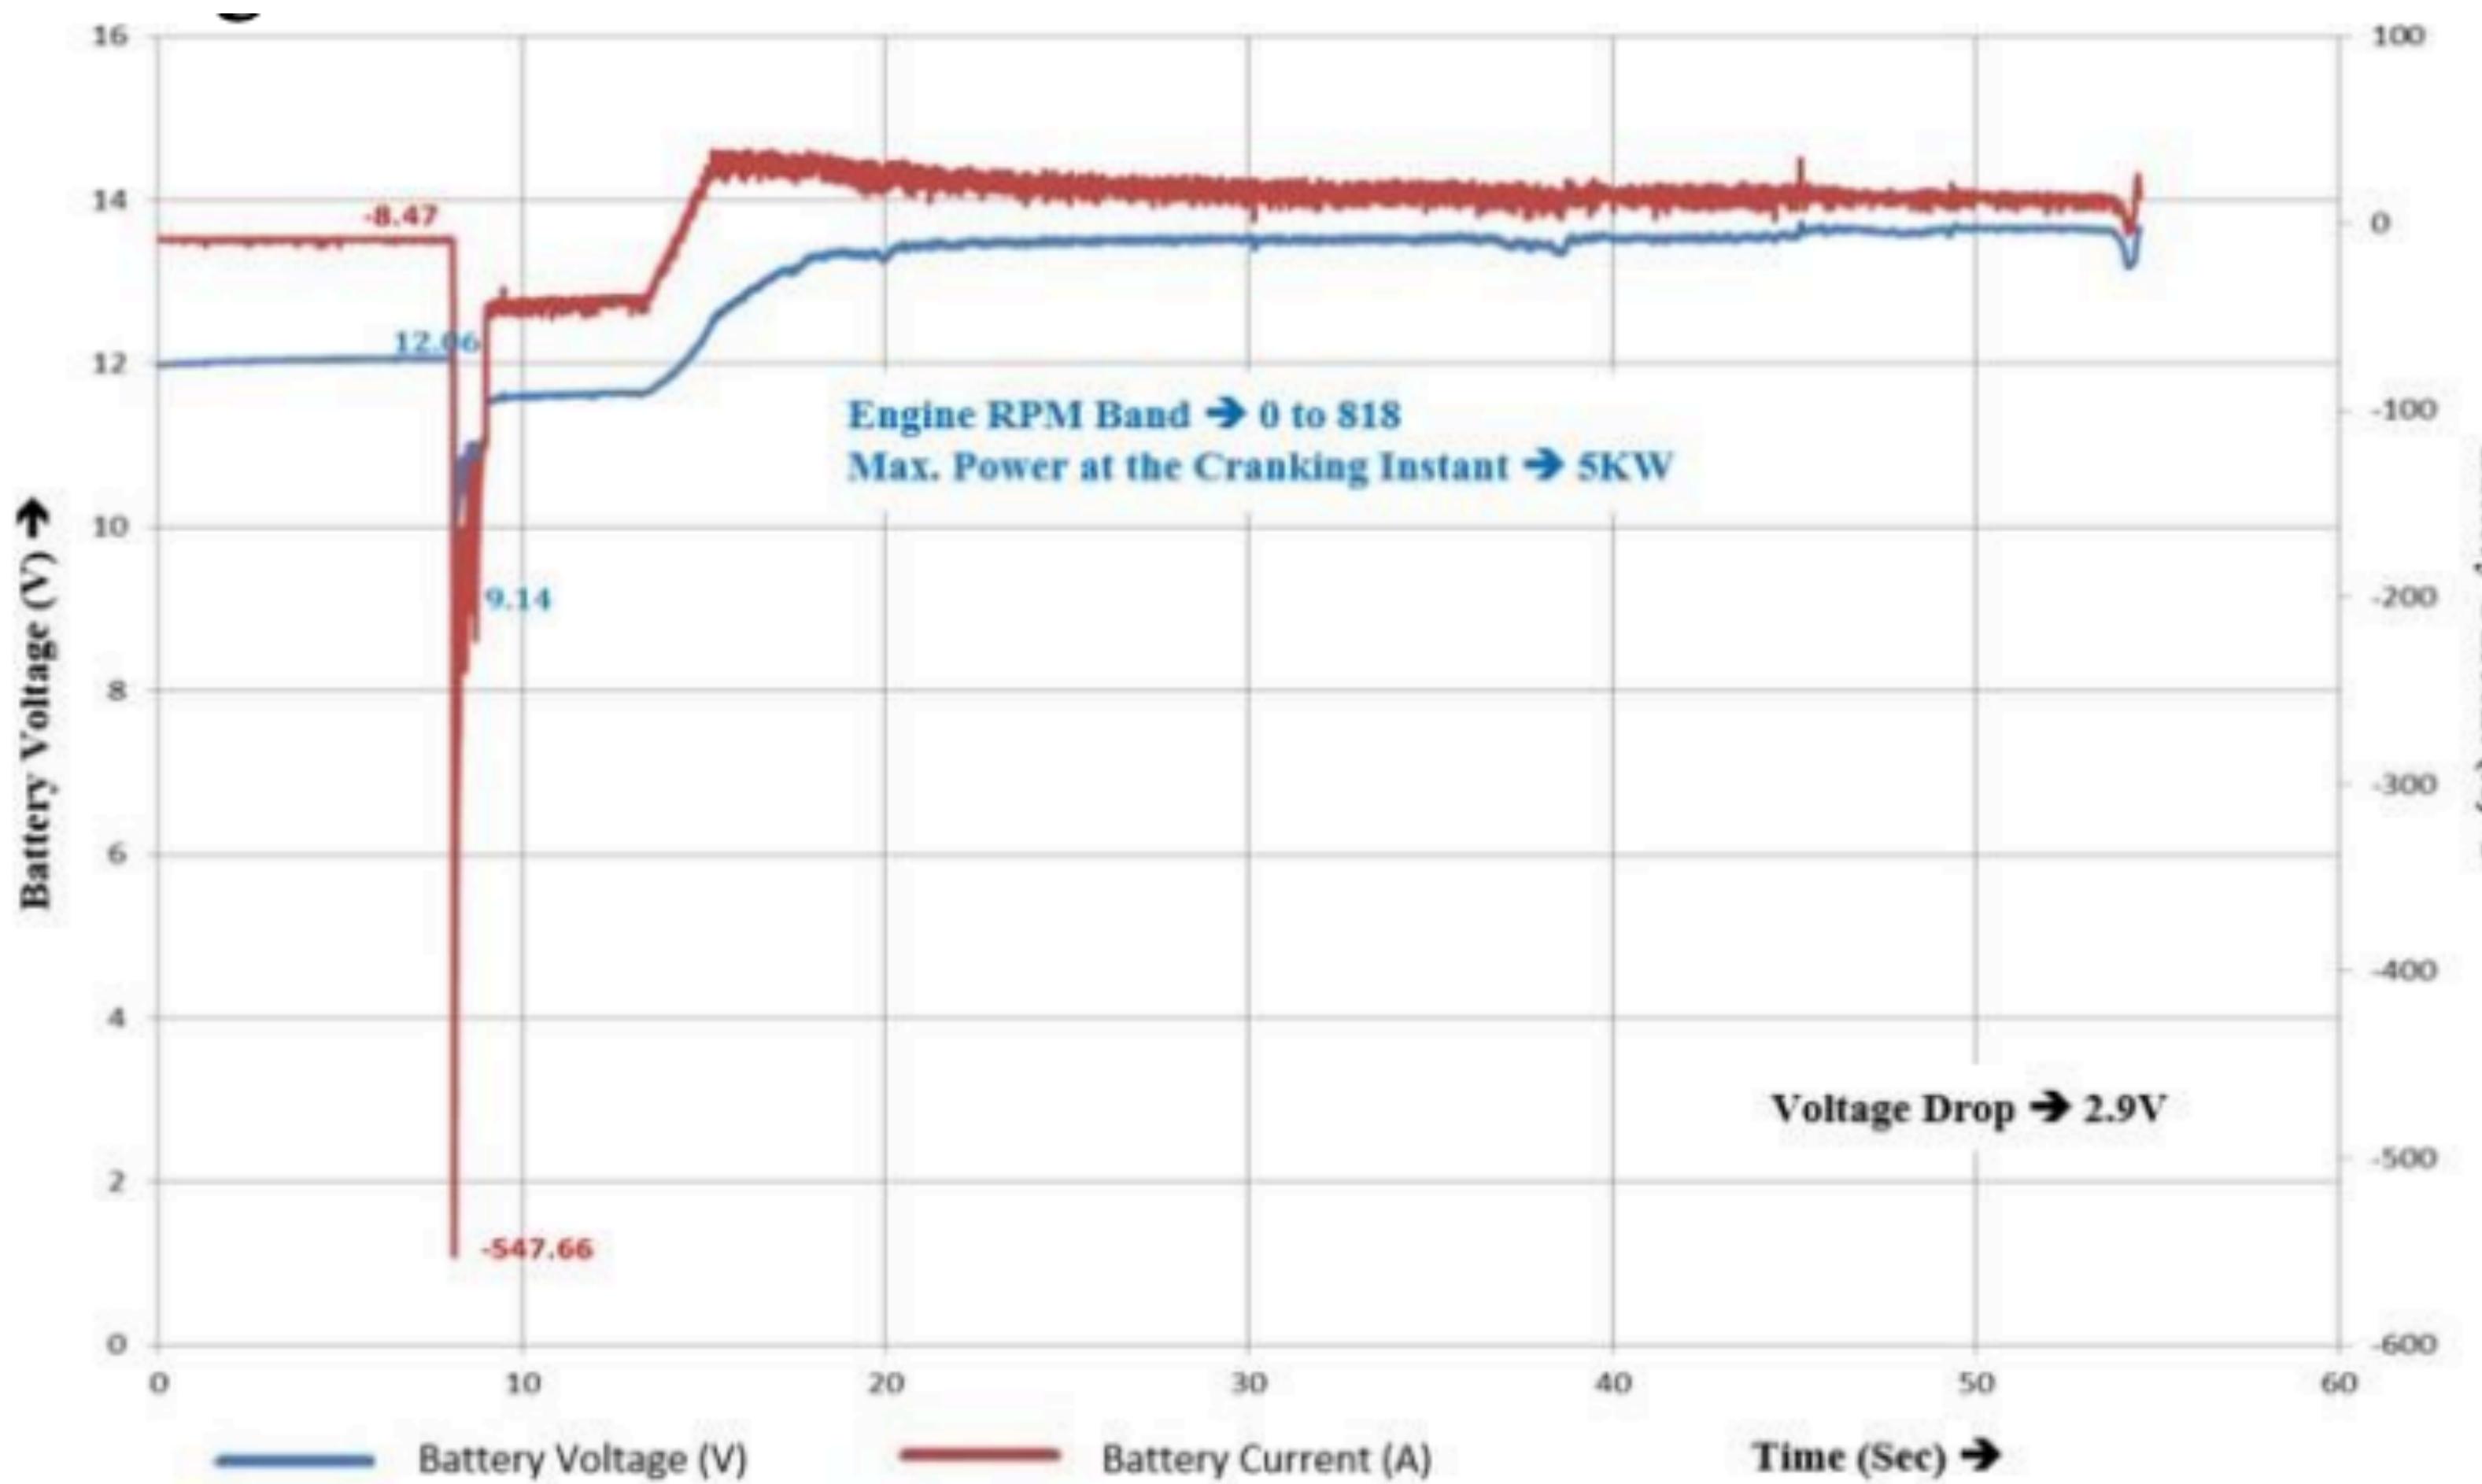


Learn more here: <https://otexts.com/fpp2/index.html>

Do you like features?

abs_energy(x)	Returns the absolute energy of the time series which is the sum over the squared values	linear_trend(x, param)	Calculate a linear least-squares regression for the values of the time series versus the sequence from 0 to length of the time series.
absolute_sum_of_changes(x)	Returns the sum over the absolute value of consecutive changes in the series x	linear_trend_timewise(x, param)	Calculate a linear least-squares regression for the values of the time series versus the sequence from 0 to length of the time series minus one.
agg_autocorrelation(x, param)	Calculates the value of an aggregation function (e.g.	longest_strike_above_mean(x)	Returns the length of the longest consecutive subsequence in x that is bigger than the mean of x
agg_linear_trend(x, param)	Calculates a linear least-squares regression for the values of the time series versus the sequence from 0 to length of the time series.		quence in x that is smaller than the mean of x
approximate_entropy(x, m, r)	Implements a vectorized version of approximate entropy		$\sum_{i=0}^{n-m} \sum_{j=i+1}^{i+r} \exp(-d(x_i, x_j)) = 0 \}$ estimated from polynomial ,
ar_coefficient(x, param)	This feature calculator fits the unconditional maximum likelihood of an autoregressive AR(k) process.	mean(^)	Returns the mean of x
augmented_dickey_fuller(x, param)	The Augmented Dickey-Fuller test is a hypothesis test which checks whether a unit root is present in a time series sample.	mean_abs_change(x)	Returns the mean over the absolute differences between subsequent time series values which is
autocorrelation(x, lag)	Calculates the autocorrelation of the specified lag, according to the formula [1]	mean_change(x)	Returns the mean over the differences between subsequent time series values which is
binned_entropy(x, max_bins)	First bins the values of x into max_bins equidistant bins.	mean_second_derivative_central(x)	Returns the mean value of a central approximation of the second derivative
c3(x, lag)	This function calculates the value of	median(x)	Returns the median of x
change_quantiles(x, ql, qh, isabs, f_agg)	First fixes a corridor given by the quantiles ql and qh of the distribution of x.	minimum(x)	Calculates the lowest value of the time series x.
cid_ce(x, normalize)	This function calculator is an estimator for the count of unique values in a time series (valleys etc.).	number_crossing_m(x, m)	Calculates the number of crossings of x on m.
count_above_mean(x)	Returns the number of values in x that are greater than the mean of x.	number_cwt_peaks(x, n)	This feature calculator searches for different peaks in x.
count_below_mean(x)	Returns the number of values in x that are smaller than the mean of x.	number_peaks(x, n)	Calculates the number of peaks of at least support n in the time series x.
cwt_coefficients(x, param)	Calculates a Continuous wavelet transform for the Ricker wavelet, also known as the “Mexican hat wavelet” which is	percentage_of_reoccurring_values_to_all_values(x)	correlation function at the given lag.
energy_ratio_by_chunks(x, param)	Calculates the sum of squares of chunk i out of N chunks expressed as a ratio with the sum of squares over the whole series.	quantile(x, q)	Returns the ratio of unique values, that are present in the time series more than once.
fft_aggregated(x, param)	Returns the spectral centroid (mean), variance, skew, and kurtosis of the absolute fourier transform spectrum.	range_count(x, min, max)	Calulates the q quantile of x.
fft_coefficient(x, param)	Calculates the fourier coefficients of the one-dimensional discrete Fourier Transform for real input by fast	ratio_beyond_r_sigma(x, r)	Count observed values within the interval [min, max].
first_location_of_maximum(x)	Returns the first location of the maximum value of x.	ratio_value_number_to_time_series_length(x)	Ratio of values that are more than $r * \text{std}(x)$ (so r sigma) away from the mean of x.
first_location_of_minimum(x)	Returns the first location of the minimal value of x.	sample_entropy(x)	Returns a factor which is 1 if all values in the time series occur only once, and below one if this is not the case.
friedrich_coefficients(x, param)	Coefficients of polynomial , which has been fitted to	set_property(key, value)	Calculate and return sample entropy of x.
has_duplicate(x)	Checks if any value in x occurs more than once	skewness(x)	This method returns a decorator that sets the property key of the function to value
has_duplicate_max(x)	Checks if the maximum value of x is observed more than once	spkt_welch_density(x, param)	Returns the sample skewness of x (calculated with the adjusted Fisher-Pearson standardized moment coefficient G1).
has_duplicate_min(x)	Checks if the minimal value of x is observed more than once	standard_deviation(x)	This feature calculator estimates the cross power spectral density of the time series x at different frequencies.
index_mass_quantile(x, param)	Those apply features calculate the relative index i where q% of the mass of the time series x lie left of i.	sum_of_reoccurring_data_points(x)	Returns the standard deviation of x
kurtosis(x)	Returns the kurtosis of x (calculated with the adjusted Fisher-Pearson standardized moment coefficient G2).	sum_of_reoccurring_values(x)	Returns the sum of all data points, that are present in the time series more than once.
large_standard_deviation(x, r)	Boolean variable denoting if the standard dev of x is higher than 'r' times the range = difference between max and min of x.	sum_values(x)	Returns the sum of all values, that are present in the time series more than once.
last_location_of_maximum(x)	Returns the relative last location of the maximum value of x.	symmetry_looking(x, param)	Calculates the sum over the time series values
last_location_of_minimum(x)	Returns the last location of the minimal value of x.	time_reversal_asymmetry_statistic(x, lag)	Boolean variable denoting if the distribution of x looks symmetric.
		value_count(x, value)	This function calculates the value of
		variance(x)	Count occurrences of value in time series x.
		variance_larger_than_standard_deviation(x)	Returns the variance of x
			Boolean variable denoting if the variance of x is greater than its standard deviation.

TALK TO SUBJECT MATTER EXPERTS



<https://ijsea.com/archive/volume7/issue8/IJSEA07081005.pdf>

Sources of features

1. Univariate features
2. Think about the problem and translate to math
3. Fast Fourier Transform
4. Timeseries features
5. Subject matter experts

MORE MODELS

Specific applications require specific models.

Kaplan Meier estimator (product limit)

The **estimator** of the survival function $S(t)$ (the probability that life is longer than t) is given by:

$$\widehat{S}(t) = \prod_{i: t_i \leq t} \left(1 - \frac{d_i}{n_i}\right),$$

with t_i a time when at least one event happened, d_i the *number of events* (i.e., deaths) that happened at time t_i and n_i the *individuals known to have survived* (have not yet had an event or been censored) up to time t_i .

https://en.wikipedia.org/wiki/Kaplan-Meier_estimator

Cox Proportional Hazards model

Hazard function

$$h(t) = \lim_{\Delta t \rightarrow 0} \frac{P(t \leq T < t + \Delta t | T \geq t)}{\Delta t}$$

Baseline hazard

$$h_0(t)$$

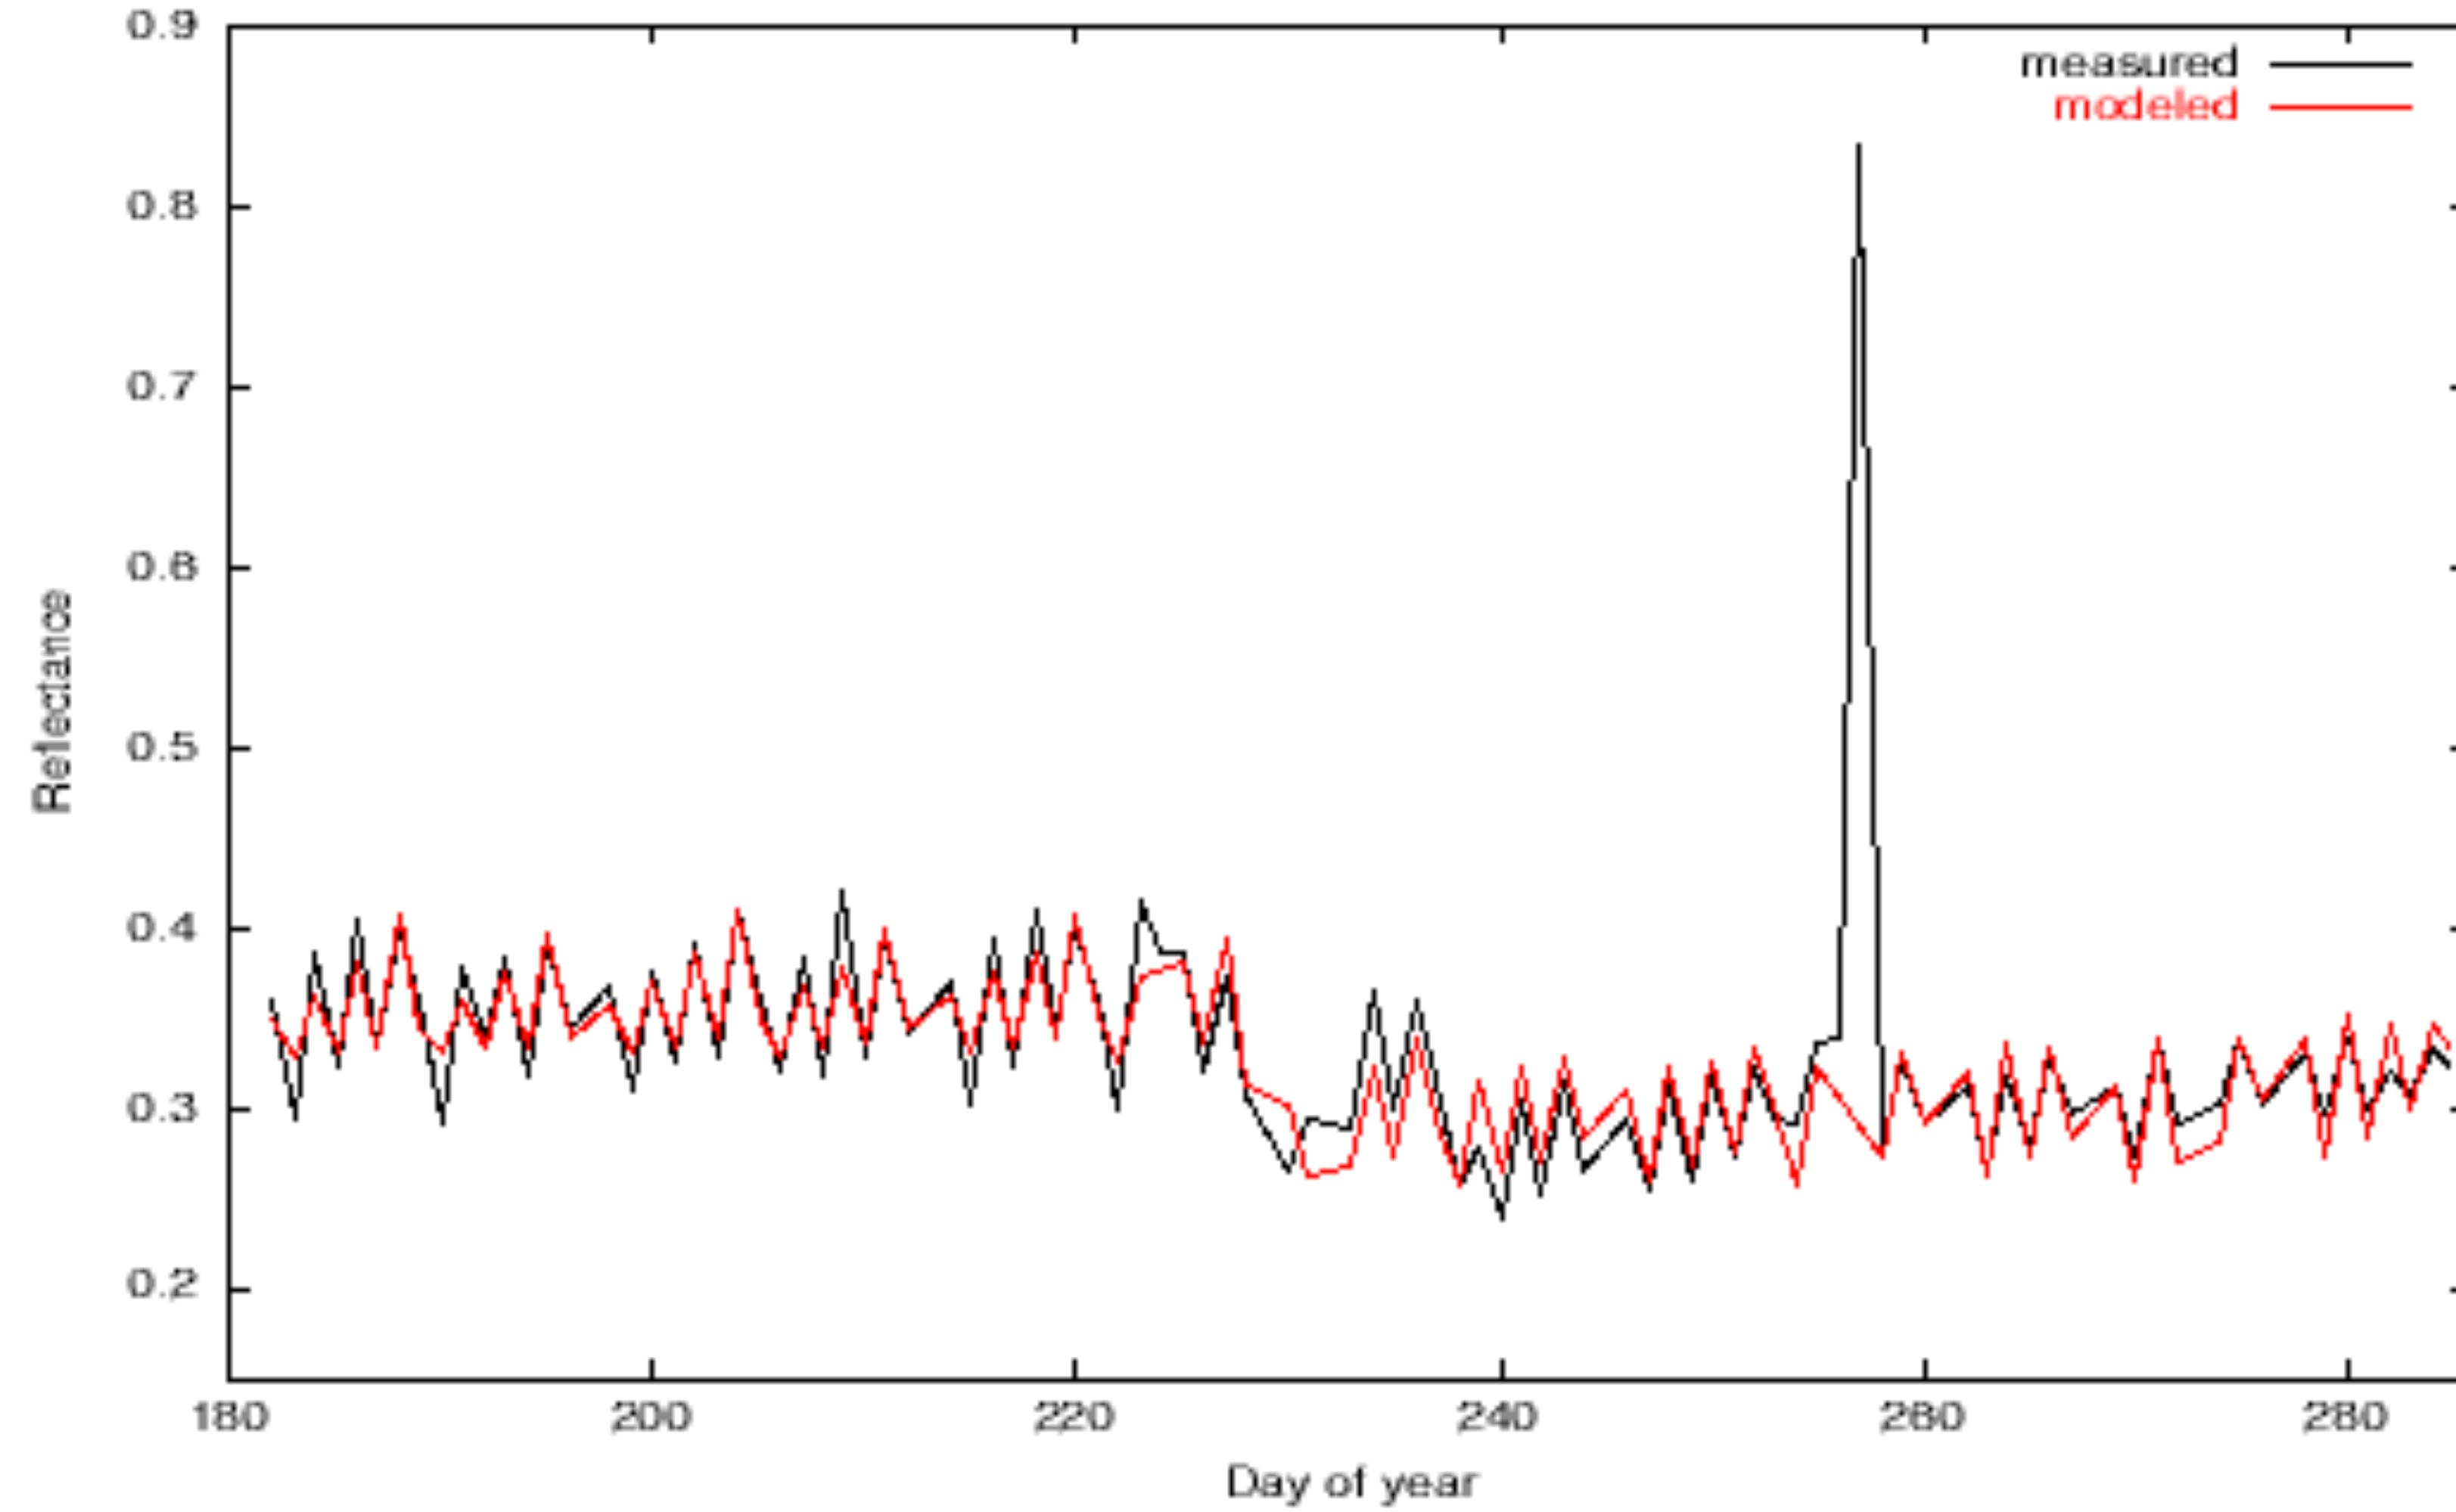
Hazard of individual i

$$h_i(t) = h_0(t) e^{\beta_1 X_{i1} + \dots + \beta_p X_{ip}}$$

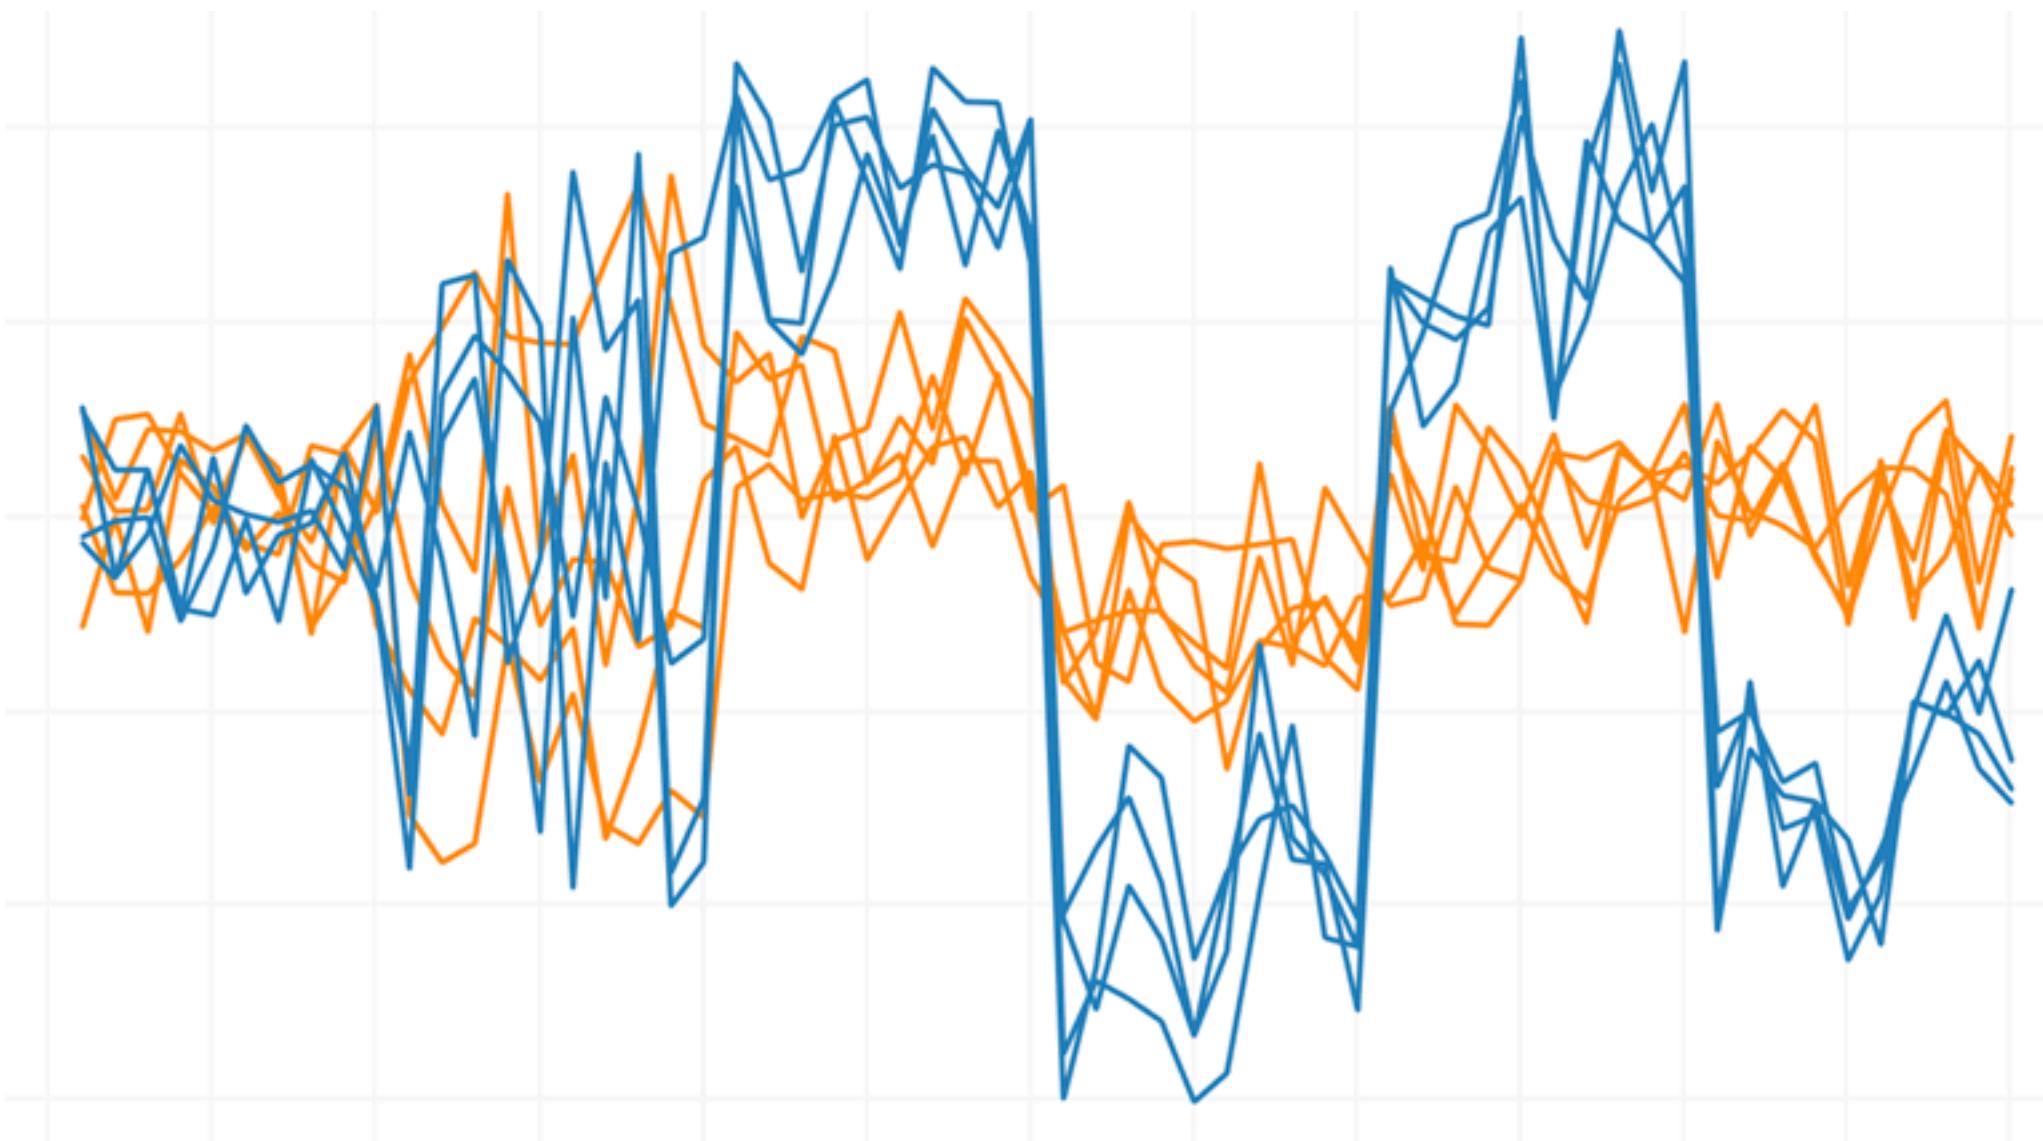
Time independent covariates

Advanced anomaly detection

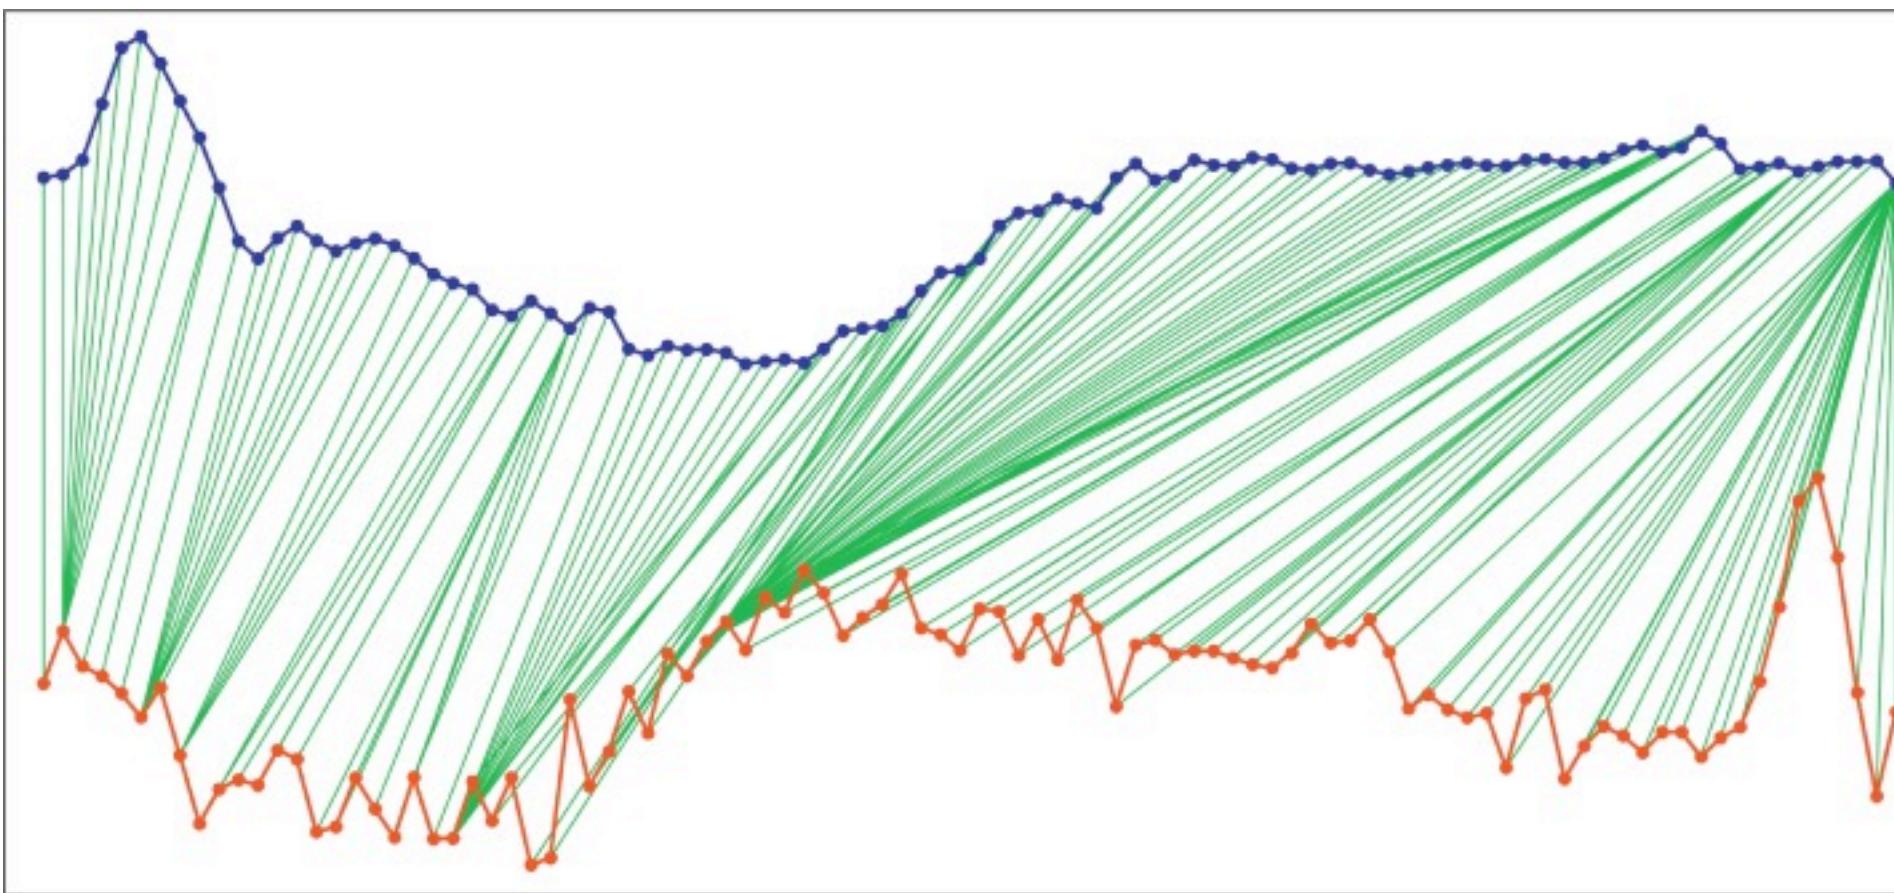
It's just forecasting! But with an extra step



Time Series Clustering and Regression



- k-means but with an L2 norm as the distance metric!
- Dynamic Time Warping
- Kernel prediction methods
- ...



DEEP LEARNING BASED APPROACHES

Basic Neuron Model In A Feedforward Network

» Inputs x_i arrive through pre-synaptic connections

» Synaptic efficacy is modeled using real **weights** w_i

» The response of the neuron is a **nonlinear function** f of its weighted inputs

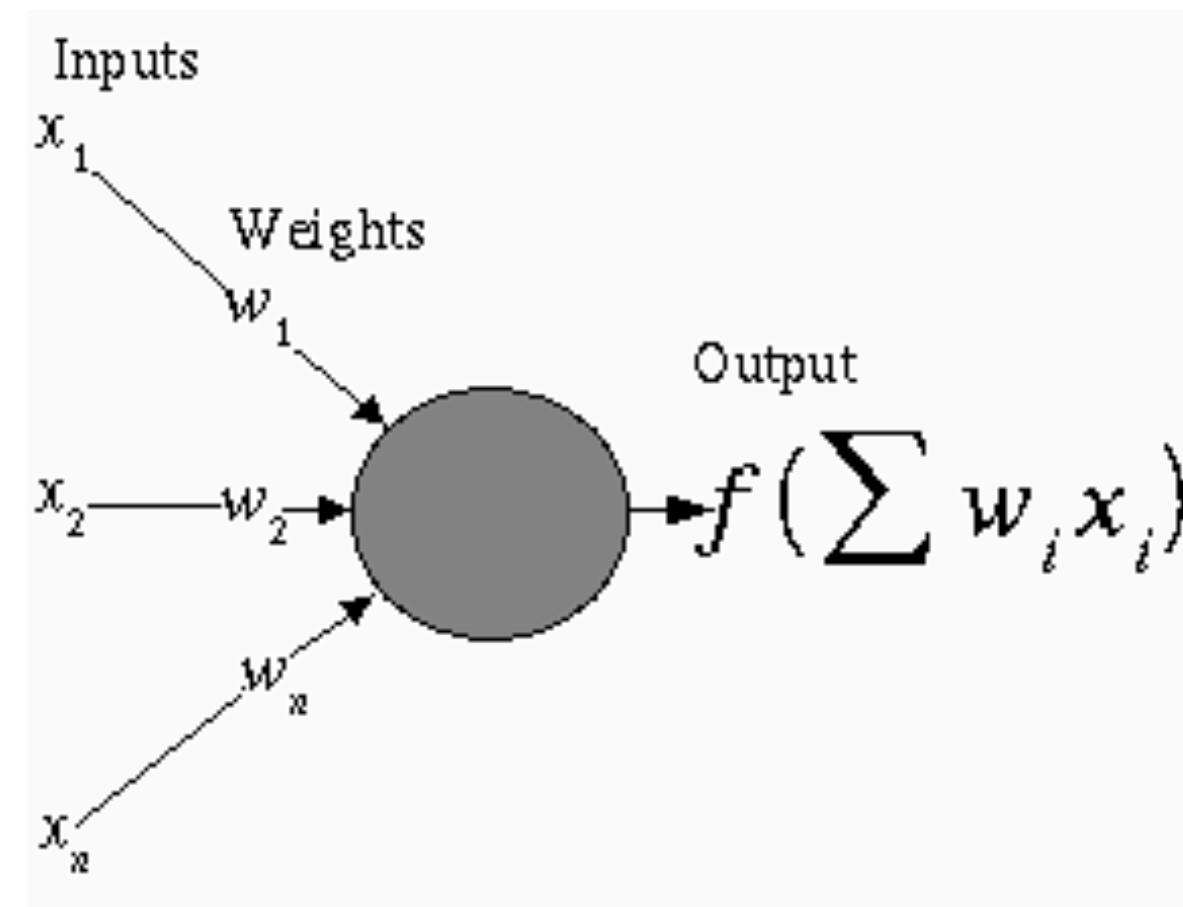
$$TSSE = \frac{1}{2} \sum_{\text{patterns}} \sum_{\text{outputs}} (\text{desired} - \text{actual})^2$$

Loss function: $RMSE = \sqrt{\frac{2 * TSSE}{\# \text{patterns} * \# \text{outputs}}}$

Neuron error $\delta_{pj} = (T_{pj} - O_{pj}) O_{pj} (1 - O_{pj})$

T_{pj} is the target value of output neuron j for pattern p

O_{pj} is the actual output value of output neuron j for pattern p



$$f(x) = \frac{1}{1 + e^{-\lambda x}}$$

High level algorithm

1. Calculate loss with current weights W_{ji}
2. Compute weight adjustments ΔW_{ji} at time t by

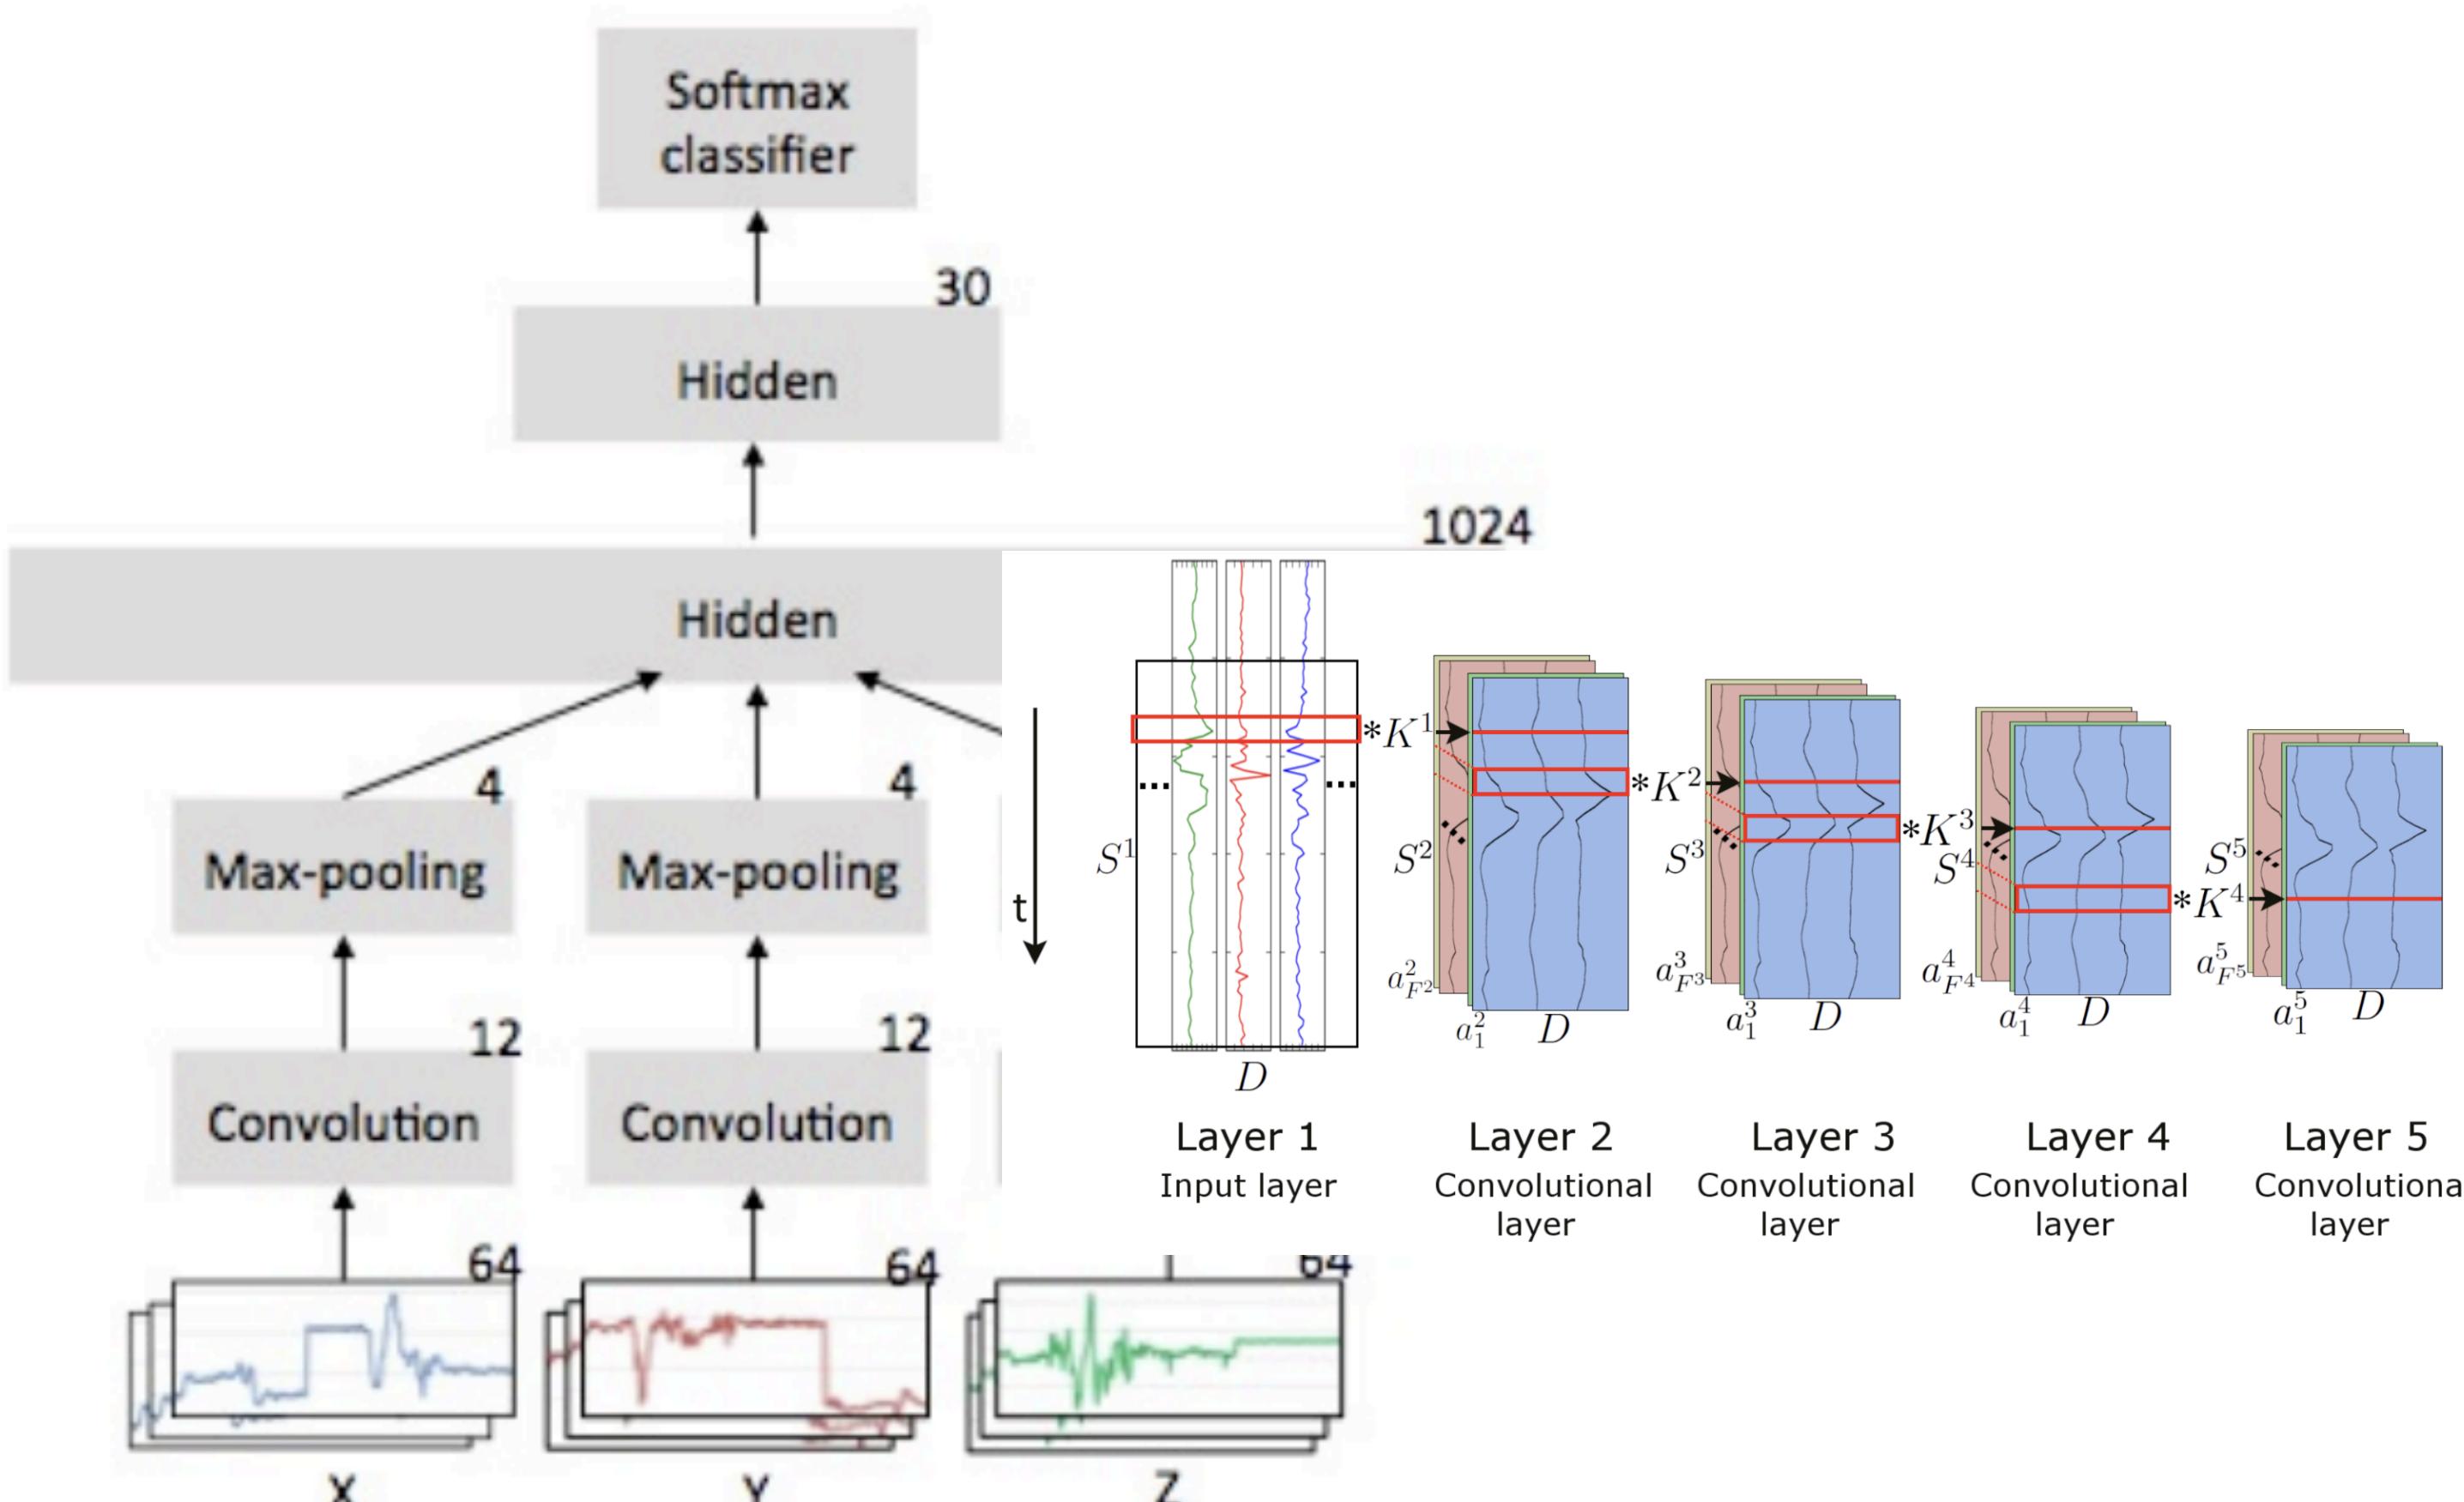
$$\Delta W_{ji}(t) = \eta \delta_{pj} O_{pi}$$

3. Apply weight adjustments according to

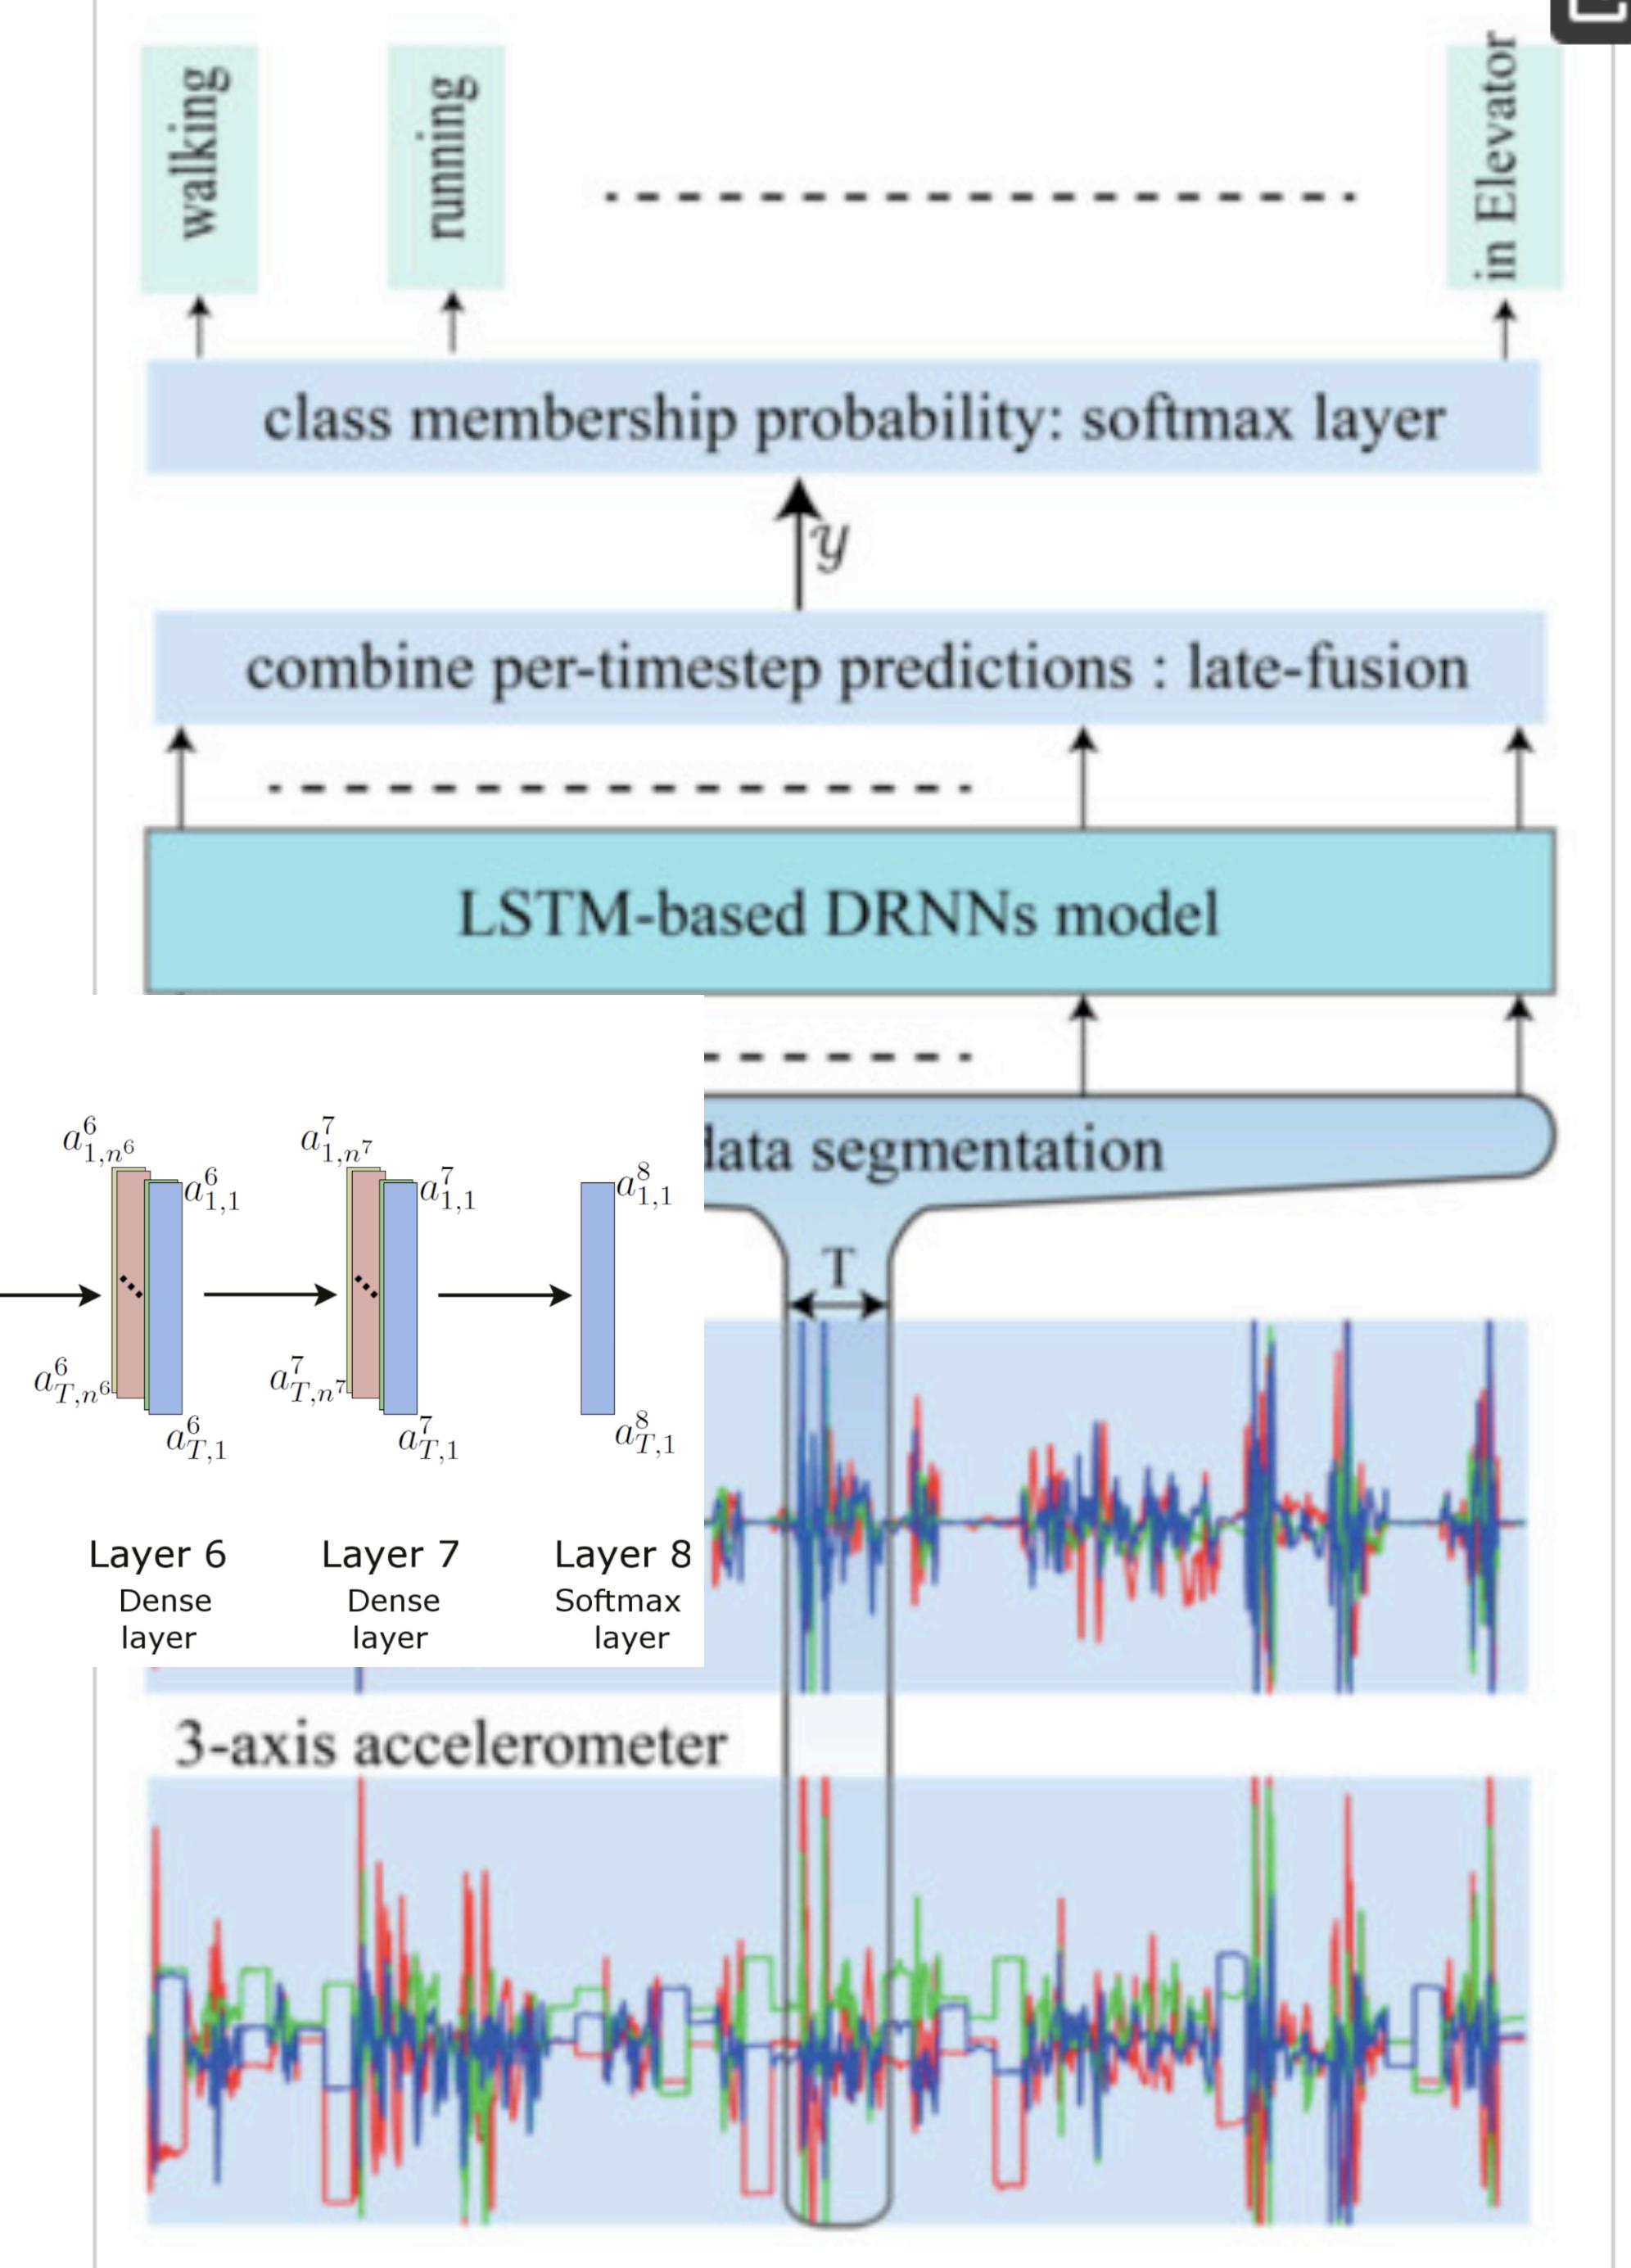
$$W_{ji}(t+1) = W_{ji}(t) + \Delta W_{ji}(t)$$

4. Iterate until loss is acceptable or convergence

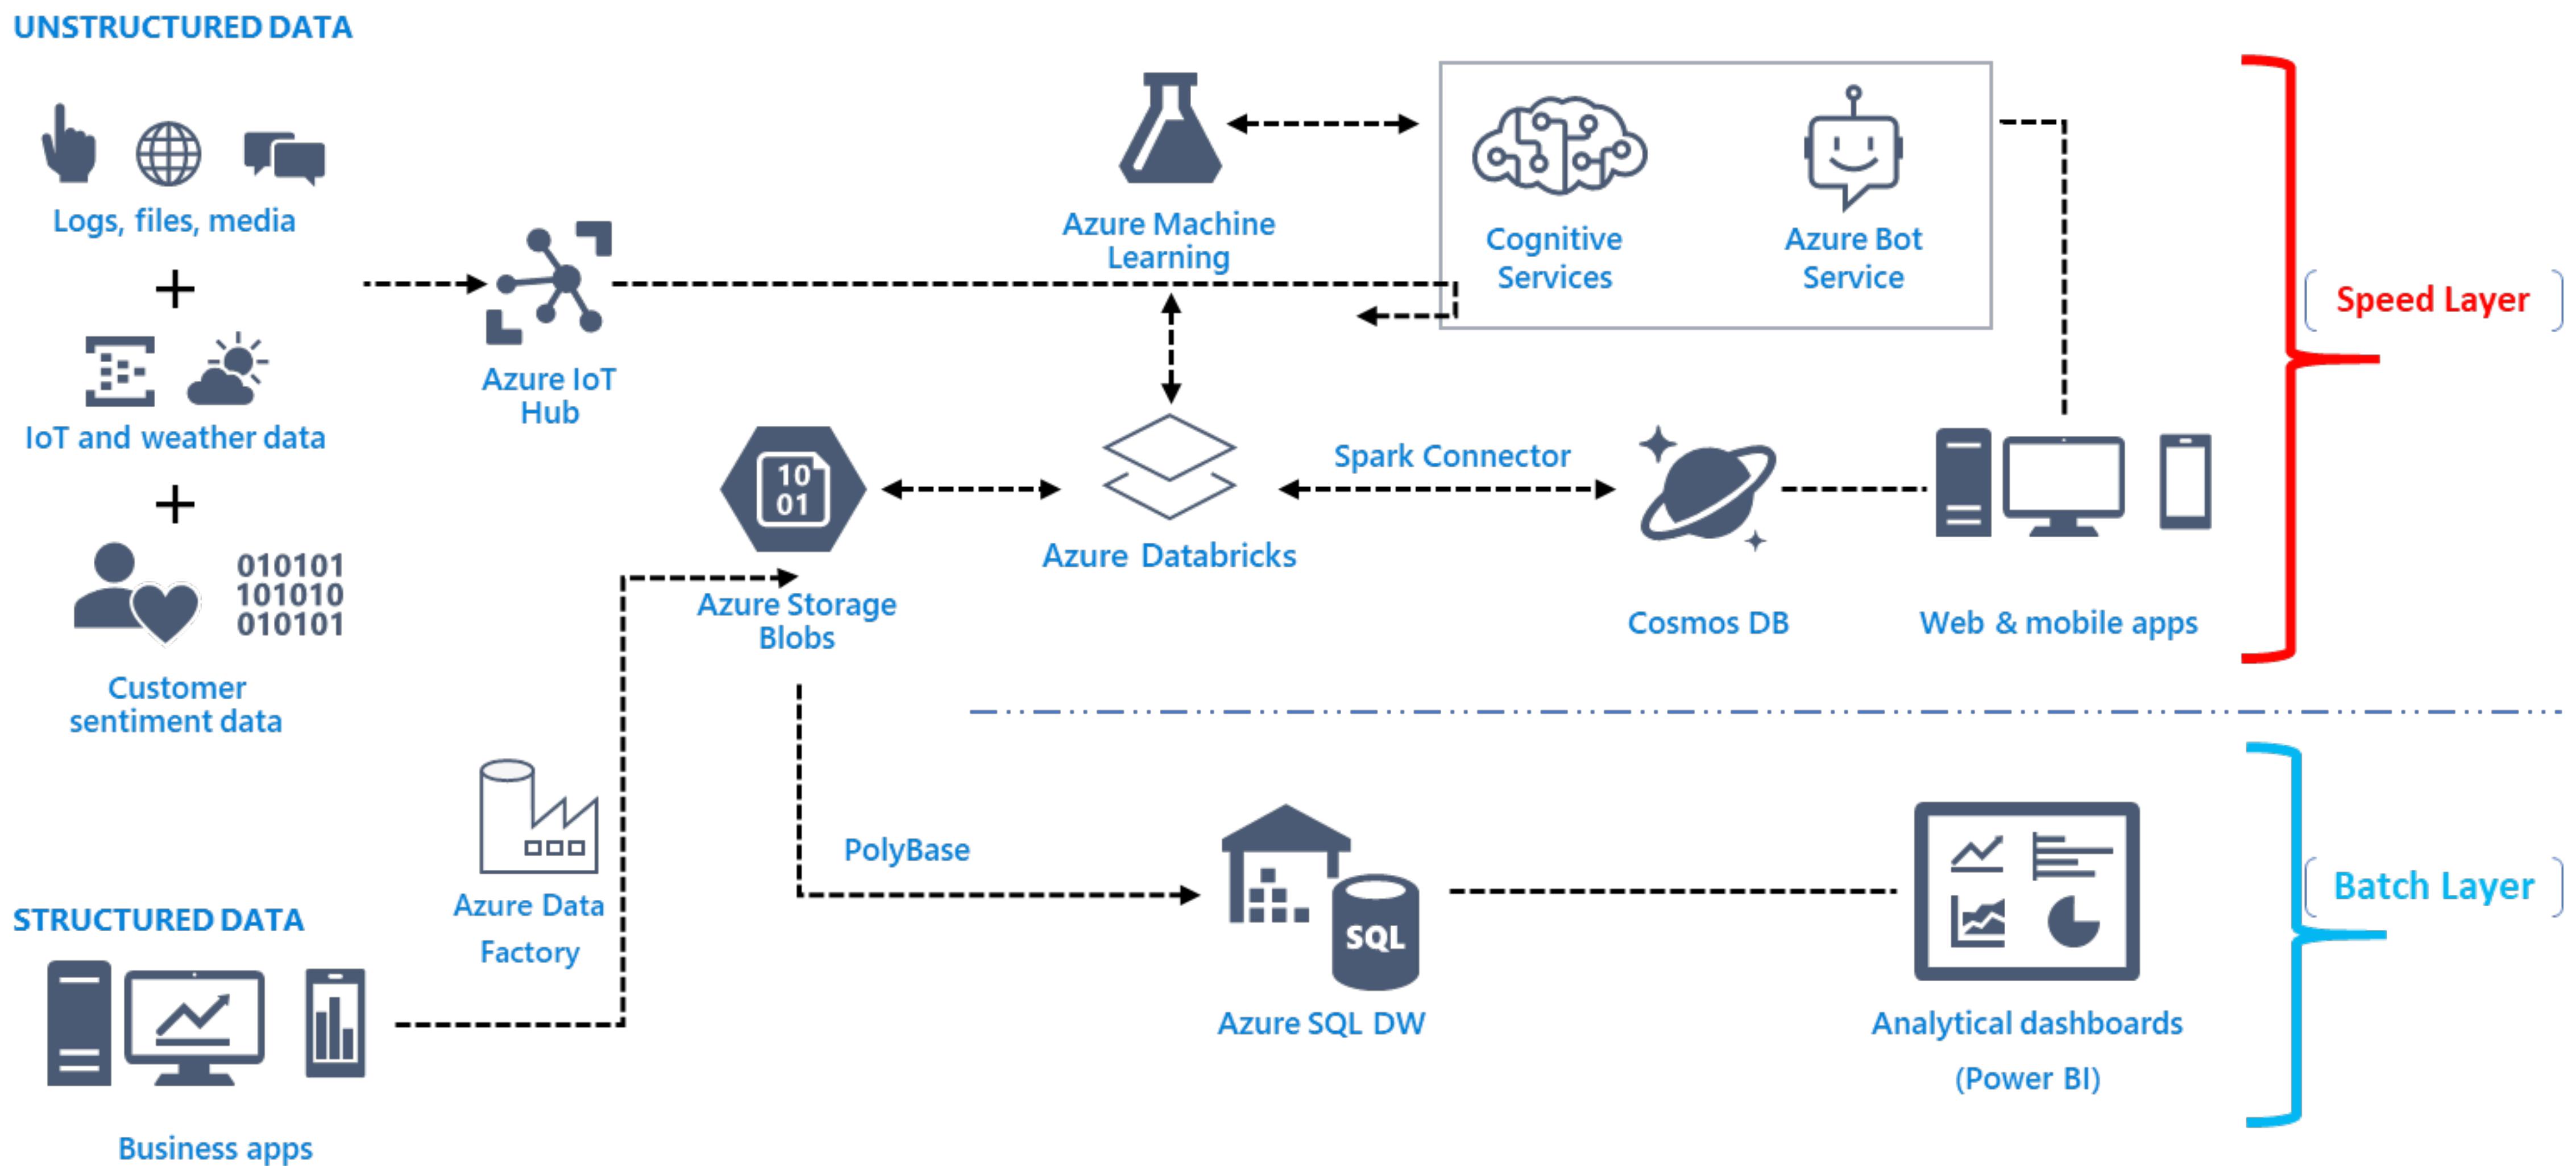
CONVOLUTIONAL AND RECURRENT



<https://machinelearningmastery.com/deep-learning-models-for-human-activity-recognition/>



Technical Architecture



THIS IS ONLY TOP OF THE ICEBERG.
ARE YOU GETTING MORE
CURIOUS?

CAN YOU CODE?

START AN AMAZING CAREER AT FAKTION

Get in touch!



Oudeleeuwenrui 39
2000 Antwerpen
www.faktion.com

ARE YOU GRADUATING

AND DO YOU LIKE MACHINE LEARNING?

Get in touch! Seriously, we can't imagine why you would miss the opportunity to get to know us better in a face-to-face conversation. It's the only way to see how we could fit. And for sure, at least we'll have an interesting chat.

ARE YOU A STUDENT

AND HAVE AN INTEREST IN MACHINE LEARNING?

Work with real-world data and gain experience on how the industry uses Machine Learning right now. This means we won't let you do academic research and we go beyond the typical applications of the big tech players. Instead, you'll work with experts and help implement Machine Learning models, projects and products that will help real people and companies in Belgium and beyond.

MASTER/BACHELOR THESIS

INTERNSHIP

PAID SUMMER JOB

THANK YOU



FAKTION
May 14, 2019